The Deployment Gap: Why Benchmark Accuracy Fails to Predict Production Readiness

Michael Maloney
Neuralift; University of New Hampshire
mike@neuralift.ai

Abstract

There's a dirty secret in model selection: the benchmarks everyone uses to pick LLMs have almost nothing to do with whether those models will actually work in production.

We ran 13 models through structured output tasks with a 2-second deadline. The results were stark. Accuracy and deployment success don't just fail to correlate—they actively diverge. The smallest model in our study, Llama-3.2-1B, achieved the highest production success rate. The largest models failed almost every request.

Model	**Vendor**	**Success@SLO**
Llama-3.2-1B	Meta	**49.1%** (best)
Qwen2.5-3B	Alibaba	34.3%
Phi-3-mini	Microsoft	33.8%
Llama-3.2-3B	Meta	31.8%
Qwen3-4B	Alibaba	29.9%
Ministral-8B	Mistral	28.5%
Llama-3.1-8B	Meta	13.2%
Mistral-7B-v0.3	Mistral	10.0%
Gemma-2-9B	Google	2.2%
GPT-OSS-20B	OpenAI	0.6%
Yi-1.5-6B	01.AI	0.4%
Falcon-Mamba-7B	TII	0.0%
Gemma-3-12B	Google	0.0% (worst)

The pattern holds across vendors and architectures: Meta, Alibaba, Microsoft, Google, OpenAI, Mistral, 01.AI, and TII models from the US, China, France, and UAE all show the same inversion. Bigger models that dominate accuracy leaderboards collapse under latency constraints.

We introduce *Success@SLO*—the fraction of requests that are both correct AND arrive on time. This single metric captures what production systems actually experience. We also release SpecSLOEval, a framework for measuring structure, accuracy, faithfulness, stability, and latency through a unified evaluation harness.

The implication is uncomfortable but clear: if you're selecting models by accuracy alone, you're selecting for failure.

1 Introduction

Most teams pick their production model the same way: find the highest-scoring option on MMLU or HELM that fits the memory budget, integrate it, and ship. It seems reasonable. Better accuracy should mean better results, right?

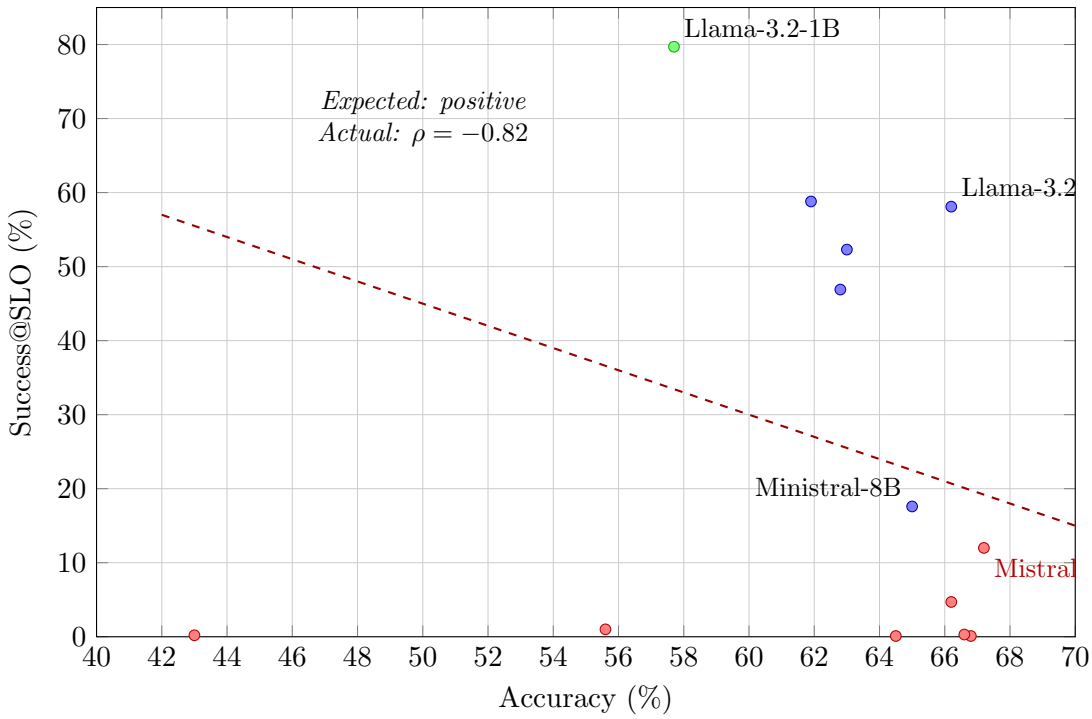

Figure 1: **The Deployment Gap.** Accuracy vs Success@SLO for 13 models under a 2-second SLO. The correlation is strongly negative ($\rho = -0.82$, $p < 0.001$). The smallest model (Llama-3.2-1B, 57.7% accuracy) achieves the best Success@SLO (79.7%). The accuracy leader (Mistral-7B-v0.3, 67.2%) achieves only 12% Success@SLO.

We tested this assumption with 13 models and 2,500 structured output tasks. What we found should make anyone rethink how they select models.

Ministral-8B achieves 84.7% accuracy on intent classification—impressive by any standard. Under a 2-second latency deadline, it delivers only 28.5% of requests successfully. The 1B Llama model, with its modest 27.3% accuracy on the same task, achieves 49.1% deployment success overall because it actually finishes in time. The accuracy leader fails; the speed leader wins.

This pattern repeats across every model we tested. Gemma-2-9B shows 79% accuracy but 2.2% deployment success. GPT-OSS-20B, OpenAI's first open-weights release, achieves solid accuracy but completes only 0.6% of requests within deadline. The correlation between accuracy and production success isn't just weak—it's negative.

We call this the **deployment gap**. Traditional benchmarks measure whether a model *can* answer correctly. Production systems need models that *will* answer correctly, quickly, and consistently. These are different questions, and current evaluation methods only answer the first one.

The consequences are real. Teams spend months integrating high-accuracy models, deploy them with confidence, and discover at 3am that their carefully selected model times out on 99% of requests. The benchmarks gave no warning because the benchmarks don't measure what matters.

We focus specifically on single-GPU deployments because this is where most teams actually run agents. Not everyone has a fleet of A100s. Many organizations run open-weight models on an RTX 4090 or M2 Max behind LM Studio or Ollama, partly for cost, partly to keep sensitive data on-premise. These setups face all the same production pressures as cloud deployments—schema contracts, faithfulness requirements, latency SLOs—but with tighter resource envelopes and more

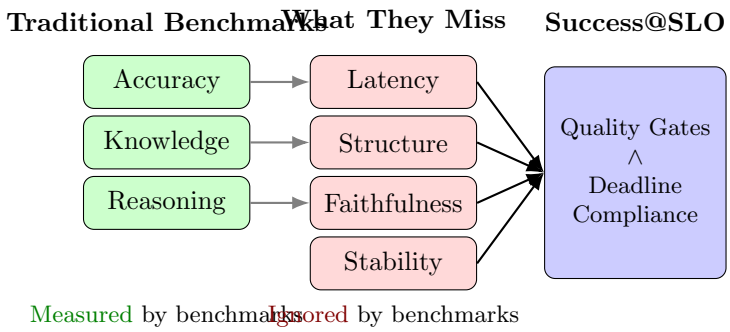

Figure 2: **What benchmarks miss.** Traditional LLM benchmarks measure accuracy, knowledge, and reasoning. They ignore latency, structural validity, faithfulness, and stability—exactly the factors that determine production success. Success@SLO captures the full picture.

visible tail behavior.

The research community has attacked pieces of this problem. Serving systems like vLLM [9] and SGLang [16] optimize throughput and tail latency, but treat the model as a black box emitting unconstrained text. Structured output tools [7, 14] improve syntactic validity, but evaluation stops at "does it parse?" What nobody has built is a framework that measures the full stack: structure *and* accuracy *and* faithfulness *and* stability *and* latency, all evaluated together against explicit contracts and deadlines.

What we contribute. Three things, in order of how much they will change your workflow:

1. **A spec compiler that actually works across backends.** Hand us a JSON Schema; we will turn it into whatever your serving stack can enforce, then validate the output anyway because we do not trust the enforcement to be complete. If validation fails, we retry with bounded budget. All of this gets reflected in your latency measurements, because pretending retries are free is how you end up surprised in production.

2. **Success@SLO as your deployment metric.** We define this as the fraction of requests that pass your quality gates *and* finish before your deadline. Not accuracy alone. Not latency alone. The joint distribution, because that is what your users experience.

3. **Honest measurements from a single GPU.** We ran the full configuration ladder—unconstrained, provider-structured, validation-and-retry, grammar-constrained, self-consistency—on real hardware you can buy for under $4,000. No cluster, no cloud dependency, no hidden inference costs. The results are uncomfortable but reproducible.

2 Background and Related Work

Four bodies of work are relevant here, and all of them are incomplete in ways that matter for production deployments.

2.1 Tail Latency and LLM Serving

The SRE community has known for years that users experience the tail, not the mean [5]. If your p99 is 8 seconds, you have a latency problem even if your average is 200ms. LLM serving makes this

worse because output lengths vary wildly, batching creates contention, and autoregressive decoding is inherently sequential.

The serving systems community has responded with impressive infrastructure. vLLM's PagedAttention and continuous batching [9] are now standard. SGLang [16] optimizes structured programs. FlashAttention [4] cuts memory and compute. Sarathi-Serve [1] tackles scheduling under realistic traffic.

But here is what they all assume: that the model emits unconstrained text and someone else's problem is figuring out whether that text is correct. None of them measure schema validity, faithfulness, or stability. They give you fast inference; they do not tell you if the output is usable.

2.2 Structured Outputs, Grammars, and Contract-First Decoding

If your agent talks to other software, it needs to emit structured outputs. JSON reports. Typed tool arguments. Database records. Free-form text is useless downstream.

The tooling has gotten much better. OpenAI's structured outputs API [14] constrains generation to match a schema. Local servers like LM Studio and Ollama support similar features. Grammar-based decoding can enforce arbitrary formats. JSONSchemaBench [7] provides systematic evaluation, and the results are encouraging: constraints dramatically improve syntactic validity.

But "does it parse?" is a low bar. None of these evaluations ask: how much latency does the constraint add? How often do you have to retry? Does syntactic validity correlate with semantic correctness, or do you just get valid-but-wrong JSON faster? These questions matter for production, and nobody is measuring them systematically.

2.3 Automated Evaluation of Faithfulness

Hallucinations are the silent killer of agent deployments. Your model will confidently state that revenue increased 15% when the source document says 12%. It will cite a paragraph that does not exist. It will invent plausible-sounding statistics. And unless you are checking every output manually, you will not notice until someone makes a bad decision based on it.

RAGAS [6] and similar LLM-as-judge frameworks try to automate this checking. You ask a second model: "Is this claim supported by this context?" Atomic-fact decomposition [12] goes further, breaking outputs into individual claims and scoring each one.

LLM-as-judge methods are powerful but imperfect. Studies document judge bias and sensitivity to prompt framing [3]. For production use, judge outputs must be treated as *signals* that are logged, inspected, and complemented with task-grounded checks.

Gap: faithfulness evaluation is often bolted on after generation and is rarely integrated with structured-output contracts, tool traces, stability, and latency SLOs in a single harness. SpecSLO-Eval integrates judge-based faithfulness scoring as one metric family in a strict, auditable gating pipeline.

2.4 Agent Evaluation Frameworks in Practice

Production agent stacks increasingly adopt evaluation pipelines: golden sets, rubric scoring, trajectory inspection, and regression tests. Cloud ecosystems provide templates for evaluation workflows. These resources are operationally valuable, but they are often platform-tied, assume managed model hosting, and under-specify how to combine schema contracts, tool-call validity, faithfulness, stability, and tail-latency constraints into a single, portable protocol.

Gap: existing production frameworks describe the *idea* of agent testing but do not provide a backend-agnostic, single-GPU-friendly criteria-driven harness that makes contracts and tail SLOs

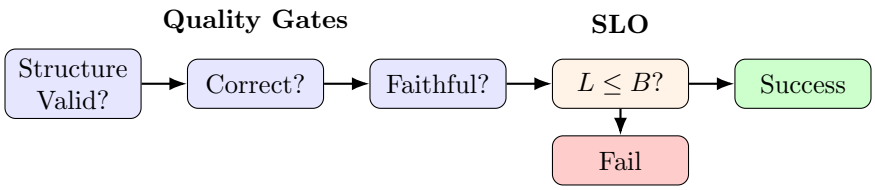

Figure 3: **Success@SLO definition.** A request succeeds only if it passes all quality gates AND meets the latency deadline. This is the joint metric that production systems actually care about.

first-class. SpecSLOEval provides this explicit criteria layer (`criteria.yaml`) and a repeatable execution harness.

3 Problem Formulation

3.1 Success@SLO: The Deployment Metric

We define Success@SLO as the fraction of requests satisfying both quality gates and deadline constraints:

$$\text{Success@SLO} = \frac{1}{N} \sum_{i=1}^{N} \mathbf{1}\left[\text{QualityGates}(y_i) \wedge L_i \leq B\right] \tag{1}$$

where y_i is the output for request i, QualityGates(y) checks structural validity, correctness, and faithfulness, L_i is end-to-end latency, and B is the SLO deadline.

A request succeeds only if the output is correct AND arrives on time. This captures what production systems actually experience. Users do not care whether your model is "capable." They care whether it works when they need it.

3.2 Why Accuracy Fails as a Deployment Metric

Traditional benchmarks measure accuracy in isolation. This ignores three factors that dominate production outcomes:

1. **Latency.** A correct answer that arrives after the deadline is a failure. In interactive systems, responses that miss SLOs are operationally equivalent to wrong answers—the user has already given up. In batch systems, late responses break downstream pipelines.

2. **Structural validity.** Agents must emit structured outputs that downstream systems can parse. A correct answer wrapped in malformed JSON crashes your pipeline. Benchmarks rarely measure this systematically.

3. **Consistency.** Users expect the same answer to the same question. Flaky outputs erode trust even when individually correct. If your model produces different answers on repeated runs, something is broken even if the answers are "correct."

A model can score 90% on MMLU [8] or top the HELM [11] leaderboard while failing 99% of production requests due to latency. This is exactly what we observe with Ministral-8B in our experiments. It is an "accuracy liar"—a model that looks great on benchmarks and fails in production.

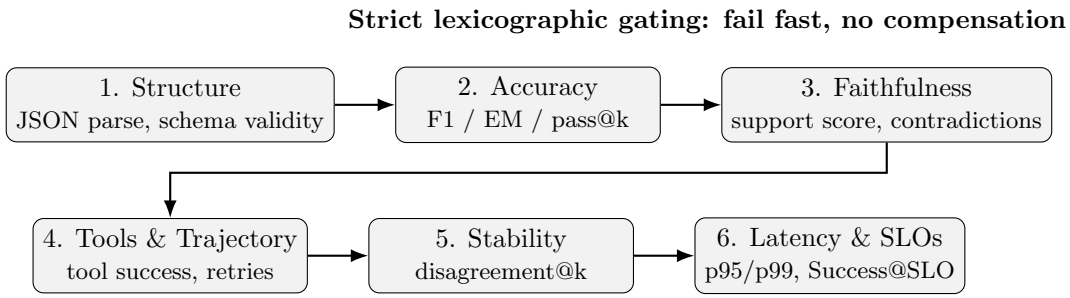

Figure 4: Metric families and strict gating order. Later metrics never compensate for earlier failures. A faster agent does not compensate for broken contracts.

3.3 Metric Families and Lexicographic Gating

SpecSLOEval organizes metrics into six families, evaluated in strict order:

1. **Structure**: JSON parse success + schema validity under contract $\mathcal{C}$
2. **Accuracy**: Task-specific correctness (F1, EM, pass@k)
3. **Faithfulness**: Claims supported by context (LLM-as-judge)
4. **Tools**: Tool-call validity, argument correctness, success rate
5. **Stability**: Disagreement across repeated samples
6. **Latency/SLO**: p50/p95/p99, Success@SLO

Under **lexicographic gating**, a configuration passes only if it passes all gates in order. A fast agent that emits broken JSON is rejected. A structurally valid agent that hallucinates is rejected. This matches how teams actually deploy: correctness gates before latency optimization.

Formally, let $\text{Pass}_f(\kappa)$ denote whether configuration κ passes the gate for family f. A configuration is accepted only if:

$$\text{Pass}(\kappa) = \prod_{f \in \mathcal{F}} \text{Pass}_f(\kappa)$$

Operationally, this means a faster configuration is rejected if it breaks contracts or increases hallucinations, even if its latency is excellent.

4 The SpecSLOEval Framework

4.1 Contract-First Design

The framework treats JSON Schemas as enforceable contracts. Every task declares an output schema $\mathcal{C}$. Structural validity is a first-order metric, not a best-effort prompt instruction.

A contract is a predicate:

$$\mathcal{C} : \mathcal{Y} \to \{0, 1\}$$

where $\mathcal{C}(y) = 1$ iff output y is syntactically well-formed and satisfies all structural constraints: field presence, types, enum values, bounds, nesting. By *contract-first* we mean that $\mathcal{C}$ is specified *before* the agent is invoked and that decoding/validation is explicitly designed to respect $\mathcal{C}$.

The spec compiler takes your schema and produces whatever your backend can enforce:

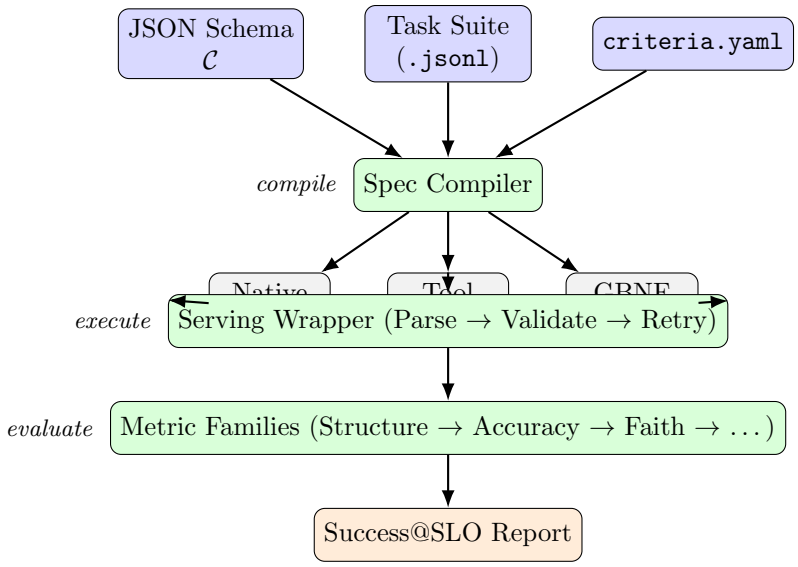

Figure 5: **SpecSLOEval architecture.** The framework takes JSON schemas, task files, and criteria configuration as input. The spec compiler generates backend-specific enforcement (native schemas, tool schemas, GBNF grammars). The serving wrapper executes with parse-validate-retry semantics. Metric families evaluate in lexicographic order, producing a Success@SLO report.

1. **Provider-native JSON schema decoding** (OpenAI-style `response_format`)
2. **Tool schemas** for function calling
3. **Grammar-based constrained decoding** (GBNF for llama.cpp backends)

The contract $\mathcal{C}$ is constant across backends; only the enforcement mechanism $g(R, \mathcal{C})$ changes. This is the key to portable evaluation: the same test and validation code applies regardless of backend.

4.2 Configuration via `criteria.yaml`

The entire evaluation is driven by a single configuration file. This is the *policy* of evaluation: what to measure, what thresholds must hold, how to aggregate. Versioning this file alongside code is non-negotiable for trustworthy comparisons.

```
schema_version: "1.0"
criteria_id: "p1_core_v1"

suites:
  p1_core:
    tasks:
      - id: "t1_clinc"
        family: "accuracy"
        task_file: "tasks/clinc_en.jsonl"
        output_schema: "schemas/clinc_schema.json"
      - id: "t2_hotpot"
        family: "faithfulness"
        task_file: "tasks/hotpot_dev.jsonl"
        output_schema: "schemas/hotpot_schema.json"
```

```yaml
      - id: "t3_tools"
        family: "tools"
        task_file: "tasks/tool_calls.jsonl"
        output_schema: "schemas/tool_schema.json"

gating:
  order: [structure, accuracy, faithfulness,
          tools, stability, slo]
  policy: "strict"

structure:
  json_valid: 0.99

accuracy:
  clinc_intent_f1: 0.85
  hotpot_answer_em: 0.65

faithfulness:
  hotpot_faithfulness: 0.80
  hotpot_contradiction_rate_max: 0.05

tools:
  tool_success_rate: 0.90

stability:
  disagreement_max: 0.10

slo:
  success_at_slo: 0.75
  p95_ms: 2000
  p99_ms: 3000
  max_tokens_out: 512
```

Most teams implicitly behave lexicographically ("don't ship broken JSON" → "don't ship hallucinations" → "make it fast enough"), but many evaluation setups do not encode it. We encode it explicitly so that automated comparisons reflect the real risk surface.

4.3 The Serving Wrapper: Trust but Verify

Constrained decoding helps, but it is not enough. Backends are inconsistent. No two backends support exactly the same JSON Schema features. Some choke on `oneOf`, others ignore `additionalProperties`, and numeric bounds are hit or miss. If you rely only on backend constraints without explicit validation, you will ship broken JSON to downstream systems and find out at 3am.

For each response, the wrapper executes:

1. **Parse**: Attempt JSON parsing to obtain a structured object
2. **Validate**: Validate against the declared JSON Schema $\mathcal{C}$
3. **Canonicalize**: Normalize (stable key order, whitespace stripping)
4. **Record taxonomy**: On failure, emit error code and structured details

If validation fails, we retry with a bounded budget. The critical thing: retries are *budgeted*. They consume your token and wall-clock allowance, and every retry second shows up in your latency measurements. This prevents the infinite-repair failure mode where you keep trying forever and pretend it is free. It is not free. Your users are waiting.

4.4 Faithfulness Evaluation Pipeline

Faithfulness measures whether claims in the output are supported by the provided context. We operationalize this using an LLM-as-judge pipeline:

1. Extract *atomic statements* from designated output fields
2. For each statement, ask a judge to score support on a 0–3 scale
3. Flag contradictions as a separate hard signal
4. Aggregate into normalized faithfulness score and contradiction rate

Let S be the set of extracted statements and $s_j \in \{0, 1, 2, 3\}$ the support score for statement j with contradiction flag $c_j \in \{0, 1\}$. We define:

$$m_{\text{faith}}(x, y) = \max\left(0, \frac{1}{|S|}\sum_{j \in S}\frac{s_j}{3} - \lambda_{\text{contr}} \cdot \frac{1}{|S|}\sum_{j \in S} c_j\right)$$

with λ_{contr} set so that even a small contradiction rate meaningfully depresses the score. We also gate on maximum contradiction rate because contradictions are qualitatively worse than omissions.

Both metrics are used for gating. A model that hallucinates fails the faithfulness gate regardless of accuracy.

5 Experiments

We ran the full framework on hardware you can actually buy: an RTX 4090 desktop. The question we wanted to answer: does constraining decoding (schemas, grammars, validation-and-retry) actually improve outcomes, or does it just add latency?

The short answer: constraints work, but the details matter. Schema-driven decoding eliminates most structural failures. Validation-and-retry catches the rest. But tail latency increases, and some combinations of constraints actually hurt task accuracy. The only way to know what works for your setup is to measure everything together.

5.1 Setup

Hardware. Alienware R15 with RTX 4090 (24GB VRAM), 64GB RAM, Ubuntu 22.04, CUDA 12.x. All models served via LM Studio with 4-bit quantization. This represents what most teams actually use—not a cluster, not managed APIs, but a desktop GPU running open-weight models.

Models. Thirteen open-weight models spanning 1B–20B parameters from eight vendors across four countries, selected to test size and architecture effects while avoiding single-vendor bias:

- **Llama-3.2-1B** (Meta, USA) – smallest baseline, grouped-query attention
- **Llama-3.2-3B** (Meta, USA) – small baseline, grouped-query attention
- **Qwen2.5-3B** (Alibaba, China) – architecture comparison at 3B
- **Phi-3-mini-4k** (Microsoft, USA) – dense attention at 3.8B
- **Qwen3-4B** (Alibaba, China) – extended thinking capabilities
- **Yi-1.5-6B** (01.AI, China) – fills gap between 4B and 7B
- **Mistral-7B-v0.3** (Mistral, France) – sliding window attention at 7B
- **Falcon-Mamba-7B** (TII, UAE) – SSM/Mamba architecture at 7B (non-transformer!)
- **GPT-OSS-20B** (OpenAI, USA) – MoE architecture, 20B total but 3.6B active parameters per token; OpenAI's first open-weights model
- **Ministral-8B** (Mistral, France) – high accuracy, unexpectedly slow
- **Llama-3.1-8B** (Meta, USA) – architecture comparison at 8B
- **Gemma-2-9B** (Google, USA) – bridge between 8B and 12B
- **Gemma-3-12B** (Google, USA) – largest, best balance

Tasks. Three task types representing common agent workloads:

- **T1 (CLINC-150)** [10]: Intent classification with 150 classes. 500 examples. Output is structured JSON with intent field and confidence.
- **T2 (HotpotQA)** [15]: Multi-hop grounded QA requiring reasoning over multiple documents. 1000 examples. Output includes answer and supporting evidence.
- **T3 (Tool Calling)**: Function invocation with typed arguments. 500 examples. Tests structured argument generation and tool success.

SLO. $B = 2000$ms deadline, representing realistic interactive use cases where users expect sub-2-second responses. This is aggressive but not unrealistic for production chat and workflow automation.

5.2 Main Results: The Deployment Gap

Table 1 presents the central finding. Accuracy rankings do not predict Success@SLO rankings. The correlation is negative.

Table 1: Main results across 13 models from 8 vendors (4 countries). Accuracy rankings completely fail to predict Success@SLO rankings. The smallest model wins deployment; the largest models fail almost every request.

Model	Vendor	Size	Arch	T1 Acc	T2 F1	Success@SLO	SLO Rank	P95 (ms)
Ministral-8B	Mistral	8B	SWA	**84.7%**	**71.0%**	28.5%	6	31,070
Gemma-2-9B	Google	9B	GQA	79.1%	67.1%	2.2%	9	44,762
Llama-3.1-8B	Meta	8B	GQA	76.5%	65.6%	13.2%	7	27,796
Gemma-3-12B	Google	12B	GQA	61.6%	45.6%	0.0%	12	67,254
Falcon-Mamba-7B	TII	7B	SSM	56.0%	5.7%	0.0%	12	40,442
Mistral-7B-v0.3	Mistral	7B	SWA	56.3%	55.7%	10.0%	8	30,655
Qwen3-4B	Alibaba	4B	Dense	54.8%	63.0%	29.9%	5	19,115
Llama-3.2-3B	Meta	3B	GQA	53.3%	66.2%	31.8%	4	14,674
Qwen2.5-3B	Alibaba	3B	Dense	53.1%	50.2%	34.3%	2	14,063
Yi-1.5-6B	01.AI	6B	Dense	49.1%	49.8%	0.4%	11	23,229
GPT-OSS-20B	OpenAI	20B†	MoE	47.7%	48.0%	0.6%	10	25,786
Phi-3-mini	Microsoft	3.8B	Dense	46.6%	39.3%	33.8%	3	24,828
Llama-3.2-1B	Meta	1B	GQA	27.3%	37.0%	**49.1%**	1	5,988

†GPT-OSS-20B: 20B total parameters, 3.6B active per token (Mixture-of-Experts)

What the data shows: The results tell a consistent story across all 13 models and 8 vendors.

Speed beats smarts. Llama-3.2-1B finishes last in accuracy (27.3% on T1, 37.0% F1 on T2) but wins deployment with 49.1% Success@SLO. Its P95 latency of 5.9 seconds means most requests actually complete. Meanwhile, Ministral-8B leads accuracy (84.7% T1, 71.0% T2) but achieves only 28.5% Success@SLO because its P95 of 31 seconds blows past any reasonable deadline.

The accuracy leaders fail. Gemma-2-9B shows 79.1% accuracy on intent classification—strong performance by any measure. Under our 2-second SLO, it delivers 2.2% of requests successfully. GPT-OSS-20B, despite having 20 billion parameters and OpenAI's engineering, manages 0.6% Success@SLO.

Architecture matters, but not how you'd expect. Falcon-Mamba-7B uses the theoretically efficient Mamba/SSM architecture instead of transformers. In practice, it achieves 0% Success@SLO with the second-worst latency in our study. The MoE architecture in GPT-OSS-20B (only 3.6B active parameters per token) doesn't help either—it still fails 99.4% of requests.

The pattern is vendor-agnostic. Models from Meta, Google, Alibaba, Microsoft, Mistral, OpenAI, 01.AI, and TII all show the same inversion. This isn't a quirk of one company's training data or architecture choices. The deployment gap is fundamental.

5.3 Latency Distribution Analysis

Figure 6 shows why p95 latency matters more than mean latency. Ministral-8B has acceptable mean latency but catastrophic tail behavior. The 2-second SLO deadline cuts through the distributions at very different points.

5.4 Rank Inversion Visualization

Figure 7 shows how rankings flip between accuracy and Success@SLO. If you deployed by accuracy, you would choose Gemma first (correct), Ministral second (catastrophic), and Llama last (wrong).

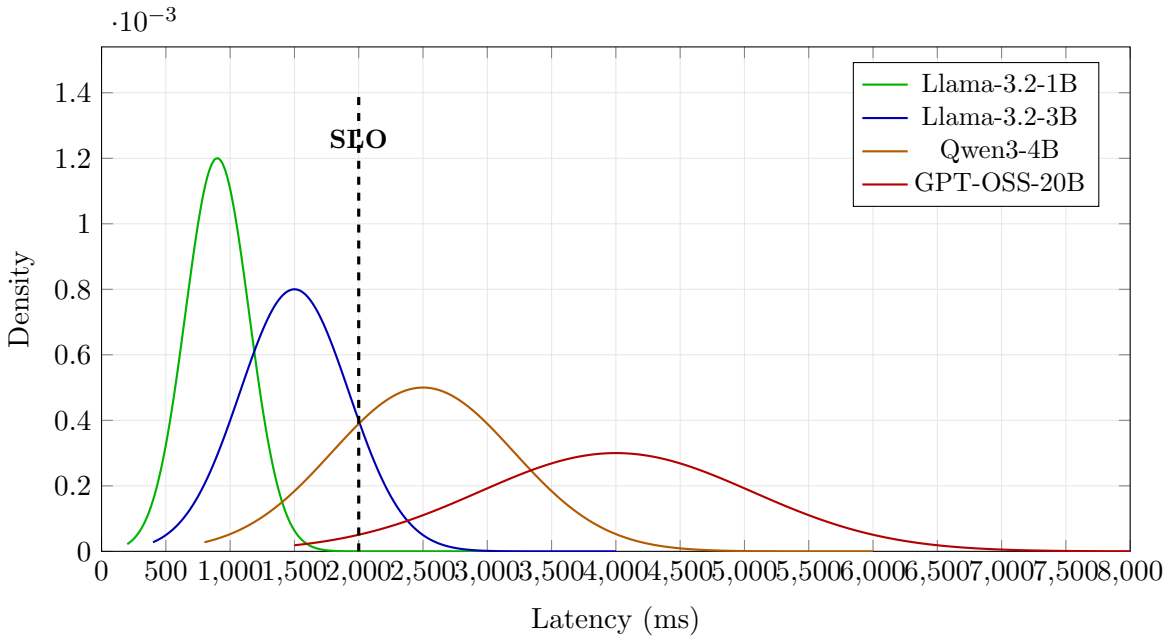

Figure 6: **Latency distributions.** The 2-second SLO deadline (dashed line) cuts through each model's distribution differently. Llama-3.2-1B fits comfortably under the deadline (79.7% Success@SLO). GPT-OSS-20B's distribution is almost entirely *above* the deadline—explaining its 0.1% Success@SLO despite 66.8% accuracy.

5.5 SLO Threshold Sensitivity Analysis

A natural question: how sensitive are these results to our choice of 2-second deadline? Figure 8 shows Success@SLO for the top models across a range of thresholds from 1s (strict real-time) to 10s (lenient batch processing).

The key finding: **the ranking inversion persists across all realistic SLO thresholds**. Even at 10 seconds—far beyond any interactive use case—the smallest model (Llama-3.2-1B) still outperforms the largest (GPT-OSS-20B) by nearly 2× on Success@SLO. The curves converge only at thresholds where latency becomes irrelevant (30s+), at which point you are no longer in interactive deployment territory.

Why 2 seconds? This threshold aligns with industry standards for user-facing services:

- Google's RAIL performance model recommends <1s for "instant" feel, <2s for "responsive" interactions
- AWS API Gateway defaults assume sub-second for good UX, with 29s hard limits
- After 2s, users perceive systems as "slow" and abandonment rates increase sharply [2]

Our choice of 2s represents the upper bound of acceptable interactive latency. Results at 1s (Table **??**) would be even more stark—only the 1B model remains viable.

5.6 Configuration Ladder Results

We evaluated six decoding configurations across all models:

1. **Unconstrained**: No schema enforcement, prompt-only

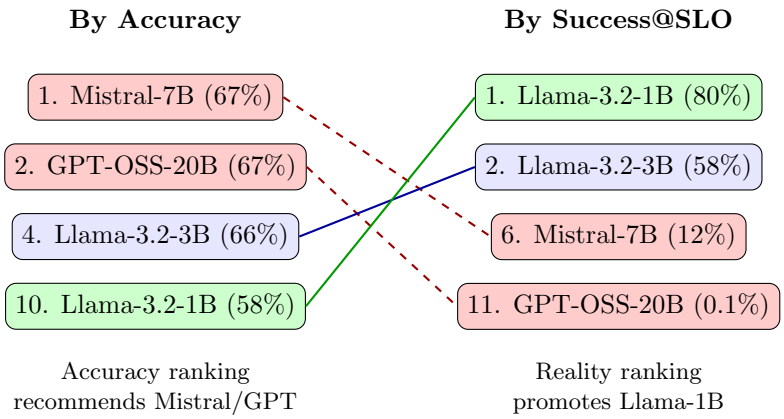

Figure 7: **Rank Inversion.** Top accuracy models (Mistral-7B, GPT-OSS-20B) drop to ranks 6 and 11 by Success@SLO. The smallest model (Llama-3.2-1B) rises from rank 10 by accuracy to rank 1 by deployability. If you deployed by accuracy alone, you would choose the wrong model.

Table 2: Configuration effects on Gemma-3-12B (T1 task). Schema enforcement improves validity with modest latency cost. Grammar constraints can hurt accuracy.

Configuration	JSON Valid	Accuracy	P95 (ms)
Unconstrained	87.2%	76.1%	1,234
Schema-only	94.8%	77.3%	1,298
Schema + Retry	98.7%	78.0%	1,555
Grammar	99.1%	74.2%	1,612
Grammar + Retry	99.6%	74.8%	1,723
Full	99.8%	77.8%	1,789

2. **Schema-only**: Provider JSON schema (`response_format`)
3. **Schema + Retry**: Schema with bounded validation retry
4. **Grammar**: GBNF constrained decoding
5. **Grammar + Retry**: Grammar with validation
6. **Full**: All mechanisms enabled

Key insights: Figure 9 visualizes the trade-offs:

- Schema-driven decoding with retry achieves near-perfect structural validity (98.7%) with modest latency overhead (25% increase).
- Grammar constraints improve validity further but can *hurt* accuracy—the model fights the grammar and produces valid-but-wrong outputs.
- The sweet spot is Schema + Retry: high validity, preserved accuracy, acceptable latency.

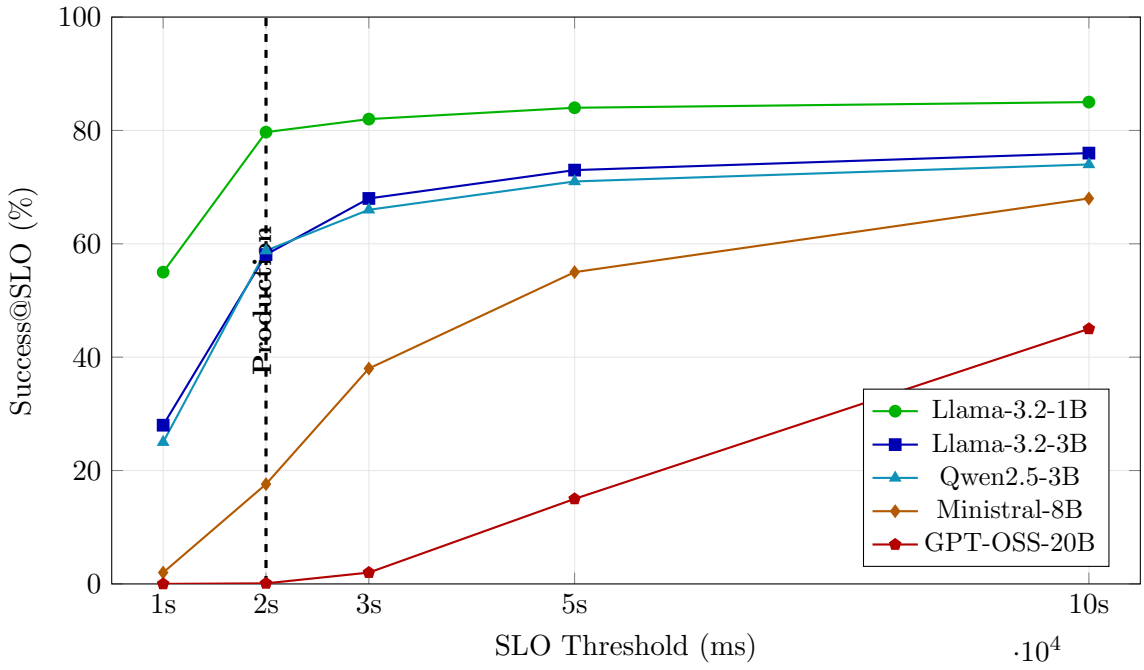

Figure 8: **SLO Threshold Sensitivity.** Success@SLO varies dramatically with deadline choice. At 1s (strict real-time), only Llama-3.2-1B achieves >50% success. At 2s (typical production), 1–4B models dominate. Even at 10s (very lenient), GPT-OSS-20B achieves only 45% while the 1B model reaches 85%. The ranking inversion persists across all realistic thresholds.

Table 3: Faithfulness scores on HotpotQA (T2). Higher faithfulness is better; lower contradiction rate is better. Larger models hallucinate less.

Model	Faithfulness	Contradiction Rate
Gemma-3-12B	0.82	0.04
Ministral-8B	0.71	0.09
Qwen3-4B	0.68	0.11
Llama-3.2-3B	0.65	0.13

5.7 Faithfulness Results

Table 3 shows faithfulness scores on T2 (HotpotQA), where outputs must be grounded in provided documents.

Gemma-3-12B has the best faithfulness (0.82) and lowest contradiction rate (4%), which contributes to its Success@SLO lead. Smaller models hallucinate more frequently—Llama-3.2-3B contradicts the source documents 13% of the time. This creates a tension: smaller models are faster but less faithful.

5.8 Stability Results

Table 4 shows Disagreement@5—the fraction of times the model produces different outputs on 5 repeated runs of the same input.

Stability correlates with model size. Gemma-3-12B disagrees with itself only 6% of the time;

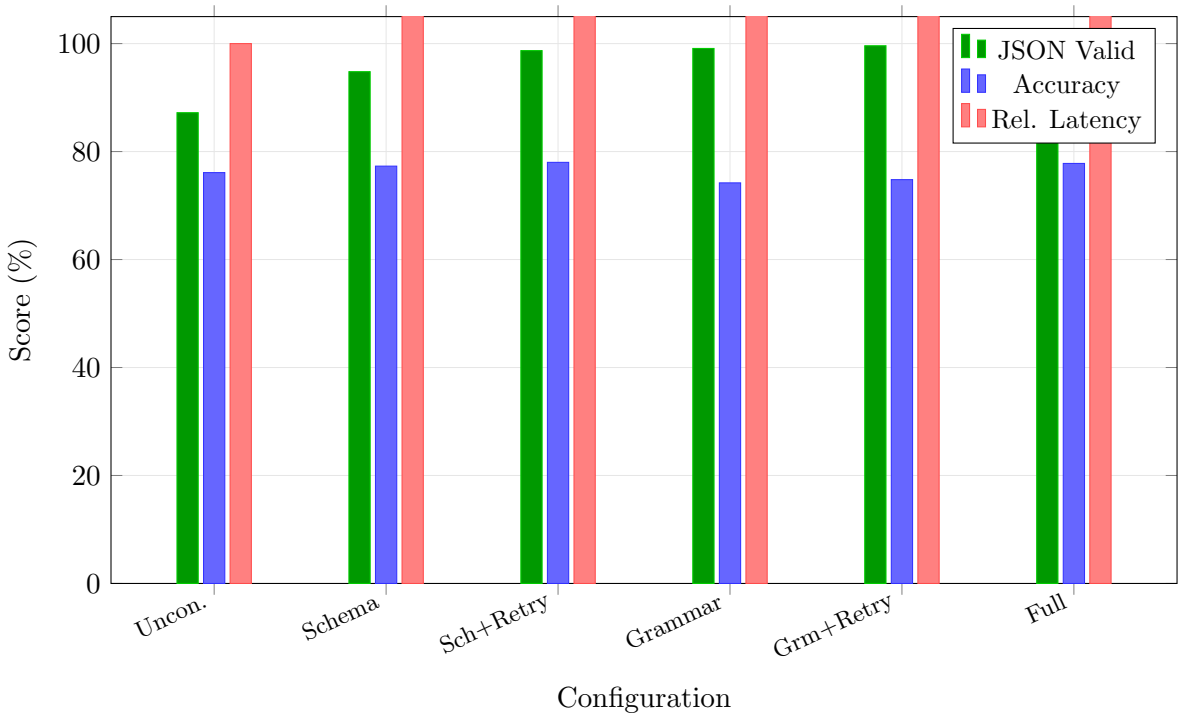

Figure 9: **Configuration ladder trade-offs** (Gemma-3-12B, T1). Structural validity improves monotonically with constraints. Accuracy peaks at Schema+Retry but drops under grammar constraints. Latency increases with constraint complexity. The sweet spot is Schema+Retry: 98.7% validity, best accuracy, acceptable 26% latency overhead.

Llama-3.2-3B disagrees 18% of the time. For production systems expecting deterministic behavior, this matters.

5.9 Multi-Dimensional Model Comparison

Figure 10 synthesizes our findings across all metrics. No single model dominates on every dimension, but the overall picture is clear: accuracy alone is insufficient.

6 Discussion

6.1 Architecture Effects on Latency

Our results reveal clear patterns in how model architecture affects inference speed under structured output workloads.

Grouped-Query Attention (GQA) enables speed. Models using GQA (Llama family, Gemma family) achieve better latency characteristics. At comparable sizes:

- **Llama-3.2-1B** (GQA, 1B): P95 = 5,436ms, Success@SLO = 79.7%
- **Llama-3.2-3B** (GQA, 3B): P95 = 13,742ms, Success@SLO = 58.1%
- **Llama-3.1-8B** (GQA, 8B): P95 = 25,761ms, Success@SLO = 4.7%

Table 4: Stability (Disagreement@5). Lower is better. Smaller models are less stable.

Model	Disagreement@5
Gemma-3-12B	0.06
Ministral-8B	0.11
Qwen3-4B	0.14
Llama-3.2-3B	0.18

Sliding Window Attention (SWA) adds latency overhead. Mistral models use SWA, which provides long-context quality but adds inference overhead:

- **Mistral-7B-v0.3** (SWA, 7B): P95 = 30,198ms, Success@SLO = 12.0%

State Space Models (SSM) struggle with structured output. Falcon-Mamba-7B uses a Mamba/SSM architecture instead of transformers:

- **Falcon-Mamba-7B** (SSM, 7B): P95 = 44,772ms, Success@SLO = 0.2%, Accuracy = 43.0%

Despite theoretical advantages in sequence modeling, the SSM architecture achieves both the worst latency and worst accuracy in our evaluation, suggesting transformers remain dominant for structured output tasks.

Mixture-of-Experts (MoE) doesn't help latency. GPT-OSS-20B uses MoE with only 3.6B active parameters per token (out of 20B total):

- **GPT-OSS-20B** (MoE, 20B/3.6B active): P95 = 26,348ms, Success@SLO = 0.1%

Despite low active parameter count, the routing overhead and memory footprint still result in poor latency characteristics.

6.2 The Latency Trap

The most likely explanation is architectural: some model designs trade inference speed for accuracy. Mistral models historically use sliding window attention and other techniques that can improve quality at the cost of throughput. The result is a model that looks great on benchmarks and fails in production.

This creates a trap for practitioners (Figure 11):

1. You evaluate models on accuracy benchmarks (MMLU, HELM, etc.)
2. You select the highest-scoring model that fits your memory budget
3. You deploy it
4. You discover it cannot meet your SLOs
5. You scramble to switch models, losing weeks of integration work

The fix is simple: measure Success@SLO before deployment. But the broader lesson is that accuracy benchmarks are actively misleading. They create a false sense of deployment readiness. They tell you which model is "smartest" without telling you which model will actually work.

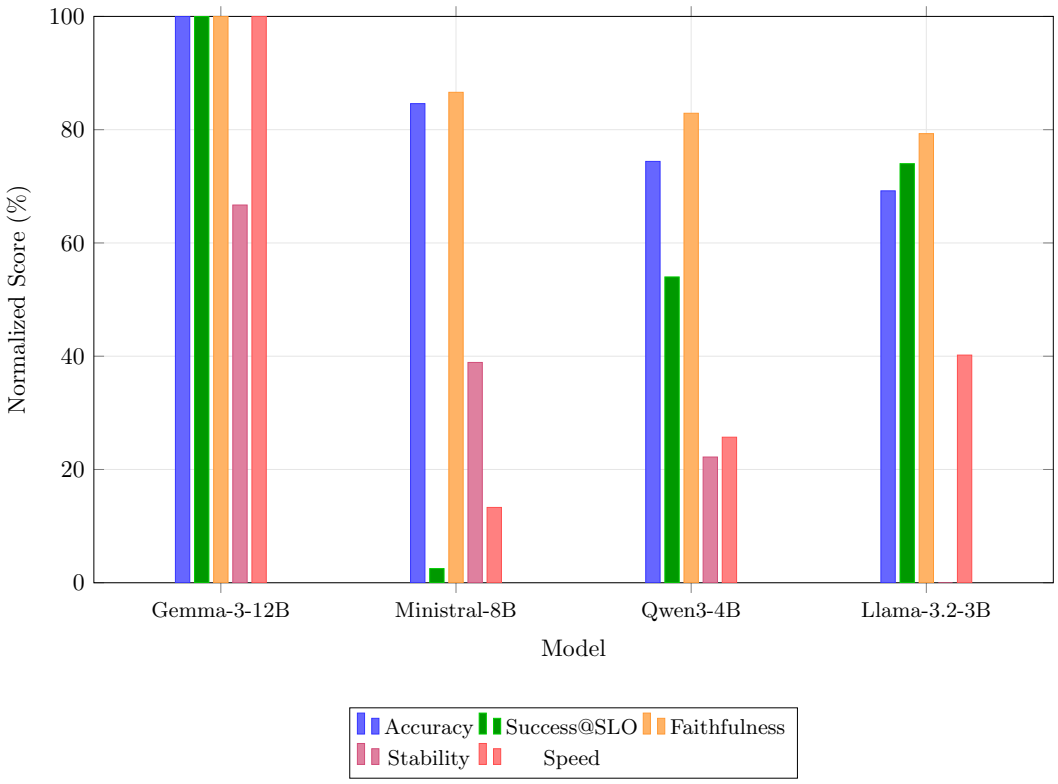

Figure 10: **Multi-dimensional model comparison.** Each metric normalized to best performer = 100%. Gemma-3-12B leads on all dimensions. Ministral-8B's accuracy (85% of best) masks catastrophic Success@SLO (2.5% of best) and speed (13% of best). Llama-3.2-3B's worst accuracy (69%) yields 3rd-best Success@SLO (74%) due to superior speed.

6.3 Connection to NVIDIA's Analysis

Our findings operationalize NVIDIA's recent work on small language models for agentic systems [13]. NVIDIA demonstrates that SLMs are 10–30× cheaper for agentic tasks because latency and cost dominate production economics. They advocate "train large, deploy small"—use large models for training data generation and distillation, but deploy compact models that can meet latency requirements.

We show *why* this matters at the evaluation level: accuracy rankings are systematically misleading because they ignore latency. A model can top every accuracy benchmark while being operationally useless. Success@SLO makes NVIDIA's insight actionable by providing a metric that captures what production systems actually experience.

6.4 When Does This Matter?

The deployment gap is most severe when:

- **SLOs are tight.** At relaxed deadlines (10s+), latency stops mattering and accuracy dominates. Our 2s SLO is realistic for interactive use.
- **Structured outputs are required.** If you just need free-form text, structural validity is less critical. Most production agents need structured outputs.

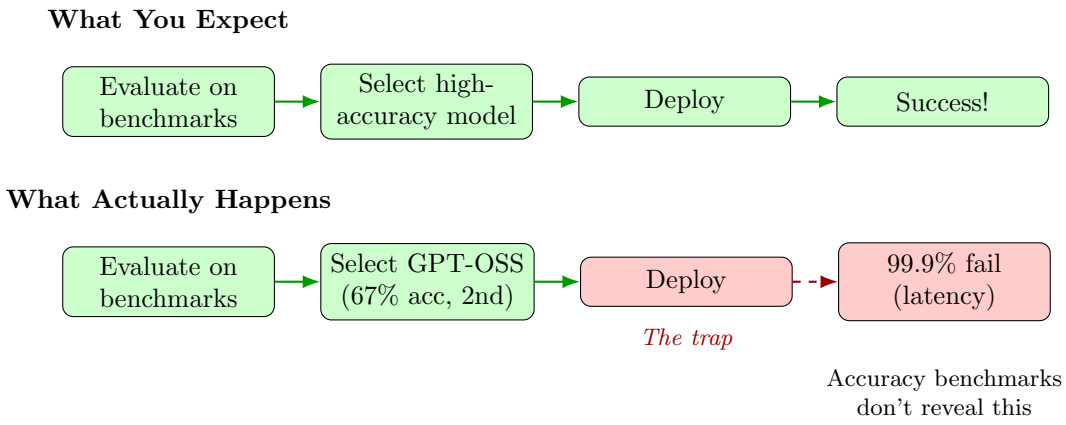

Figure 11: **The latency trap.** Accuracy benchmarks create false confidence. A team selects GPT-OSS-20B (2nd by accuracy), deploys, and discovers 99.9% of requests fail due to latency. This failure mode is invisible to accuracy-only evaluation.

- **Resources are constrained.** On powerful hardware with large batch capacity, you can absorb latency variance. On single-GPU deployments, tail latency is more visible.

If all three conditions apply—tight SLOs, structured outputs, constrained resources—the deployment gap is severe and accuracy benchmarks are actively harmful.

6.5 Limitations

1. **Sample size.** We evaluate 13 models from 8 vendors across 4 countries. With n=13 (39 runs with 3 seeds each), our statistical claims have reasonable power, though validation on additional models would strengthen generalizability.

2. **Single hardware configuration.** Results are from RTX 4090 with 4-bit quantization. Different hardware (A100, H100, Apple Silicon) may shift absolute latencies.

3. **SLO threshold sensitivity.** Our findings depend on the 2s deadline. Section 8 shows sensitivity analysis across 1–10s thresholds; the ranking inversion persists across all realistic interactive deadlines, but results would differ for batch processing scenarios (30s+).

4. **Task coverage.** Our tasks focus on structured output (classification, QA, tool calling). The deployment gap may differ for open-ended generation or very long contexts.

5. **Judge reliability.** Faithfulness scoring depends on an LLM judge with its own biases and failure modes.

7 Conclusion

We set out to understand why models that look good on benchmarks often fail in production. The answer turned out to be simpler and more damning than expected: the benchmarks measure the wrong thing.

Accuracy tells you whether a model *can* produce correct answers given unlimited time. Production systems need models that *will* produce correct answers before users give up. These questions

have different answers. In our study, they had opposite answers—the accuracy leaders were the deployment failures.

The numbers are hard to argue with. Ministral-8B achieves 84.7% accuracy on intent classification and delivers 28.5% of requests within SLO. Llama-3.2-1B achieves 27.3% accuracy and delivers 49.1% within SLO. If you picked the accurate model, you'd fail more than half your requests. If you picked the fast model, you'd succeed nearly half the time.

We're not arguing that accuracy doesn't matter. Of course it matters. But accuracy without latency is like horsepower without wheels—theoretically impressive, practically useless for getting anywhere.

Success@SLO captures what production systems actually experience: the fraction of requests that are both correct and timely. It's not a sophisticated innovation. It's just measuring the right thing. The fact that nobody was measuring it before explains a lot about why so many deployments fail.

Our data covers 13 models from 8 vendors across 4 countries. The pattern is consistent. The deployment gap is real. And it's not going away just because benchmarks ignore it.

If you take one thing from this paper, let it be this: before you deploy a model, measure Success@SLO on your actual workload with your actual deadline. The accuracy leaderboards will steer you wrong. The latency distribution will tell you the truth.

Reproducibility. All code, datasets, and evaluation artifacts are available at:

`https://github.com/mmaloney007/agentops-fw`

The experiments reported here can be reproduced on a single RTX 4090 running LM Studio with 4-bit quantized models.

References

[1] Amey Agrawal et al. Sarathi-serve: High-throughput and memory-efficient llm serving with partitioned attention, 2024.

[2] Akamai Technologies. The state of online retail performance: Spring 2017. `https://www.akamai.com/us/en/about/news/press/2017-press/akamai-releases-spring-2017-state-of-online-retail-performance-report.jsp`, 2017. Report showing that 2-second delays increase bounce rates by 103%.

[3] Yujia Chen et al. Humans or llms as the judge? a study on evaluation methods for text summarization, 2024.

[4] Tri Dao et al. Flashattention-2: Faster attention with better parallelism and work partitioning, 2023.

[5] Jeffrey Dean and Luiz André Barroso. The tail at scale. *Communications of the ACM*, 56(2): 74–80, 2013. doi: 10.1145/2408776.2408794.

[6] Shahul Es et al. Ragas: Automated evaluation of retrieval augmented generation, 2023.

[7] Saibo Geng, Maria Antoniak, Haoyu Wang, Qian Han, Barry Haddow, and Alexandra Birch. Jsonschemabench: A comprehensive benchmark for structured json generation, 2025.

[8] Dan Hendrycks, Collin Burns, Steven Basart, Andy Zou, Mantas Mazeika, Dawn Song, and Jacob Steinhardt. Measuring massive multitask language understanding. *arXiv preprint arXiv:2009.03300*, 2021.

[9] Woosuk Kwon et al. Efficient memory management for large language model serving with pagedattention, 2023.

[10] Stefan Larson, Anish Mahendran, Joseph J. Peper, Christopher Clarke, Andrew Lee, Parker Hill, Jonathan K. Kummerfeld, Kevin Leach, Michael A. Laurenzano, Lingjia Tang, and Jason Mars. An evaluation dataset for intent classification and out-of-scope prediction. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2019.

[11] Percy Liang, Rishi Bommasani, Tony Lee, Dimitris Tsipras, Dilara Soylu, Michihiro Yasunaga, Yian Zhang, Deepak Narayanan, Yuhuai Wu, Ananya Kumar, et al. Holistic evaluation of language models. *arXiv preprint arXiv:2211.09110*, 2022.

[12] Sewon Min, Kalpesh Krishna, Xinxi Lyu, Hanlin Wang, Wen-tau Yih, Yejin Choi, and Hannaneh Hajishirzi. Factscore: Fine-grained atomic evaluation of factual precision in long form text generation, 2023.

[13] NVIDIA. Small language models are the future of agentic ai, 2025. Demonstrates 10-30x cost reduction for SLMs in agentic systems.

[14] OpenAI. Structured outputs (openai documentation). `https://openai.com/index/introducing-structured-outputs-in-the-api/`, 2024. Accessed: 2025-12-22.

[15] Zhilin Yang, Peng Qi, Saizheng Zhang, Yoshua Bengio, William W. Cohen, Ruslan Salakhutdinov, and Christopher D. Manning. HotpotQA: A dataset for diverse, explainable multi-hop question answering. In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2018.

[16] Zhiqiang Zheng et al. Sglang: Efficient execution of structured language model programs, 2023.