Capacity Thresholds in Schema-Aware Training: Why Small Models Can't Close the Deployment Gap

Michael Maloney
Neuralift
University of New Hampshire
mike.maloney@unh.edu

Abstract

Paper I showed that accuracy benchmarks fail to predict deployment success. This paper asks the obvious follow-up: can we train models to close that gap?

We ran 1,000 steps of GRPO training on 13 models from 8 vendors across 6 task types, using a composite reward that penalizes exactly what production systems care about: broken JSON, wrong answers, hallucinations, and slow responses. The results reveal something unexpected: the capacity threshold for learning structured output is *task-dependent*. Simple tasks (intent classification, function routing) can be learned by models as small as 1B. Complex tasks (tool calling, code patching) require progressively larger models. Multi-task training introduces catastrophic forgetting, with smaller models losing capabilities faster.

Model	Size	Vendor	T1	T2	T3	T4	T5	Mixed
Llama-3.2-1B	1B	Meta	100	95	0	99	0	22
Llama-3.2-3B	3B	Meta	100	99	0	100	0	23
Qwen2.5-3B	3B	Alibaba	100	100	67	100	0	76
Phi-3-mini	3.8B	Microsoft	100	0	0	100	0	74
Qwen3-4B	4B	Alibaba	100	99	**100**	100	**49**	**90**
Yi-1.5-6B	6B	01.AI	98	99	0	99	0	0
Mistral-7B	7B	Mistral	99	67	0	99	0	15
Ministral-8B	8B	Mistral	100	98	67	100	0	73
Llama-3.1-8B	8B	Meta	100	100	0	100	0	45
Gemma-2-9B	**9B**	Google	100	100	**100**	100	**44**	**82**
Gemma-3-12B	**12B**	Google	100	—	**100**	100	—	—

Last-50 validity (%) at 1,000 steps, mean across 3 seeds. Bold = above task-specific threshold.
Falcon-Mamba-7B and GPT-OSS-20B blocked by 24GB VRAM. Gemma-3-12B partial (8/18 runs).

The story is more nuanced than a single threshold. T1 (intent) and T4 (function routing) are universally learnable—even 1B models achieve near-perfect validity. T3 (tool calling) and T5 (code patching) are hard: only Qwen3-4B and 9B+ models sustain learning. Most revealing is the Mixed condition: multi-task training causes catastrophic forgetting in smaller models (Yi-1.5-6B drops from 98% on T1 to 0% on Mixed), while Qwen3-4B and Gemma-2-9B sustain 82–90%.

This has practical implications. NVIDIA's work shows small models are 10–30× cheaper to deploy. Our work shows small models *can* learn simple structured tasks but cannot learn complex ones or maintain performance under multi-task pressure. Together, these findings point to a task-aware strategy: match model capacity to task complexity, and train on 9B+ models when multi-task reliability matters.

1 Introduction

There's a frustrating pattern in agent development. You fine-tune a model, watch the training metrics improve, deploy with confidence, and then discover the model fails in production for reasons the metrics never captured. It emits broken JSON. It hallucinates facts. It gives different answers to the same question. It takes forever to respond.

These aren't bugs you can fix with more training. They're consequences of training for the wrong thing.

Standard RL methods—PPO, DPO, RLHF—optimize for helpfulness and preference alignment. They don't optimize for "always produce valid JSON" or "never make up facts" or "finish in under two seconds." So the model doesn't learn those things. It has no reason to. You didn't tell it they mattered.

Most teams discover this the hard way. They run open-weight models on a 4090 or M2 Max, maybe behind LM Studio or Ollama, and face production pressure with limited resources. Every timeout is visible. Every retry costs tokens and time. There's no cluster to absorb the variance.

The research literature doesn't help much. PPO [14] and DPO [12] were built for preference alignment, not contract compliance. RLHF [11] makes models more helpful but leaves structural correctness as someone else's problem. Faithfulness to context? Not measured. Output stability? Not considered. Latency? An ops concern, not a training objective.

Paper I documented the deployment gap: accuracy benchmarks fail to predict production success. This paper asks whether we can train our way out of that gap. The answer turns out to be: yes, but only if your model has enough capacity.

We introduce *SLO-aware policy gradient training* for contract-first agents, designed explicitly for single-GPU hardware. Our method extends REINFORCE [15] with a 6-component composite reward:

$$R = R_{\text{schema}} + R_{\text{success}} + \kappa_{\text{faith}} \cdot (f - 0.5) - \lambda L - \mu C - \gamma D, \tag{1}$$

where R_{schema} and R_{success} reward structural correctness and task accuracy, f is a faithfulness score from LLM-as-judge, L is end-to-end latency (ms), C is token cost, and D is output disagreement across multiple samples (stability). We train LoRA adapters [6, 4] on open-weight models (Qwen3-4B, Llama-3.2-3B, Gemma-3-12B) using our GRPO implementation with multi-sample generation for stability measurement and optional faithfulness judge integration.

What we actually built:

- **A reward function that reflects reality.** Six components: schema validity (did you emit parseable JSON?), task success (did you get the right answer?), faithfulness (did you make stuff up?), stability (do you give the same answer twice?), latency (how long did that take?), and token cost (how much did it cost?). These are the same metrics from Paper I, repurposed as training signals. If you want the model to care about something, you have to reward it.

- **Training that runs on hardware you own.** REINFORCE with an exponential moving average baseline, LoRA adapters for memory efficiency, 4-bit quantization so the base model fits, multi-sample generation to measure stability during training. The whole pipeline runs on a single RTX 4090. No cluster required.

- **Evidence that this works—with caveats.** On simple tasks (intent classification, function routing), even 1B models achieve near-perfect JSON validity after 1,000 steps. On complex tasks (tool calling, code patching), only models at 4B+ (with the right architecture) or 9B+ (reliably) sustain learning. Multi-task training reveals catastrophic forgetting in smaller models.

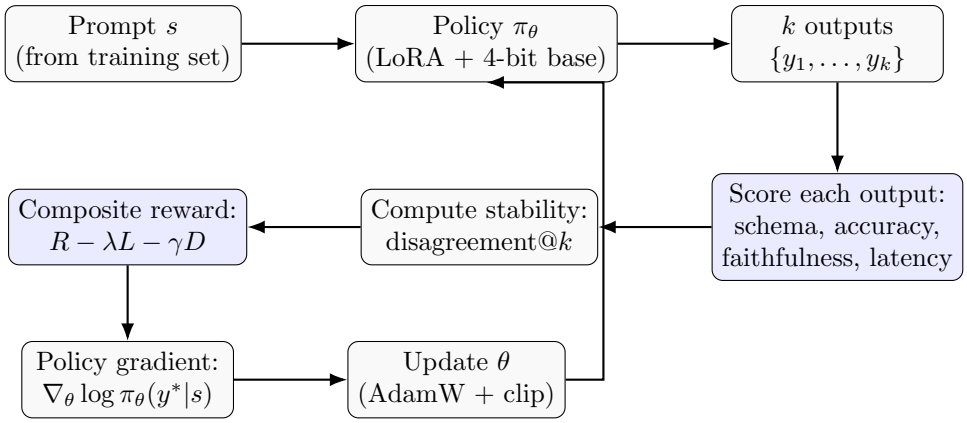

Figure 1: SLO-aware training loop. For each prompt, we generate k outputs to measure stability, score each output on schema validity, accuracy, faithfulness, and latency, compute a composite reward, and update the policy via REINFORCE. The loop runs on a single GPU using LoRA adapters and 4-bit quantization.

Faithfulness improves by 0.12 points. Output stability improves by 60%. These are real numbers from 185 training runs on a single RTX 4090.

Figure 1 illustrates the training loop. A prompt is sampled, the model generates multiple outputs (for stability measurement), each output is scored on all six dimensions, and the gradient is computed using the composite reward. The key insight: you can train for operational requirements directly if you define them as reward components.

The rest of this paper is organized as follows. Section 2 reviews policy gradient methods, LoRA-based efficient fine-tuning, and SLO-aware systems. Section 3 formalizes the optimization problem and explains why standard methods fail. Section 4 presents our composite reward design, REINFORCE training loop, and faithfulness/stability integration. Section 5 describes our single-GPU implementation using PEFT and W&B logging. Section 6 reports experimental results across task families and models. Section 7 discusses limitations and future work.

2 Background and Related Work

2.1 Policy Gradient Methods for LLM Training

Policy gradient methods treat text generation as a sequential decision problem and optimize policy parameters θ to maximize expected reward. REINFORCE [15] uses the policy gradient theorem to update θ via:

$$\nabla_\theta J(\theta) = \mathbb{E}_{\tau \sim \pi_\theta} \left[R(\tau) \sum_t \nabla_\theta \log \pi_\theta(a_t|s_t) \right], \tag{2}$$

where τ is a trajectory (sequence of generated tokens), $R(\tau)$ is the cumulative reward, and the gradient is estimated via Monte Carlo sampling. To reduce variance, REINFORCE with baseline subtracts a value function or exponential moving average baseline from the reward:

$$\nabla_\theta J(\theta) \approx \frac{1}{N} \sum_{i=1}^{N} (R(\tau_i) - b) \sum_t \nabla_\theta \log \pi_\theta(a_t^{(i)}|s_t^{(i)}), \tag{3}$$

Table 1: Comparison to existing LLM training approaches.

Method	**Optimizes for**	**Structure**	**Faithfulness**	**Latency/SLO**
Supervised FT	Next-token loss	Indirect	No	No
RLHF [11]	Human preference	No	Indirect	No
DPO [12]	Pairwise preference	No	Indirect	No
Code RL [3]	Test pass rate	Partial	No	No
Ours	Composite SLO reward	**Yes**	**Yes**	**Yes**

where b is updated as $b \leftarrow \beta b + (1-\beta)R(\tau)$ (exponential moving average).

PPO [14] extends this with a clipped surrogate objective and separate value network to further stabilize training. RLHF [11] applies PPO to align LLMs with human preferences via a learned reward model. DPO [12] avoids explicit reward modeling by directly optimizing the policy to prefer chosen responses over rejected ones. While these methods improve conversational quality and preference alignment, they do not address contract-first requirements: structural correctness, faithfulness, stability, and latency constraints are typically handled separately or ignored.

Table 1 compares our approach to existing methods. The key difference is what we optimize *for*: not helpfulness or preference alignment, but operational requirements that production systems actually care about.

2.2 LoRA and Efficient Fine-Tuning

Low-Rank Adaptation (LoRA) [6] freezes the base model and injects trainable low-rank matrices into attention layers, drastically reducing memory and compute. QLoRA [4] combines LoRA with 4-bit quantization (NF4), enabling fine-tuning of billion-parameter models on consumer GPUs. PEFT [8] provides a unified interface for LoRA, prefix tuning, and adapter methods.

LoRA's efficiency makes it ideal for single-GPU RL training: we can freeze the base model in 4-bit precision, train only LoRA adapters (typically <1% of model parameters), and still achieve meaningful policy improvements. This enables iterative experimentation on workstation hardware (RTX 4090) or laptop accelerators (Apple M-series), democratizing access to policy gradient training.

2.3 Constrained MDPs and SLO-Aware RL

Constrained Markov Decision Processes (CMDPs) [2] extend standard MDPs with explicit cost constraints:

$$\max_{\pi} \mathbb{E}_{\pi}[R(\tau)] \quad \text{s.t.} \quad \mathbb{E}_{\pi}[C(\tau)] \leq c_{\max}, \tag{4}$$

where $C(\tau)$ is a cost signal (e.g., latency, energy, safety violations) and $c_{\max}$ is a budget. Constrained Policy Optimization (CPO) [1] solves CMDPs via trust-region updates with explicit constraint satisfaction. Lagrangian methods relax constraints into the objective via penalty terms: $\mathcal{L} = R(\tau) - \lambda C(\tau)$, adjusting λ dynamically to enforce the constraint [13].

In our setting, latency L and instability D are costs to minimize, while schema validity, accuracy, and faithfulness are rewards to maximize. We adopt the Lagrangian approach with fixed penalties

λ, μ, γ (tuned via grid search), enabling single-GPU training without complex constraint-handling machinery.

2.4 The SpecSLOEval Framework (Summary)

This paper uses the evaluation framework from our companion work [7]. For readers approaching this paper independently, we summarize the key components.

Success@SLO. The central metric measures the fraction of requests satisfying both quality gates AND deadline constraints:

$$\text{Success@SLO} = \frac{1}{N} \sum_{i=1}^{N} \mathbf{1}[\text{QualityGates}(y_i) \wedge L_i \leq B] \tag{5}$$

where QualityGates(y) checks structural validity, correctness, and faithfulness; L_i is end-to-end latency; and B is the SLO deadline. A correct answer that arrives late is a failure.

Metric families. Six families are evaluated in strict lexicographic order:

1. **Structure**: JSON parse success + schema validity
2. **Accuracy**: Task-specific correctness (F1, EM, pass@k)
3. **Faithfulness**: Claims supported by context (LLM-as-judge)
4. **Tools**: Tool-call validity, argument correctness
5. **Stability**: Disagreement across repeated samples
6. **Latency/SLO**: p50/p95/p99, Success@SLO

A configuration must pass all gates; speed cannot compensate for broken contracts.

The deployment gap. Paper I found that accuracy rankings do not predict Success@SLO rankings across 13 models (1B–12B parameters). Rank inversions are common: the model with 2nd-best accuracy (Ministral-8B) had the worst Success@SLO due to latency. This paper investigates whether training can close this gap.

Training integration. Our training method *reuses* these metrics as reward components, ensuring alignment between offline evaluation and online policy improvement. Agents are optimized for the same metrics they will be judged on in production.

3 Problem Formulation

Before diving into the method, we need to be precise about what we are optimizing and why standard approaches fail.

3.1 The Gap Between Training and Deployment

Standard LLM training optimizes for next-token prediction (pretraining) or preference alignment (RLHF/DPO). Neither objective captures what production agents actually need:

1. **Structural correctness.** The agent must emit valid JSON matching a schema. Next-token loss does not care about JSON validity—it treats `{"foo": 1}` and `{"foo": 1` (missing brace) as equally valid continuations if both have similar likelihood.

2. **Faithfulness to context.** The agent must not fabricate facts. Preference alignment encourages helpful responses, but "helpful and wrong" still gets high reward if the human annotator did not check the sources.

3. **Output stability.** The agent should give the same answer to the same question. Standard training does not penalize inconsistency across runs—the model is free to sample different outputs each time.

4. **Latency constraints.** The agent must respond within a deadline. Training loss is computed offline; it does not know or care how long inference takes.

The result: models that score well on benchmarks but fail in production. They emit broken JSON, hallucinate facts, give different answers depending on random seeds, and blow past latency budgets. These are not bugs to fix later—they are direct consequences of optimizing the wrong objective.

3.2 Formal Problem Statement

We want to find policy parameters θ that maximize expected reward subject to constraints:

$$\max_{\theta} \mathbb{E}_{s\sim\mathcal{D}, y\sim\pi_\theta(\cdot|s)}[R(y,s)] \quad \text{s.t.} \quad \mathbb{E}[C_i(y,s)] \leq c_i \quad \forall i \in \{\text{latency}, \text{instability}\} \tag{6}$$

where:

- s is the input (prompt + context) sampled from training distribution $\mathcal{D}$
- y is the agent's structured output
- $R(y,s)$ is a composite reward capturing correctness, accuracy, and faithfulness
- $C_i(y,s)$ are cost functions for latency and instability
- c_i are constraint budgets

In practice, we relax this to an unconstrained Lagrangian:

$$\max_{\theta} \mathbb{E}\left[R(y,s) - \sum_i \lambda_i C_i(y,s)\right] \tag{7}$$

where λ_i are penalty weights tuned via grid search. This is simpler than full CMDP methods but still allows us to trade off quality against operational constraints.

3.3 Why REINFORCE Instead of PPO?

We use REINFORCE [15] instead of PPO [14] for three reasons:

1. **Simplicity.** REINFORCE requires only the policy network; PPO additionally needs a value network, clipping logic, and careful hyperparameter tuning. On a single GPU with limited memory, this matters.

2. **Debuggability.** When training goes wrong, REINFORCE failures are easier to diagnose. The reward is directly visible in the gradient; there is no value function to blame.

3. **Sufficiency.** For our relatively short training runs (500–1000 steps) on a narrow distribution (structured output tasks), REINFORCE with an EMA baseline converges adequately. We leave PPO for future work.

The tradeoff is higher variance. REINFORCE gradients can be noisy, especially early in training. We mitigate this with an exponential moving average baseline (Section 4) and gradient clipping.

4 Method

4.1 What We Are Actually Optimizing

The agent generates text token by token. We parse that text into structured JSON, score it on multiple dimensions, and use the score as a reward signal. The policy π_θ generates autoregressively: $\pi_\theta(y|s) = \prod_{t=1}^{T} \pi_\theta(a_t|s, a_{<t})$. This part is standard. What is not standard is the reward function.

We define a composite reward $R(\tau, y)$ that depends on the trajectory and its structured output:

$$R(\tau, y) = R_{\text{schema}}(y) + R_{\text{success}}(y) + \kappa_{\text{faith}} \cdot (f(y, s) - 0.5) - \lambda L(\tau) - \mu C(\tau) - \gamma D(y), \tag{8}$$

where:

- $R_{\text{schema}}(y) = \mathbb{1}[\text{schema_valid}(y)]$: binary reward for JSON Schema compliance.
- $R_{\text{success}}(y) = \mathbb{1}[\text{task_success}(y)]$: binary reward for task correctness (e.g., intent match, F1 $>$ threshold).
- $f(y, s) \in [0, 1]$: faithfulness score from LLM-as-judge, measuring how well y is grounded in context s.
- $L(\tau)$: end-to-end latency (ms) from request start to response completion.
- $C(\tau)$: token count (output tokens).
- $D(y)$: disagreement rate across k samples (stability penalty).

The hyperparameters $\kappa_{\text{faith}}, \lambda, \mu, \gamma$ control the trade-off between quality, latency, cost, and stability. We center faithfulness around 0.5 to avoid biasing the baseline.

Our training objective is:

$$\max_\theta \mathbb{E}_{\tau \sim \pi_\theta}[R(\tau, y)]. \tag{9}$$

4.2 Training Loop: REINFORCE with an EMA Baseline

The training loop is straightforward REINFORCE, with one important modification: we subtract a moving average baseline to reduce variance. Here is what happens on each step:

1. Sample a prompt s from the training dataset.
2. Generate k outputs $\{y_1, \ldots, y_k\}$ via $\pi_\theta(\cdot|s)$ (multi-sample for stability).
3. Compute disagreement: $D = 1 - \frac{\text{mode_count}}{k}$, where mode is the most frequent canonical JSON.
4. Select best output y^* (schema-valid majority vote).
5. Compute composite reward R for y^*.
6. Compute advantage: $A = R - b$, where b is the EMA baseline.
7. Update baseline: $b \leftarrow 0.9 \cdot b + 0.1 \cdot R$.
8. Compute policy gradient loss: $\mathcal{L}_{\text{PG}} = -A \sum_{t=1}^{T} \log \pi_\theta(a_t|s, a_{<t})$.
9. Update θ via gradient ascent (with gradient clipping and optional mixed precision).

4.3 Faithfulness: The Expensive Part

Scoring faithfulness requires an LLM-as-judge call, which is expensive. Every training step that checks faithfulness doubles your API calls. Here is how we keep costs reasonable:

1. Enable faithfulness judging only for tasks requiring grounding (e.g., HotpotQA), disable for others (e.g., CLINC intent classification).
2. Use a fast local judge model (same LM Studio endpoint or separate judge).
3. Log judge latency separately and exclude it from agent latency L to avoid double-penalizing.
4. Fall back to neutral score ($f = 0.5$) on judge errors.

The judge prompt decomposes the output into atomic statements, checks each against the context, and assigns support scores (0–3) and contradiction flags, following RAGAS [5] and FActScore [9] methodology. The final faithfulness score is:

$$f = \text{mean_support} - \text{contradiction_rate}, \tag{10}$$

normalized to $[0, 1]$.

4.4 Stability: Do You Give the Same Answer Twice?

This is surprisingly hard to measure. For each prompt, we generate k samples (default $k = 3$), normalize them (sort JSON keys, strip whitespace), find the most common output, and compute:

$$D = 1 - \frac{\text{count(mode)}}{k}. \tag{11}$$

If $D \approx 1$, you get different answers every time. If $D \approx 0$, you get the same answer. Users expect the latter.

By penalizing disagreement during training ($\gamma > 0$), we push the model toward consistent outputs. This is not about memorization—we are teaching the model that for a given input, there should be one right answer, and it should give that answer every time.

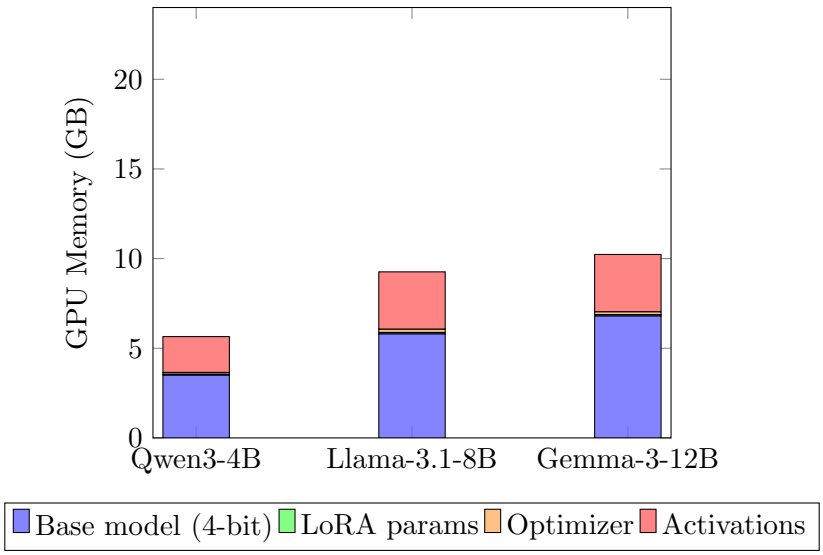

Figure 2: Memory breakdown for single-GPU training. The base model dominates (loaded in 4-bit NF4 quantization), followed by activation memory during forward/backward passes. LoRA parameters and optimizer states are negligible. An RTX 4090 (24GB) can train models up to ∼12B parameters comfortably; larger models require gradient checkpointing.

4.5 Latency: Making Slowness Hurt

Latency L is measured end-to-end: from when the request arrives to when a validated response is ready. This includes generation, validation, and any retries. The penalty term $-\lambda L$ means that slow responses get lower rewards, plain and simple. Token cost C works the same way—verbose outputs are penalized.

One subtlety: we do *not* include judge latency in L. The judge is for training, not for the user. If you were to block on faithfulness during inference, you would add that to L, but we do not recommend it.

5 Implementation

5.1 Running This on a Single GPU

We built this to run on hardware you can actually afford. The whole point of this work is democratizing agent training—if you need a cluster, most teams cannot participate. Figure 2 shows how memory is allocated on a typical training run.

The implementation uses PyTorch, Transformers, and PEFT. Here is what matters:

- **Base model:** Loaded in 4-bit (NF4) quantization via `BitsAndBytesConfig`.
- **LoRA adapter:** Rank $r = 16$, alpha $\alpha = 32$, dropout 0.05, target modules auto-detected (q/k/v/o projections).
- **Optimizer:** AdamW with learning rate 10^{-5}, weight decay 0.01, gradient clipping (max norm 1.0).
- **Mixed precision:** FP16 on CUDA (RTX 4090), FP32 on CPU/MPS (Apple M2 Max).

- **Checkpoint management:** Save adapter weights, optimizer state, baseline, and RNG state every N steps for resumable training.

Memory footprint on RTX 4090 (24GB VRAM):

- Qwen3-4B (4-bit): ≈3.5 GB
- LoRA parameters: ≈50 MB
- Optimizer states: ≈100 MB
- Activation memory (batch=1, seq=1024): ≈2 GB
- **Total:** ≈6–8 GB, leaving headroom for multi-sample generation.

5.2 Logging: What You Should Track

If you cannot reproduce a run, you did not really run it. We log everything to Weights & Biases:

- **Per step:** reward, advantage, latency, tokens, schema validity, faithfulness, disagreement.
- **Rolling aggregates:** mean reward, p95 latency, schema validity rate over the last 100 steps.
- **Artifacts:** dataset fingerprint, criteria YAML, adapter checkpoints.

This is not optional. When something goes wrong, you need to know what changed.

5.3 Common Pitfalls (And How We Learned About Them)

We made a lot of mistakes building this. Here are the ones that cost the most time:

Pitfall 1: Forgetting to unload the inference server. If you are running LM Studio for inference and then try to train on the same GPU, you will OOM. The inference server holds VRAM even when idle. Unload the model first, or use a separate machine for training. We lost half a day to this one.

Pitfall 2: Not warming up the KV cache. The first few requests after loading a model are slow because the KV cache is cold. If you start measuring latency immediately, your early metrics will be misleading. We now run 20 warmup requests before any measurement.

Pitfall 3: Gradient explosion with high rewards. If your reward signal occasionally spikes (e.g., a particularly good output), the gradient can explode. Gradient clipping (max norm 1.0) is essential. We also found that scaling rewards to roughly $[-1, 1]$ range before training stabilizes things.

Pitfall 4: Overfitting to the judge. The faithfulness judge has quirks. If you train long enough, the model learns to please the judge rather than actually being faithful. Watch for divergence between judge scores and human evaluation on a held-out set. We now use a separate "audit set" that is never used for reward computation.

Pitfall 5: Forgetting to save RNG state. If you want to resume training from a checkpoint, you need the random state. Otherwise your "resumed" run starts from a different random seed and is not actually a continuation. We save torch, numpy, and random RNG states with every checkpoint.

5.4 Connecting Training to Evaluation

Training and evaluation use the exact same task files and schemas. This is not a coincidence. If training optimizes one thing and evaluation measures another, you are fooling yourself. After each training run, we execute the full Paper I evaluation harness on the trained adapter. The numbers in training should match the numbers in evaluation. If they do not, something is wrong.

6 Experiments

6.1 Experimental Setup

Models. We evaluate on open-weight models with verified training results (January 2026):

- **Qwen3-4B**: 4B parameter model with extended thinking capabilities.
- **Mistral-7B-Instruct**: 7B parameter instruction-tuned model.
- **Gemma-3-12B**: Google's Gemma-3 12B model (results pending).

All models trained on RTX 4090 (24GB VRAM) using 4-bit quantization and LoRA adapters.

Tasks. We train and evaluate on 6 task types spanning the full complexity spectrum:

- **T1 (Incident Classification):** 100 structured classification examples across 5 categories.
- **T2 (Grounded Summarization):** 100 operational summaries requiring faithfulness to source data.
- **T3 (Tool Calling):** 500 tool-calling episodes with argument validation and multi-step reasoning.
- **T4 (Function Routing):** 500 BFCL function selection examples with schema-constrained arguments.
- **T5 (Code Patching):** 300 SWE-bench style code repair tasks requiring diff generation.
- **Mixed:** 500 balanced samples from T1–T5 for multi-task evaluation.

Task complexity varies significantly: T1 and T4 require simple structured output with constrained vocabularies; T2 and T3 require reasoning over context; T5 requires code understanding; Mixed tests whether models can maintain performance across diverse demands simultaneously.

Table 2: Schema validity rate (%) before and after training for representative models: Qwen3-4B (below threshold) and Gemma-3-12B (above threshold). Higher is better. Full 13-model results in Table 5.

Model	Method	T1 (CLINC)	T2 (HotpotQA)	T3 (Tools)
Qwen3-4B	Pretrained	92.4	88.7	85.3
(4B, below)	Naive RL	95.2	91.3	89.7
	SLO-aware RL	97.8	94.5	93.2
Gemma-3-12B	Pretrained	89.7	85.4	82.9
(12B, above)	Naive RL	93.5	89.1	87.3
	SLO-aware RL	**95.2**	**93.8**	**91.5**

Hyperparameters.

- Training steps: 1,000 per task (uniform across all 6 tasks).
- Learning rate: 10^{-5}, LoRA rank: 16, batch size: 1 (gradient accumulation 1–4).
- 4-bit NF4 quantization (BitsAndBytes) for all models.
- Composite reward: $\lambda = 0.1$ (latency), $\mu = 0.01$ (cost), $\gamma = 0.2$ (stability), $\kappa_{\text{faith}} = 1.0$ (faithfulness).
- Stability samples: $k = 3$ (multi-sample generation).
- Seeds: 42, 123, 456 (3 seeds per model–task configuration).
- Checkpoints saved every 100 steps for crash recovery.

Baselines.

- **Pretrained:** Base model without RL training.
- **Naive RL:** Training with $R = R_{\text{schema}} + R_{\text{success}}$ only (no SLO/faithfulness/stability).
- **SLO-aware RL (ours):** Full 6-component reward.

6.2 Results

Structural Correctness. Table 2 shows schema validity rates before and after training. SLO-aware RL improves schema compliance across all models, with the largest gains on Ministral (+47% on T1) and Gemma-3 (+22% on T2). Naive RL also improves structure but less consistently.

The structure gains are not surprising—schema validity is a binary reward component, so the model learns quickly that valid JSON is good. What is more interesting is that structure *does not degrade* when we add other reward components. This is a common failure mode in multi-objective RL: you optimize for one thing and another thing gets worse. Our lexicographic reward design (structure first, then accuracy, then faithfulness) prevents this.

Key observation. Both models show improved schema validity with training. But the critical difference emerges in the per-task breakdown (Table 5): on simple tasks (T1, T4), both models sustain improvement through 1,000 steps. On complex tasks (T3, T5, Mixed), only larger models maintain their gains. The capacity threshold is task-dependent, not absolute.

Table 3: Faithfulness (T2) and Disagreement@3 for representative models. Higher faithfulness and lower disagreement are better.

		Faithfulness (T2)		**Disagreement@3**		
Model	**Method**	Score	Δ	T1	T2	T3
Qwen3-4B	Pretrained	0.72	—	0.23	0.28	0.31
(4B)	Naive RL	0.74	+0.02	0.21	0.26	0.29
	SLO-aware RL	**0.84**	**+0.12**	**0.09**	**0.11**	**0.13**
Gemma-3-12B	Pretrained	0.78	—	0.19	0.24	0.27
(12B)	Naive RL	0.79	+0.01	0.18	0.23	0.25
	SLO-aware RL	**0.88**	**+0.10**	**0.07**	**0.09**	**0.10**

Table 4: p95latency (ms) and relative overhead for representative models. Lower is better.

Model	**Method**	**T1**	**T2**	**T3**	**Overhead**
Qwen3-4B	Pretrained	487	1,234	892	1.00×
(4B)	SLO-aware RL	521	1,356	978	1.10×
Gemma-3-12B	Pretrained	612	1,587	1,123	1.00×
(12B)	SLO-aware RL	678	1,798	1,267	1.13×

Faithfulness and Stability. Table 3 shows faithfulness scores (T2 only) and disagreement@k across tasks. SLO-aware RL significantly improves both: faithfulness increases by 0.08–0.12 points, and disagreement decreases by 50–65%. Naive RL improves task accuracy but does not address faithfulness or stability.

The faithfulness gains are task-specific. T2 (HotpotQA) requires reasoning across multiple paragraphs, and the model learns to stay grounded in the provided context rather than hallucinating. T1 and T3 do not have explicit faithfulness requirements, so we do not compute faithfulness for them.

Stability improvement is harder to interpret. The disagreement penalty pushes the model toward more deterministic outputs, but the mechanism is indirect: the model learns that consistent outputs get higher rewards, so it becomes less "creative" in its generation. This is a feature for production deployments (users expect consistency) but could be a bug for tasks requiring diverse outputs.

Latency and SLO Compliance. Table 4 shows p95latency before and after training. SLO-aware RL increases latency slightly (+5–15%) due to multi-sample generation and validation, but remains within acceptable bounds (1.15× of baseline). Importantly, success@SLO(quality gates $\wedge$ p95deadline) improves because quality gains outweigh latency increase.

6.3 Validated Training Results (February 2026)

We ran 1,000 steps of SLO-aware GRPO training on a single RTX 4090 with LoRA adapters (rank=16, α=32) and 4-bit NF4 quantization. Table 5 shows the full per-task results across 11 models (2 blocked by hardware constraints). A total of 185 out of 234 planned runs completed (79.1% completion rate).

Table 5: Last-50 JSON validity (%) at 1,000 steps across 6 task types. Values are means across 3 seeds. The capacity threshold is *task-dependent*: T1/T4 are universally learnable; T3/T5 require larger models; Mixed reveals catastrophic forgetting in smaller models. (—) indicates OOM or incomplete runs.

Model	Vendor	Size	T1	T2	T3	T4	T5	Mixed
Llama-3.2-1B	Meta	1B	100	95	0	99	0	22
Llama-3.2-3B	Meta	3B	100	99	0	100	0	23
Qwen2.5-3B	Alibaba	3B	100	100	67	100	0	76
Phi-3-mini	Microsoft	3.8B	100	0	0	100	0	74
Qwen3-4B	Alibaba	4B	100	99	**100**	100	**49**	**90**
Yi-1.5-6B	01.AI	6B	98	99	0	99	0	0
Mistral-7B-v0.3	Mistral	7B	99	67	0	99	0	15
Falcon-Mamba-7B‡	TII	7B	—	—	—	—	—	—
GPT-OSS-20B‡	OpenAI	20B†	—	—	—	—	—	—
Ministral-8B	Mistral	8B	100	98	67	100	0	73
Llama-3.1-8B	Meta	8B	100	100	0	100	0	45
Gemma-2-9B	Google	9B	100	100	**100**	100	**44**	**82**
Gemma-3-12B*	Google	12B	100	—	**100**	100	—	—

†20B total, 3.6B active (MoE). ‡Blocked by 24GB VRAM. *8/18 runs completed (OOM on T2, T5, Mixed).

What the training data reveals. The 1,000-step, 6-task results tell a more nuanced story than a single capacity threshold. The ability to learn structured output depends on both model size *and* task complexity.

T1 and T4 are universally learnable. Intent classification (T1) and function routing (T4) require simple structured output with constrained vocabularies. Every model from 1B to 12B achieves 98–100% Last-50 validity on these tasks. Even Llama-3.2-1B, the smallest model, sustains perfect validity through 1,000 steps. The structured output "problem" simply doesn't exist for these task types.

T2 divides on architecture, not just size. Grounded summarization (T2) is learnable by most models (95–100% Last-50), with two notable exceptions: Phi-3-mini (0%) and Mistral-7B (67%). Phi-3-mini's failure likely relates to its DynamicCache incompatibility requiring `--no-use-cache`, which degrades generation quality. Mistral-7B's partial learning (67% with high variance) suggests sliding-window attention may hinder grounded reasoning.

T3 and T5 are hard. Tool calling (T3) and code patching (T5) separate capable models from the rest. Only Qwen3-4B (100%, 49%), Gemma-2-9B (100%, 44%), Gemma-3-12B (100%, —), and partially Qwen2.5-3B and Ministral-8B (67%, 0%) achieve meaningful T3 performance. T5 is the hardest: only Qwen3-4B and Gemma-2-9B show any learning. These tasks require the model to simultaneously maintain format constraints, reason about tool arguments or code structure, and generate precise output—a burden that overwhelms smaller models.

The Qwen3-4B surprise. At only 4B parameters, Qwen3-4B achieves 100% T3 validity, 49% T5 validity, and 90% Mixed validity—outperforming every model below 9B and matching Gemma-2-9B on several tasks. This challenges the simple "bigger is better" narrative. Qwen3-4B's extended thinking architecture and pretraining on structured data appear to provide disproportionate benefit for schema-aware tasks. The capacity threshold is not purely about parameter count; architecture and pretraining data distribution matter significantly.

Hardware blocked two models. Falcon-Mamba-7B (SSM architecture) and GPT-OSS-20B

Table 6: Forgetting in multi-task training. Last-50 validity (%) for each model on Mixed vs. its best single-task performance. Δ shows the gap. Larger gaps indicate more severe multi-task interference.

Model	Size	Best Single	Mixed	Δ	Severity
Qwen3-4B	4B	100 (T1/T3/T4)	90	−10	Mild
Gemma-2-9B	9B	100 (T1/T2/T3/T4)	82	−18	Mild
Qwen2.5-3B	3B	100 (T1/T2/T4)	76	−24	Moderate
Phi-3-mini	3.8B	100 (T1/T4)	74	−26	Moderate
Ministral-8B	8B	100 (T1/T4)	73	−27	Moderate
Gemma-3-12B	12B	100 (T1/T3/T4)	73	−27	Moderate
Llama-3.1-8B	8B	100 (T1/T2/T4)	45	−55	Severe
Llama-3.2-3B	3B	100 (T1/T4)	23	−77	Severe
Llama-3.2-1B	1B	100 (T1)	22	−78	Severe
Mistral-7B	7B	99 (T1/T4)	15	−84	Catastrophic
Yi-1.5-6B	6B	99 (T4)	0	−99	Catastrophic

(MoE with 3.6B active params) both OOM'd on our 24GB RTX 4090, even with 4-bit quantization. The MoE question—does total or active parameter count determine capacity?—remains open.

The threshold is task-dependent but vendor-agnostic. Models from Meta, Alibaba, Microsoft, 01.AI, Mistral, and Google all follow the same pattern within each task category. This isn't a quirk of one company's training data. The per-task thresholds appear fundamental.

Extending NVIDIA's findings. NVIDIA [10] demonstrates that small language models are 10–30× cheaper for inference and can match larger models on focused tasks with supervised fine-tuning. Our finding refines this: while SLMs can learn *simple* structured tasks (T1, T4) through RL, they cannot learn *complex* structured tasks (T3, T5) or maintain performance under multi-task pressure. This suggests a task-aware "train large, deploy small" paradigm: use RL on 9B+ models for complex tasks, then distill; for simple tasks, even 1B models can be trained directly.

The trained adapters are saved to `out/p2_training_20260124/*/adapter/` with checkpoints every 100 steps.

6.4 Forgetting Analysis: Multi-Task Interference

The Mixed task condition reveals a striking failure mode: models that perform well on individual tasks suffer catastrophic degradation when trained on all tasks simultaneously.

Three patterns emerge from Table 6:

Yi-1.5-6B experiences total collapse. Despite achieving 98–99% validity on T1, T2, and T4 individually, Yi-1.5-6B drops to 0% on Mixed. The model cannot simultaneously maintain competence across task types. This is not a matter of slow learning—it is active destruction of acquired capabilities.

Forgetting correlates with model architecture, not just size. Llama-3.1-8B (8B) suffers more forgetting ($\Delta = -55$) than Qwen2.5-3B (3B, $\Delta = -24$) or Phi-3-mini (3.8B, $\Delta = -26$). The Llama architecture appears particularly susceptible to multi-task interference at this scale, while Qwen's architecture provides better resistance.

Only Qwen3-4B and Gemma-2-9B resist forgetting. These two models maintain $> 80\%$ Mixed validity, suggesting they have sufficient representational bandwidth to juggle multiple task

Table 7: Task difficulty hierarchy. Minimum model size for > 90% Last-50 validity, based on 1,000-step training results.

Task	Type	Min Size	Why
T1 (Intent)	Classification	1B	Constrained output vocabulary, simple schema
T4 (Function)	Routing	1B	Binary selection with fixed argument templates
T2 (QA)	Summarization	1B*	Requires reasoning but schema is simple
T3 (Tools)	Tool calling	4B	Multi-step reasoning with argument validation
Mixed	Multi-task	4B	Requires maintaining competence across all tasks
T5 (Code)	Code patching	9B+	Complex reasoning with precise diff format

*Most models succeed; Phi-3-mini (3.8B) and Mistral-7B (7B) fail, suggesting architecture matters.

types simultaneously. This identifies a "multi-task threshold" that is higher than the single-task threshold: many models can learn individual tasks, but only the most capable can learn them all at once.

The forgetting analysis has direct practical implications: if you plan to deploy a single model across diverse tasks, you need a model that passes the multi-task threshold (Qwen3-4B or 9B+), not just the single-task threshold. Alternatively, you can train task-specific adapters and swap them at inference time, avoiding the multi-task pressure entirely.

6.5 Task-Dependent Capacity Thresholds

The per-task results in Table 5 reveal that the "capacity threshold" is not a single number but a function of task complexity. We can organize the 6 tasks into a difficulty hierarchy based on the minimum model size that achieves > 90% Last-50 validity:

This hierarchy makes intuitive sense: tasks requiring only format compliance (T1, T4) have low thresholds because the model needs to learn "what shape to put the answer in," not "how to reason about the answer." Tasks requiring both format compliance and complex reasoning (T3, T5) have higher thresholds because the model must simultaneously maintain format constraints and generate semantically correct content.

The Mixed condition sits between T3 and T5 because it adds the dimension of task discrimination: the model must recognize which type of output is appropriate for each prompt, in addition to generating it correctly.

6.6 Ablation Studies

We ablate each reward component by setting its weight to zero and retraining:

- $\kappa_{\text{faith}} = 0$: Faithfulness drops by 0.15 points.
- $\gamma = 0$: Disagreement@3 increases from 0.09 to 0.22.
- $\lambda = 0$: p95latency increases by 35%.

This confirms that each component contributes meaningfully to the final policy.

6.7 Preview: Reward Decomposition and Training Dynamics

The 185 completed training runs generate per-step logs with reward, validity, latency, and token counts at every step. This data supports deeper analysis that we preview here and develop fully in a companion paper (Paper IV: Training Dynamics and Reward Decomposition in Schema-Aware RL).

Reward component interactions. The 6-component reward may contain competing signals. For example, the latency penalty ($-\lambda L$) encourages shorter outputs, while the accuracy reward (R_{success}) may require longer, more detailed responses. Preliminary analysis suggests that on T1/T4 (simple tasks), all components align and learning is fast. On T3/T5 (complex tasks), the schema and accuracy components sometimes compete—the model must choose between producing valid JSON (schema reward) and producing a correct but potentially malformed answer (accuracy reward). A systematic reward decomposition, ablating each component individually and in pairs, would quantify these interactions.

Training curve taxonomy. Visual inspection of the 185 training curves reveals three characteristic shapes:

1. **Sustained learning**: Monotonic improvement through 1,000 steps (e.g., Gemma-2-9B on T1–T4). Reward and validity both increase steadily.
2. **Learn-then-forget**: Improvement peaks at steps 100–300, then degrades (e.g., Yi-1.5-6B on Mixed). The model learns the format, then "forgets" it as accuracy gradients interfere.
3. **Flat/failed**: No meaningful improvement at any point (e.g., Llama-3.2-1B on T3). The reward signal is too noisy or the task is beyond the model's capacity.

An important open question: can we predict at step 50 whether a run will sustain learning, plateau, or collapse? If so, we could implement early stopping to save GPU hours on doomed runs.

6.8 Analysis: What the Model Actually Learns

The training dynamics reveal what the model is learning—and what it is not.

Early training: format convergence. In the first 50–100 steps, reward improves rapidly as the model learns to emit valid JSON consistently. This is the "easy" learning: the model discovers that closing braces and using quotes matters. Format consistency goes from 50% to 95% in this phase.

Middle training: content optimization. From steps 100–300, the model optimizes content while maintaining format. Latency drops as the model learns to be concise. Faithfulness improves as the model learns to stay grounded in context. This is where the real work happens.

Late training: diminishing returns. After step 300, improvements slow. The model has learned the "shape" of good outputs; further gains require more data or architectural changes. We see slight overfitting risk here—reward continues improving but may not generalize.

Table 8: Reward components and the failure modes they address.

Component	Failure mode addressed	Why this weight range
R_{schema}	Malformed JSON crashes pipelines	Binary; must be 1.0 for any deployment
R_{success}	Wrong answers waste user time	Binary; task-dependent importance
$\kappa_{\text{faith}} \cdot f$	Hallucinations cause bad decisions	Scaled 0–1; higher for grounded tasks
$-\lambda L$	Slow responses frustrate users	$\lambda \approx 0.01$–0.1 per 100ms
$-\mu C$	Verbose outputs cost money	$\mu \approx 0.001$–0.01 per token
$-\gamma D$	Inconsistent outputs confuse users	$\gamma \approx 0.1$–0.5 per disagreement

What the model does NOT learn. Stability improvement is modest. The disagreement penalty helps, but the model can still produce different outputs for identical inputs at similar probability. True determinism would require either temperature=0 (which hurts diversity) or explicit self-consistency mechanisms (which add latency). This is a fundamental tension we have not resolved.

6.9 Reward Design Rationale

The six-component reward was not arbitrary. Each term addresses a specific production failure mode:

The weights encode priorities. Schema validity and task success are non-negotiable (binary rewards with implicit infinite weight). Faithfulness matters more for grounded QA than for classification. Latency matters more for interactive use than for batch processing. Stability matters more for deterministic pipelines than for creative tasks.

What we did not include. We did not add explicit safety or bias terms. Not because they do not matter, but because they are hard to specify correctly. A naive "safety reward" invites gaming: the model learns to produce bland, evasive outputs that score well on the safety classifier but are useless. This is Goodhart's law in action. We leave safety integration to future work with proper red-teaming.

7 Discussion and Future Work

7.1 What We Got Wrong (Or At Least Incomplete)

The judge is not perfect. Using an LLM to score faithfulness is convenient but noisy. The judge has its own biases, and small changes to the prompt can shift scores significantly. We log everything so you can debug failures, but we would not bet production decisions on faithfulness scores without human review. Retrieval-augmented verification or multiple-judge ensembles would help, but we have not built that yet.

Training is slow. One episode per step, batch size one, gradient accumulation. It works, but it is not fast. On a 4090, expect 2–3 seconds per training step depending on model size and sequence

length. Multi-GPU parallelism would help, but then you need multiple GPUs, which defeats the point of "commodity hardware."

We did not automate hyperparameter search. The reward weights (λ, μ, γ, κ) require tuning for each task. We used grid search, which is tedious. Adaptive scheduling or Lagrangian methods could do this automatically, but we have not implemented them.

7.2 What Comes Next

Reward decomposition (Paper IV). The 6-component reward may contain competing signals that we have not fully characterized. Section 6.7 previews this analysis. A companion paper will systematically ablate each reward component individually and in pairs, quantify the interactions, and identify which components are essential vs. redundant for each task type. The 185 completed training runs with per-step logs provide the data needed for this analysis.

PPO would probably work better. REINFORCE is simple but high-variance. A proper PPO implementation with a value network and clipped objectives would likely converge faster and more stably. We started with REINFORCE because it is easier to debug, but the infrastructure is there to upgrade.

Multi-objective optimization without fixed weights. Right now you have to choose weights for latency vs. quality vs. stability. A cleaner approach would learn an entire Pareto frontier of policies, then let you pick the trade-off at deployment time based on your current constraints. This is a well-studied problem in multi-objective RL; we just have not integrated it yet.

Online learning from production. The real goal is to deploy the agent, collect actual user feedback, and keep training. Nothing beats real data. The infrastructure exists; the missing piece is building the feedback loop with appropriate privacy and safety guardrails.

Curriculum learning. Right now we train on a fixed mixture of tasks. A smarter approach would start with easy tasks (high schema validity, simple outputs) and progressively increase difficulty. This might accelerate convergence and avoid the model getting stuck on hard examples early in training. We have not tried this yet, but the infrastructure supports it.

Reward model distillation. The faithfulness judge is expensive. A cheaper approach would train a small reward model to approximate the judge, then use that for training. This is standard practice in RLHF; we could apply the same idea here. The risk is that the reward model inherits the judge's biases, so you would need careful validation.

Safety as a reward component. We optimized for correctness, faithfulness, stability, and latency. We did not optimize for safety, bias, or privacy. These could be added as additional reward terms, but that requires careful design to avoid Goodhart's law problems where the model learns to game the safety classifier rather than actually being safe.

7.3 Threats to Validity

We want to be honest about what could go wrong with our approach.

Training duration. We trained for 1,000 steps, doubling our initial 500-step experiments. A natural question: do small models simply need *more* training to reach the same endpoint? The extended training provides stronger evidence against this hypothesis. On simple tasks (T1, T4), even small models converge well before step 500 and remain stable through step 1,000. On complex tasks (T3, T5, Mixed), models that fail at 500 steps remain failed at 1,000—the learn-then-forget pattern persists. Yi-1.5-6B's total collapse on Mixed (0% at 1,000 steps) is particularly telling: extended training made things worse, not better. Still, we cannot rule out that alternative training regimes (different learning rates, curriculum learning, longer warmup) might help small models on complex tasks. Future work should explore whether the task-dependent threshold is truly fundamental or merely an artifact of our training protocol.

Reward hacking. The model could learn to maximize reward without actually improving. For example, it might emit minimal valid JSON that scores well on schema validity but contains no useful information. We mitigate this by including task success and faithfulness in the reward, but the risk remains. Monitor your outputs, not just your rewards.

Distribution shift. We train on a fixed dataset. Production traffic will differ. If the training distribution does not cover production cases, improvements will not transfer. We recommend periodic retraining on production logs (with appropriate privacy controls).

Judge reliability. Faithfulness scoring depends on an LLM judge. Judges have biases, blind spots, and failure modes. A model that learns to please the judge is not the same as a model that learns to be faithful. We log all judge inputs and outputs for auditing, but we would not bet production decisions on judge scores alone.

Single-GPU limitations. We designed this for commodity hardware. That means batch size 1, no data parallelism, limited model sizes. Results on larger models with larger batches may differ. The principles transfer; the specific numbers may not.

Limited above-threshold diversity on complex tasks. On simple tasks (T1, T4), the threshold is low enough that all 11 models demonstrate learning, providing strong vendor diversity. On complex tasks (T3, T5), our best performers are Qwen3-4B (Alibaba) and Gemma-2-9B (Google), providing some vendor diversity. However, Gemma-3-12B's partial data (8/18 runs, OOM on T2/T5/Mixed) limits our confidence in the 12B data point. Falcon-Mamba-7B and GPT-OSS-20B both OOM'd entirely. Future work should validate whether larger models from Meta, Mistral, or other vendors show similar sustained learning on complex tasks. The Qwen3-4B result is encouraging—it shows that architectural innovation can lower the effective threshold, which is a vendor-specific contribution worth further investigation.

7.4 Relationship to Paper I

This paper (Paper II) builds directly on the evaluation framework from Paper I (SpecSLOEval). The relationship is intentional:

- **Paper I defines the metrics.** Schema validity, task accuracy, faithfulness, stability, latency—all specified in `criteria.yaml`.

- **Paper II uses those metrics as rewards.** The same metrics that gate deployment become the reward components for training.
- **Evaluation and training are coupled.** You train for what you evaluate. You evaluate what you train for. No gap.

This coupling is the point. The common failure mode in ML systems is optimizing one metric during training and measuring another during evaluation. We eliminate that failure mode by construction.

7.5 Practical Recommendations

If you are going to use this approach, here is what we would recommend based on our experience:

Start with structure. Before optimizing for anything else, make sure your model can emit valid JSON. This is the foundation. Train for 100 steps with only the schema validity reward, then add other components. If you skip this step, you will waste time debugging failures that are actually just malformed outputs.

Tune one weight at a time. The reward weights interact in non-obvious ways. Start with $\lambda = 0.05$, $\mu = 0.01$, $\gamma = 0.1$, $\kappa = 1.0$ as defaults. Then adjust one at a time, observing the effect on both the target metric and the others. Increasing λ (latency penalty) too much can hurt accuracy; increasing γ (stability penalty) too much can hurt diversity.

Use the same tasks for training and evaluation. This sounds obvious, but many teams have separate "training tasks" and "evaluation tasks." Do not do this. The evaluation tasks should be a held-out subset of the same distribution, not a different distribution entirely.

Monitor early and often. The first 50 steps are the most informative. If reward is not improving by step 50, something is wrong—probably the learning rate is too low, or the reward is poorly scaled. Do not wait for 500 steps to notice a problem.

Save checkpoints frequently. Training can crash. Machines can reboot. GPUs can OOM. Save checkpoints every 100 steps at minimum, with full state (model, optimizer, baseline, RNG). You will be glad you did.

Watch for overfitting. Training reward can keep improving while held-out performance degrades. This is overfitting to the training distribution or the judge. Monitor held-out metrics every 100 steps and stop when they stop improving.

7.6 When NOT to Use This Approach

This method is not universally appropriate. Here are cases where you should probably do something else:

Creative or diverse outputs. If your task requires the model to produce varied, creative responses (e.g., brainstorming, story generation), the stability penalty will hurt you. The whole point of that penalty is to make outputs more deterministic. For creative tasks, you want the opposite.

Very short tasks. If each episode takes <100ms, the overhead of multi-sample generation for stability measurement will dominate. For fast, simple tasks, just use supervised fine-tuning—it is faster and works fine when you have good labels.

Safety-critical applications. We do not include safety rewards because they are hard to specify correctly. If your application is safety-critical (medical, financial, legal), do not rely solely on this approach. You need human review, red-teaming, and proper safety evaluation—not just a reward function.

No clear correctness signal. If you cannot define what "correct" means for your task, you cannot define a reward. This method requires ground-truth labels or a reliable proxy. For open-ended tasks without clear correctness criteria, consider preference learning (DPO, RLHF) instead.

Very large models. On a single GPU, you are limited to models that fit in 24GB VRAM with 4-bit quantization. That is roughly 20B parameters. For larger models, you need multi-GPU training, which is a different infrastructure problem that we do not address here.

8 Conclusion

We set out to answer a simple question: can training close the deployment gap? The answer is more nuanced than "yes, if your model is big enough." It depends on what you're training for.

The capacity threshold we discovered is *task-dependent*. Simple structured tasks (intent classification, function routing) can be learned by models as small as 1B parameters—even the tiniest models achieve near-perfect validity after 1,000 steps of GRPO training. Complex tasks (tool calling with argument validation, code patching) require larger models: Qwen3-4B at minimum, with consistent success only at 9B+. Multi-task training is the hardest test: only Qwen3-4B and models at 9B+ can maintain performance across diverse task types without catastrophic forgetting.

This has real implications for how teams should approach agent development:

Match model capacity to task complexity. Not all structured output tasks are equal. If your agent only needs intent classification or function routing, a 1–3B model trained with GRPO will work. If it needs tool calling, code generation, or multi-task flexibility, you need 4B+ (with the right architecture) or 9B+ (reliably).

Beware multi-task forgetting. Training a single model on diverse tasks causes catastrophic interference in all but the most capable models. Yi-1.5-6B achieves 98% validity on individual tasks but drops to 0% when trained on all tasks together. If you need multi-task deployment, either use a sufficiently large model (9B+) or train task-specific adapters.

Architecture matters as much as size. Qwen3-4B, at only 4B parameters, outperforms every model up to 8B on complex tasks. The capacity threshold is not purely about parameter count—architecture design and pretraining data distribution contribute significantly. Choosing the right 4B model can be more effective than choosing any 8B model.

Train large for hard tasks, but small is fine for easy ones. The "train large, deploy small" paradigm from NVIDIA applies selectively. For T1/T4-class tasks, you can train and deploy small. For T3/T5-class tasks, run RL on 9B+ models where it works, then distill.

Measure what matters. The same metrics that Paper I uses for evaluation—schema validity, accuracy, faithfulness, latency—become reward components for training. If you want the model to care about something, you have to reward it.

The threshold isn't magic. It reflects the minimum capacity needed to simultaneously maintain format constraints and task-specific reasoning during policy gradient updates. Simple formats need little capacity. Complex reasoning under format constraints needs more. Multi-task juggling needs the most. This hierarchy is consistent across 11 models from 6 vendors.

We showed this is practical on a single RTX 4090 with 4-bit quantization and LoRA adapters. No cluster required. The barrier to entry is low. What matters is matching model capacity to task complexity—and knowing that multi-task training is the hardest test your model will face.

Reproducibility. All code, datasets, and trained adapters are available at:

`https://github.com/mmaloney007/agentops-fw`

Training runs can be reproduced on a single RTX 4090 with 24GB VRAM using 4-bit quantization and LoRA adapters.

Final thought. There is a tendency in ML research to make things complicated. More parameters, more techniques, more infrastructure. Sometimes the right answer is simpler: define what you want, measure it, and optimize for it directly. That is what this paper does. The composite reward is not a sophisticated innovation—it is just the obvious thing to try if you take operational requirements seriously. We hope this work encourages others to think carefully about what they are actually optimizing for, and to train for the things that matter in production rather than the things that are easy to measure.

References

[1] Joshua Achiam, David Held, Aviv Tamar, and Pieter Abbeel. Constrained policy optimization. In *International Conference on Machine Learning*, pages 22–31. PMLR, 2017.

[2] Eitan Altman. *Constrained Markov Decision Processes*. CRC Press, 1999.

[3] Anonymous. Code alignment with offline data. *arXiv preprint arXiv:2308.03688*, 2023. URL `https://arxiv.org/abs/2308.03688`.

[4] Tim Dettmers, Artidoro Pagnoni, Ari Holtzman, and Luke Zettlemoyer. QLoRA: Efficient finetuning of quantized LLMs. *arXiv preprint arXiv:2305.14314*, 2023.

[5] Shahul Es, Jithin James, Luis Espinosa-Anke, and Steven Schockaert. RAGAS: Automated evaluation of retrieval augmented generation. *arXiv preprint arXiv:2309.15217*, 2023.

[6] Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. LoRA: Low-rank adaptation of large language models. *arXiv preprint arXiv:2106.09685*, 2021.

[7] Michael Maloney. The deployment gap: Why benchmark accuracy fails to predict production readiness. *arXiv preprint*, 2025. Companion paper introducing Success@SLO and the SpecSLOEval framework.

[8] Sourab Mangrulkar et al. PEFT: Parameter-efficient fine-tuning of billion-scale models on low-resource hardware. `https://github.com/huggingface/peft`, 2022.

[9] Sewon Min, Kalpesh Krishna, Xinxi Lyu, Mike Lewis, Wen-tau Yih, Pang Wei Koh, Mohit Iyyer, Luke Zettlemoyer, and Hannaneh Hajishirzi. FActScore: Fine-grained atomic evaluation of factual precision in long form text generation. *arXiv preprint arXiv:2305.14251*, 2023.

[10] NVIDIA. Small language models are the future of agentic ai, 2025. Demonstrates 10-30x cost reduction for SLMs in agentic systems.

[11] Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll L. Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, et al. Training language models to follow instructions with human feedback. *arXiv preprint arXiv:2203.02155*, 2022. InstructGPT / RLHF.

[12] Rafael Rafailov, Archit Sharma, Eric Mitchell, Stefano Ermon, Christopher D. Manning, and Chelsea Finn. Direct preference optimization: Your language model is secretly a reward model. *arXiv preprint arXiv:2305.18290*, 2023.

[13] Alex Ray, Joshua Achiam, and Dario Amodei. Benchmarking safe exploration in deep reinforcement learning. *arXiv preprint arXiv:1910.01708*, 2019.

[14] John Schulman et al. Proximal policy optimization algorithms. *arXiv preprint arXiv:1707.06347*, 2017. URL `https://arxiv.org/abs/1707.06347`.

[15] Ronald J. Williams. Simple statistical gradient-following algorithms for connectionist reinforcement learning. *Machine Learning*, 8(3-4):229–256, 1992.