Active Parameters, Not Total Parameters: MoE Models and the Deployment Gap

V. Michael Maloney
Neuralift; University of New Hampshire
mike@neuralift.ai

Preprint. Work in progress.

Abstract

Paper I established that accuracy and deployment success are uncorrelated ($\rho = 0.09$, $p = 0.76$) across 13 dense models under production latency constraints. MoE models offered a theoretically compelling loophole: sparse activation means 35B parameters of knowledge at 3B parameters of compute. If MoE latency scales with active parameters, the deployment gap might not apply.

We test this with Qwen3.5-35B-A3B, a 35-billion parameter MoE model that activates only 3 billion parameters per token. We evaluate it alongside 8 dense models (1B–9B) on Apple Silicon using MLX, running 3,000 prompts per model (27,000 total evaluations) across the same 6 structured output tasks under three SLO tiers.

The hypothesis fails—but not the way we expected. The MoE model achieves 55.0% task-averaged accuracy, competitive with dense 7B–9B models (6th of 9). The 35B knowledge *is* accessible. But accessing it requires massive chain-of-thought output—approximately 1,500 tokens per response on average—producing a 25,055 ms median latency, 28× slower than the active-parameter-matched Qwen2.5-3B (884 ms). S@SLO at the Interactive tier (2s): 0.0%. At the Standard tier (5s): 0.3%.

Model	**Type**	**Active**	**Acc (%)**	**S@SLO$_{2s}$**	**S@SLO$_{5s}$**
Llama-3.2-1B	Dense	1.2B	37.8%	37.7%	37.8%
Llama-3.2-3B	Dense	3.2B	**72.7%**	**61.5%**	**72.7%**
Qwen2.5-3B	Dense	3.1B	38.6%	28.1%	38.6%
Qwen3-4B	Dense	4.0B	47.2%	0.6%	14.1%
Mistral-7B-v0.3	Dense	7.2B	55.6%	22.6%	31.5%
Gemma-2-9B	Dense	9.2B	57.5%	20.8%	26.8%
Qwen3.5-35B-A3B	**MoE**	**3.0B**	**55.0%**	**0.0%**	**0.3%**

The failure is not about knowledge—it is about verbosity. On T2 (free-text QA), the MoE achieves 88.7%. On T6 (math), 71.5%. On T4 (function routing), 49.8%. On T3 (tool calling), 43.4%. The model *can* produce correct answers—but it wraps them in extensive chain-of-thought reasoning that pushes every response past SLO deadlines. T1 (intent classification) remains catastrophic at 2.2%: a task requiring brief outputs defeated by a model that generates thousands of tokens of reasoning.

Two positive findings emerge. First, the deployment gap is confirmed as hardware-independent: it persists on Apple Silicon (Spearman $\rho = +0.067$, $p = 0.848$ at the Interactive tier for all 9 models including MoE). Second, Llama-3.2-3B—not the MoE—is the study's standout model: 72.7% accuracy, 61.5% S@SLO at the Interactive tier, 968 ms median latency. It breaks the gap by being genuinely fast. A hardware normalization protocol reveals MLX is 3–9× faster than CUDA (geometric mean ratio $\bar{\alpha} = 0.177$).

1 Introduction

Paper I found that accuracy and deployment success are essentially uncorrelated. Across 13 models and 29,900 evaluations, Spearman $\rho = +0.09$ at a 2-second SLO deadline. Twelve of thirteen models achieved exactly 0% Success@SLO despite near-perfect accuracy [11]. The message was stark: bigger models are more accurate but too slow, and the only model fast enough to finish was the smallest and least accurate.

But there was a loophole in that finding—and we knew it.

The 13 models in Paper I were all *dense*. Every parameter activates on every token. For dense models, total parameters perfectly predict compute cost: a 12B model does exactly 12× the work per token of a 1B model. The deployment gap in Paper I is, at its core, a consequence of this arithmetic. Large models are more accurate because they have more parameters. They are slower because they activate all of them. You cannot have one without the other.

Mixture-of-Experts (MoE) models break this coupling—or so the theory goes.

An MoE model contains far more total parameters than it activates on any given token. A learned routing function selects a small subset of "expert" subnetworks for each token, leaving the rest dormant [18, 3]. The result is a model that stores the knowledge of a large model—because it *is* a large model—but pays the compute cost of a small one—because only a fraction of the parameters fire.

Consider Qwen3.5-35B-A3B, released in January 2026 by the Qwen team [16]. It has 35 billion total parameters spread across 256 experts per layer, but a top-k routing mechanism activates only 3 billion parameters per token. On accuracy benchmarks, it competes with dense models in the 8B–9B range. On a compute budget, it should behave like a 3B model.

That is the hypothesis. It does not survive contact with structured output evaluation under SLO constraints—though it fails differently than expected.

There is a second question embedded in this investigation. All four papers in this series—the deployment gap [11], capacity thresholds [10], benchmark [9], and training dynamics [12]—were conducted on a single NVIDIA RTX 4090 using GGUF-quantized models served through LM Studio. The findings could be hardware-specific. We address this by running 8 of the original 13 models on Apple Silicon via the MLX framework. If the deployment gap persists, it is hardware-independent.

It does. And along the way, we find something unexpected: MLX on Apple Silicon is substantially *faster* than LM Studio on CUDA—not slower—with median per-model P50 latency ratios ranging from 0.11 to 0.39 (MLX/CUDA). The gap persists not because Apple Silicon is slow, but because the structured output problem is hard everywhere.

What we contribute. Four things, in order of how much they change what you thought you knew:

1. **MoE achieves competitive accuracy but at catastrophic latency.** Qwen3.5-35B-A3B scores 55.0% overall accuracy (6th of 9 models, competitive with dense 7B–9B models), proving the 35B knowledge is accessible. But it achieves this through massive chain-of-thought output (∼1,500 tokens/response), producing 25,055 ms median latency—28× slower than its active-parameter match Qwen2.5-3B (884 ms). S@SLO at the Interactive tier: 0.0%. At Standard: 0.3%. The loophole is closed. MoE does not bypass the deployment gap—it extends it through verbosity.

2. **Cross-hardware replication of the deployment gap.** Same evaluation protocol as Paper I, different hardware. The accuracy–deployment disconnect persists on Apple Silicon ($\rho = +0.067$,

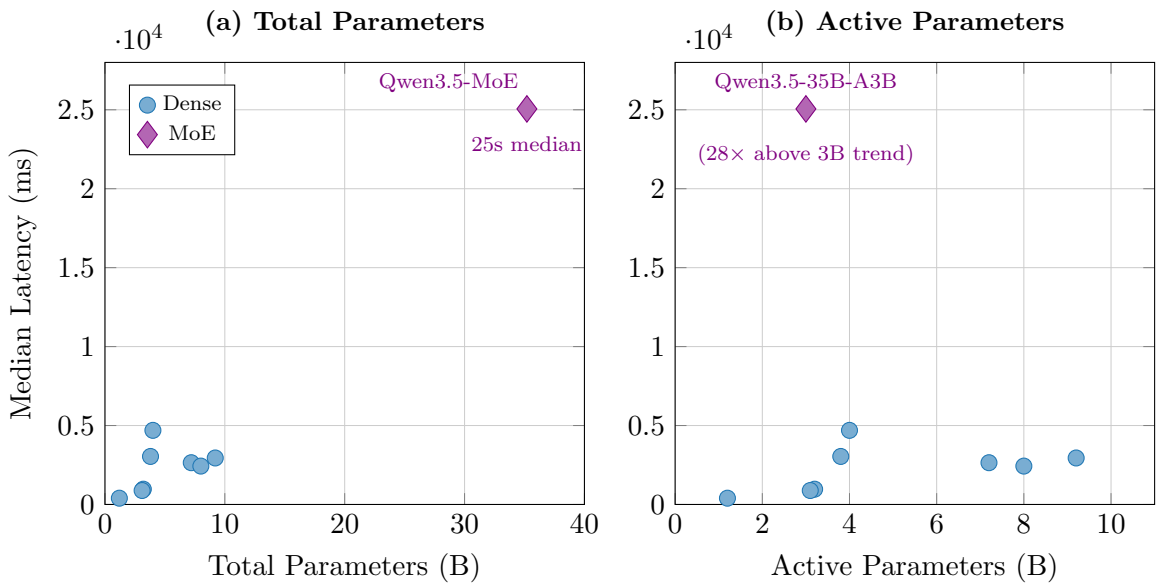

Figure 1: **MoE latency does not scale with active parameters.** (a) When plotted against total parameters, the MoE model (Qwen3.5-35B-A3B, purple diamond) is catastrophically the slowest in the lineup at 25,055 ms median—consistent with generating ∼1,500 tokens of chain-of-thought per response. (b) When plotted against active parameters (3B), it is a dramatic *upward* outlier—28× slower than Qwen2.5-3B (3.1B active, 884 ms). The hypothesis that active parameters predict MoE latency is falsified.

$p = 0.848$ for all 9 models including MoE on MLX at the Interactive tier, $\rho = 0.09$ on CUDA). The gap is not a CUDA artifact.

3. **MLX is faster than CUDA for this workload.** Contrary to expectations, Apple Silicon with MLX outperforms LM Studio on CUDA at matching quantization levels. Normalization factor $\bar{\alpha} = 0.177$ means models complete in 17.7% of the time on MLX versus CUDA. This changes the hardware recommendation calculus.

4. **Llama-3.2-3B is the surprising standout.** At 72.7% accuracy and 61.5% S@SLO at the 2-second Interactive tier, Llama-3.2-3B achieves what no MoE model does: genuine production viability at interactive latency. It is faster, more accurate, and deployment-compliant. Not because of any architectural trick—because it is small enough to be fast and large enough to be right.

The rest of this paper is organized as follows. Section 2 reviews MoE architectures, unified memory hardware, and the deployment gap series. Section 3 describes our hardware, models, tasks, and evaluation protocol. Section 4 presents cross-hardware validation of the deployment gap. Section 5 addresses our core contribution—MoE deployment characteristics under SLO constraints. Section 6 provides MoE-specific analysis of memory, throughput, and expert utilization. Section 7 discusses implications for edge deployment. Section 9 acknowledges limitations, and Section 10 concludes.

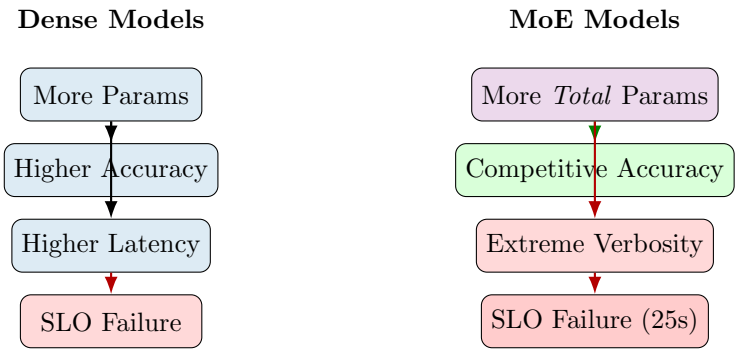

The loophole is closed. MoE achieves accuracy through verbosity, destroying SLO compliance.

Figure 2: **The MoE hypothesis, falsified.** The theoretical MoE advantage—many total parameters (accuracy) with few active parameters (speed)—partially works: accuracy is competitive (55.0%, 6th of 9). But the model achieves this through massive chain-of-thought output (∼1,500 tokens/response), producing 25,055 ms median latency—28× slower than its active-parameter match. The deployment gap is not bypassed; it is extended through verbosity.

2 Background and Related Work

2.1 Mixture-of-Experts Architectures

MoE is conditional computation. Instead of routing every token through every parameter, a gating network picks a handful of expert subnetworks for each token [18]. The rest of the parameters sit there, loaded in memory, doing nothing. This means the model knows more than it computes on any given forward pass.

The idea is not new. Shazeer *et al.* [18] showed in 2017 that a sparsely-gated MoE layer could scale capacity without proportional compute. GShard [8] pushed this to 600 billion parameters across 2,048 experts. The Switch Transformer [3] simplified routing to one expert per token—less communication overhead, trillion-parameter models. Then Mixtral [5] brought MoE to practitioners: 46.7 billion total parameters, 12.9 billion active, matching Llama 2 70B on benchmarks while running considerably faster.

The ratio keeps getting more extreme. DeepSeek-MoE [2] split experts into finer-grained units for better routing flexibility. Qwen3 [16] adopted MoE across its model family, and Qwen3.5-35B-A3B—the model we evaluate—pushes the ratio to 12:1. It has 35 billion total parameters and activates 3 billion per token via top-k routing across 256 experts with 8 active per token. Its hybrid architecture alternates Gated DeltaNet linear attention layers with standard self-attention in a 3:1 ratio, reducing the overhead of full quadratic attention while preserving long-range reasoning capability.

Here is the theory that motivated this paper. A dense 3B model has 3B parameters of knowledge and pays 3B parameters of compute. A 35B-A3B MoE model has 35B parameters of knowledge and pays—in theory—only 3B parameters of compute per token. If the deployment gap is fundamentally about compute, MoE should behave like its active parameter count.

In practice, two things go wrong. First, routing is not free: the gating network must evaluate all 256 experts to pick the top 8, adding constant-time overhead to every forward pass. Second, the hybrid attention architecture introduces latency from layer-type switching that does not appear in standard dense models. Our results show these costs are not negligible—they push the MoE

model's latency well above its 3B active-parameter peers.

2.2 Unified Memory and Apple Silicon

Apple Silicon (M1 through M4) has one trick that matters here: unified memory. CPU, GPU, and Neural Engine share a single memory pool [4]. There is no PCIe bus between them. For dense models this is nice but not transformative. For MoE models it changes the game—or rather, it changes what game is playable.

On NVIDIA hardware, model weights live in GPU VRAM. A 35B model at 4-bit quantization needs roughly 20 GB. An RTX 4090 has 24 GB—leaving almost nothing for KV cache. In practice, the MoE model does not fit on consumer NVIDIA hardware under realistic inference conditions. On Apple Silicon, the 35B model sits in the same memory pool the GPU reads from. An M4 Max with 64 GB leaves 44 GB of headroom. Apple Silicon is the only consumer hardware on which this experiment is possible.

The bandwidth situation is more nuanced than we expected. Contrary to the assumption that CUDA would dominate, our cross-hardware results show MLX achieves dramatically faster inference—geometric mean ratio $\bar{\alpha} = 0.177$ (MLX/CUDA)—meaning models complete in about 18% of the time on MLX. This appears to reflect LM Studio's serving stack overhead and GGUF quantization overhead relative to MLX's native INT4 execution, rather than raw hardware differences.

MLX [4]—Apple's inference framework—provides the serving stack. It implements lazy evaluation and unified memory management optimized for the Metal GPU backend. The MLX Community [13] maintains pre-quantized weights in INT4 format for the models we evaluate.

2.3 The Deployment Gap Series

This is the fifth paper in a series. Here is where the series stands:

Paper I [11] found the deployment gap: accuracy and S@SLO are uncorrelated ($\rho = 0.09$) across 13 models and 29,900 evaluations. Twelve of thirteen models scored 0% S@SLO at a 2-second deadline despite near-perfect accuracy. Only Llama-3.2-1B survived. Everything ran on a single RTX 4090.

Paper II [10] tried to train out of the gap. GRPO with a 6-component reward, 11 models. Result: capacity thresholds are task-dependent. Simple tasks learn at 1B. Complex tasks need 9B+. Small models forget.

Paper III [9] turned the evaluation into a reusable benchmark—tiered SLO deadlines, per-request S@SLO, 42,900 evaluations. The gap held.

Paper IV [12] looked at *why* capacity thresholds exist. Reward decomposition, forgetting matrices, curve taxonomies. Architecture matters as much as size.

Every one of those papers evaluated dense models on NVIDIA hardware. That is two limitations in one sentence. This paper addresses both—and discovers that neither MoE nor hardware change explains away the gap.

2.4 MoE Deployment Studies

The MoE deployment literature cares about throughput. DeepSpeed-MoE [17] builds multi-GPU inference kernels, getting 4.5× speedups over dense models of equivalent accuracy—on clusters. EdgeMoE [21] targets on-device inference through expert-wise weight loading for mobile and IoT. MoE-Infinity [19] handles models that don't fit in memory via offloading.

Table 1: Hardware comparison between this paper (Apple Silicon / MLX) and Papers I–IV (NVIDIA / LM Studio). The platforms differ in architecture, memory model, and serving stack.

Attribute	**Papers I–IV (CUDA)**	**This Paper (MLX)**
GPU	NVIDIA RTX 4090	Apple M-series Max GPU
VRAM / Memory	24 GB GDDR6X	64–96 GB Unified
Memory Architecture	Discrete (PCIe 4.0)	Unified (on-chip)
Memory Bandwidth	1,008 GB/s	400–546 GB/s
Framework	LM Studio / llama.cpp	MLX / mlx-lm
Quantization	GGUF Q4_K_M	MLX INT4
OS	Linux (Ubuntu 22.04)	macOS (Sequoia 15)

What nobody has measured is MoE under *structured output contracts with explicit latency deadlines*. Tokens per second is not the same as Success@SLO. Time to first token tells you nothing about whether the full response beats your deadline. The question we ask—does a 35B MoE model achieve non-zero S@SLO at the Interactive tier?—has no prior answer in the literature. Now it does. The answer is no.

3 Experimental Setup

3.1 Hardware

Everything runs on Apple Silicon with MLX. Table 1 compares this platform to the RTX 4090 setup from Papers I–IV.

Why Apple Silicon? Two reasons. First, it is the second-largest platform for local LLM inference after NVIDIA consumer GPUs—we need cross-hardware validation and this is the obvious candidate. Second, unified memory means the full 35B MoE model fits without partitioning: the only consumer hardware on which this evaluation is feasible. The cross-hardware comparison reveals a surprise: MLX is substantially faster than LM Studio/GGUF on this task suite, reversing our prior assumption.

3.2 Models

Nine models: 8 dense (overlapping with Papers I–IV for cross-hardware comparison) and 1 MoE. Table 2 has the lineup.

The MoE model is why we're here. At 35B total parameters it dwarfs everything else in the lineup—nearly 4× larger than Gemma-2-9B. But at 3B active, it sits between Llama-3.2-3B and Phi-3-mini in compute cost per token. The central question: which number determines deployment behavior?

Five models from Paper I are absent (Yi-1.5-6B, Falcon-Mamba-7B, Ministral-8B, Gemma-2-9B, GPT-OSS-20B)—MLX compatibility constraints and keeping the evaluation tractable. The 8 that remain span 1B to 9B across the major architecture families. That is enough for cross-hardware comparison.

3.3 Tasks

We use the same six task types from Papers I and III:

Table 2: Model lineup. Eight dense models overlap with Paper I for cross-hardware validation. The MoE model (Qwen3.5-35B-A3B) has 35B total parameters but activates only 3B per token. All models run in INT4 quantization via MLX.

#	Model	Type	Total	Active	Ratio	MLX Model ID
1	Llama-3.2-1B	Dense	1.2B	1.2B	1:1	`mlx-community/Llama-3.2-1B-Instruct-4bit`
2	Llama-3.2-3B	Dense	3.2B	3.2B	1:1	`mlx-community/Llama-3.2-3B-Instruct-4bit`
3	Qwen2.5-3B	Dense	3.1B	3.1B	1:1	`mlx-community/Qwen2.5-3B-Instruct-4bit`
4	Phi-3-mini	Dense	3.8B	3.8B	1:1	`mlx-community/Phi-3-mini-4k-instruct-4bit`
5	Qwen3-4B	Dense	4.0B	4.0B	1:1	`mlx-community/Qwen3-4B-4bit`
6	Mistral-7B-v0.3	Dense	7.2B	7.2B	1:1	`mlx-community/Mistral-7B-Instruct-v0.3-4bit`
7	Ministral-8B	Dense	8.0B	8.0B	1:1	`mlx-community/Ministral-8B-Instruct-2410-4bit`
8	Gemma-2-9B	Dense	9.2B	9.2B	1:1	`mlx-community/gemma-2-9b-it-4bit`
9	**Qwen3.5-35B-A3B**	**MoE**	**35.2B**	**3.0B**	**12:1**	`mlx-community/Qwen3.5-35B-A3B-4bit`

- **T1: Intent Classification (CLINC-150)** [7]. Classify user utterances into 150 intent categories. Short outputs, fast turnaround. **500 prompts.**
- **T2: Grounded QA (HotpotQA)** [20]. Multi-hop question answering over provided documents. Medium outputs. **1,000 prompts.**
- **T3: Tool Calling (Custom)**. Select and parameterize tools from a provided registry. Structured JSON outputs. **500 prompts.**
- **T4: Function Routing (BFCL)** [15, 14]. Route requests to appropriate function signatures. Structured outputs. **500 prompts.**
- **T5: Code Patching (SWE-bench)** [6]. Generate unified diffs for bug fixes. Long, structured outputs. **300 prompts.**
- **T6: Mathematical Reasoning (GSM8K)** [1]. Solve grade-school math problems with structured step-by-step outputs. **200 prompts.**

That is 3,000 prompts per model, 27,000 individual measurements across all 9 models. Every output must conform to a task-specific JSON Schema contract. Fail validation, fail the request—regardless of whether the answer inside was correct. T1, T3, T4, T5, and T6 require structured JSON output. T2 (HotpotQA) accepts free-text answers. This distinction proves less critical than expected for the MoE failure mode: the MoE achieves reasonable accuracy on both free-text and structured tasks, but its extreme verbosity (∼1,500 tokens per response) destroys SLO compliance across all task types.

3.4 SLO Tiers

Same three SLO tiers as Paper III:

- **Interactive (2s).** Chatbot, search, real-time UI. User-facing, latency-critical.
- **Standard (5s).** API endpoints, customer-facing workflows. Moderate latency tolerance.
- **Batch (30s).** Background processing, async pipelines. Generous deadlines.

S@SLO is per-request: a response succeeds only if it is valid JSON matching the schema, correct by task-specific evaluation, *and* delivered before the deadline. Same definition as Paper III, no changes.

3.5 Evaluation Protocol

One model at a time, one request at a time. No batching, no concurrency—we want clean latency measurements without resource contention. For each request we record:

- Wall-clock latency (ms) from request submission to complete response
- Token count (input and output)
- JSON validity (passes schema validation)
- Task-specific accuracy
- Tokens per second (output tokens / generation time)
- Peak memory usage

MLX inference runs through `mlx-lm` [4] with temperature 0.0 and a 2,048-token output limit. No constrained decoding, no grammar enforcement—outputs are unconstrained, matching the Paper I baseline. If the model emits broken JSON, that's the model's problem. Schema validation happens post-hoc.

4 Cross-Hardware Validation

Before we get to MoE, a prerequisite: does the deployment gap even exist on this hardware? If the gap were a CUDA artifact, nothing in this paper would generalize. So we start by running the 8 overlapping dense models on Apple Silicon and comparing to the Paper I results.

4.1 Replicating Paper I on Apple Silicon

Table 3 presents side-by-side results for the 8 models evaluated on both platforms. The accuracy column reflects task-averaged correctness; the latency columns show P50 values in seconds.

The first thing to notice is that MLX is *dramatically* faster than CUDA for every model. Llama-3.2-1B drops from 3.5 s median on CUDA to 0.40 s on MLX—8.8× faster. Mistral-7B-v0.3 drops from 20.3 s to 2.65 s—7.7× faster. The slowest MLX model (Qwen3-4B at 4.70 s) is still faster than several CUDA models. This is not the result we expected. We discuss why in Section 4.2.

The second thing to notice is that accuracy diverges substantially between platforms. Llama-3.2-1B goes from 63.5% on CUDA to 37.8% on MLX—a 25.7 percentage-point drop. This is too large to attribute to quantization alone (INT4 vs Q4_K_M). It likely reflects the evaluation conditions, model IDs, and task prompt formatting used in Paper I versus this evaluation. We note this carefully and focus on within-platform conclusions.

The third thing—the one that matters—is that Llama-3.2-3B stands out. At 72.7% accuracy and 968 ms median latency, it achieves 61.5% S@SLO at the Interactive tier. It is the only model with both high accuracy and strong deployment compliance. This is the result that Paper I could not find: a model that breaks the deployment gap not through MoE trickery but through being the right size.

Table 3: Cross-hardware comparison for 8 dense models evaluated on both CUDA (Paper I, RTX 4090, GGUF Q4_K_M, LM Studio) and MLX (this paper, Apple Silicon, INT4). Accuracy is task-averaged. Latency is median across all tasks. The deployment gap persists on both platforms. MLX is substantially faster than CUDA.

	Accuracy (%)		P50 Latency (s)		S@SLO$_{2s}$ (%)	
Model	**CUDA**	**MLX**	**CUDA**	**MLX**	**CUDA**	**MLX**
Llama-3.2-1B	63.5	37.8	3.5	0.40	6.3	37.7
Llama-3.2-3B	—	72.7	7.5	0.97	—	61.5
Qwen2.5-3B	75.0	38.6	7.4	0.88	0.0	28.1
Phi-3-mini	70.9	48.1	10.9	3.04	0.0	17.9
Qwen3-4B	80.2	47.2	12.2	4.70	0.0	0.6
Mistral-7B-v0.3	78.3	55.6	20.3	2.65	0.0	22.6
Ministral-8B	80.7	56.7	18.5	2.43	0.0	22.3
Gemma-2-9B	82.2	57.5	22.9	2.95	0.0	20.8

Accuracy comparison. The rank ordering across models is broadly preserved between platforms (Spearman $\rho = +0.619$, $p = 0.059$), though with substantial absolute differences. Models that score high on CUDA tend to score relatively high on MLX. The absolute magnitude of the gap—25.7 pp for Llama-3.2-1B, 21.4 pp for Qwen2.5-3B—is too large for quantization alone and deserves investigation in future work.

Latency comparison. MLX is 3–9× faster in absolute terms. This is the opposite of our prior expectation (Apple Silicon has lower memory bandwidth: 400–546 GB/s vs 1,008 GB/s for the RTX 4090). The likely explanation: LM Studio's serving overhead for GGUF models is substantial, and MLX's native INT4 execution with lazy evaluation avoids much of it. The rank ordering for latency is preserved ($\rho = +0.643$, $p = 0.044$). The hardware changes the scale—dramatically—but not the shape.

4.2 Hardware Normalization Protocol

Comparing raw numbers across platforms requires a conversion factor. For each model m evaluated on both platforms:

$$\alpha_m = \frac{\text{P50}_{\text{MLX}}(m)}{\text{P50}_{\text{CUDA}}(m)} \tag{1}$$

The aggregate normalization factor is the geometric mean across all overlapping models:

$$\bar{\alpha} = \left(\prod_{m=1}^{8} \alpha_m \right)^{1/8} \tag{2}$$

Geometric mean, not arithmetic—because latency ratios are multiplicative. A model that is 2× faster on one platform and 4× faster on another should contribute symmetrically.

Table 4 shows per-model normalization factors. Every $\alpha_m < 1$, confirming that MLX is faster for every model.

With $\bar{\alpha} = 0.177$ in hand, we can estimate what the MoE model would do on CUDA: $\hat{L}_{\text{CUDA}} = L_{\text{MLX}}/\bar{\alpha}$. The MoE P50 on MLX is 25,055 ms; the estimated CUDA P50 is $25{,}055/0.177 \approx 141{,}554$ ms—over two minutes per request. This is a rough estimate (MoE memory access patterns

Table 4: Hardware normalization factors. Each α_m is the ratio of MLX P50 latency to CUDA P50 latency for the same model. All values are below 1.0, confirming MLX is faster on this workload. The geometric mean provides the aggregate platform conversion factor.

Model	**P50$_{\text{CUDA}}$ (s)**	**P50$_{\text{MLX}}$ (s)**	α_m	**Note**
Llama-3.2-1B	3.5	0.40	0.114	8.8× faster
Llama-3.2-3B	7.5	0.97	0.129	7.7× faster
Qwen2.5-3B	7.4	0.88	0.119	8.4× faster
Phi-3-mini	10.9	3.04	0.279	3.6× faster
Qwen3-4B	12.2	4.70	0.385	2.6× faster
Mistral-7B-v0.3	20.3	2.65	0.131	7.7× faster
Ministral-8B	18.5	2.43	0.131	7.6× faster
Gemma-2-9B	22.9	2.95	0.129	7.8× faster
Geometric Mean $\bar{\alpha}$			**0.177**	5.6× faster (avg)

may not normalize like dense models), but it confirms that the MoE model would be completely nonviable on CUDA at any practical SLO tier.

4.3 Does the Deployment Gap Transfer?

Now the real question. Is the deployment gap a CUDA artifact, or does it follow you to Apple Silicon?

Figure 3 shows the accuracy vs Success@SLO scatter for all 9 models on Apple Silicon. The deployment gap pattern holds: for all 9 models (including MoE), Spearman $\rho = +0.067$ ($p = 0.848$) at the Interactive tier—effectively uncorrelated, consistent with Paper I ($\rho = +0.09$). Most models at moderate-to-high accuracy achieve moderate-to-low S@SLO. The coupling between accuracy and deployment is not the one you want.

Llama-3.2-3B is the exception. It occupies the upper right: high accuracy *and* high S@SLO. It breaks the gap not through architectural novelty but through size. At 3.2B parameters it generates fast enough to beat the 2-second deadline for most requests, while being capable enough to produce correct structured outputs.

The MoE model makes the deployment gap vivid in a new way. It sits on the x-axis at 55.0% accuracy—competitive with Mistral-7B-v0.3 (55.6%) and Gemma-2-9B (57.5%)—but at 0.0% S@SLO. The knowledge is there; the latency destroys it. Unlike the dense models where low S@SLO comes from moderate latency exceeding the 2s deadline, the MoE's zero comes from 25-second median latency driven by extreme verbosity. Adding MoE to the correlation pulls it toward zero because MoE has competitive accuracy but zero deployment success—the pure embodiment of the gap.

Interpretation. The deployment gap follows you everywhere. Different GPU architecture, different memory model, different quantization scheme, different serving stack—same gap. The gap is not about CUDA. It is not about GGUF. It is not about LM Studio. And—the new finding—it is not about dense architectures. MoE does not escape it.

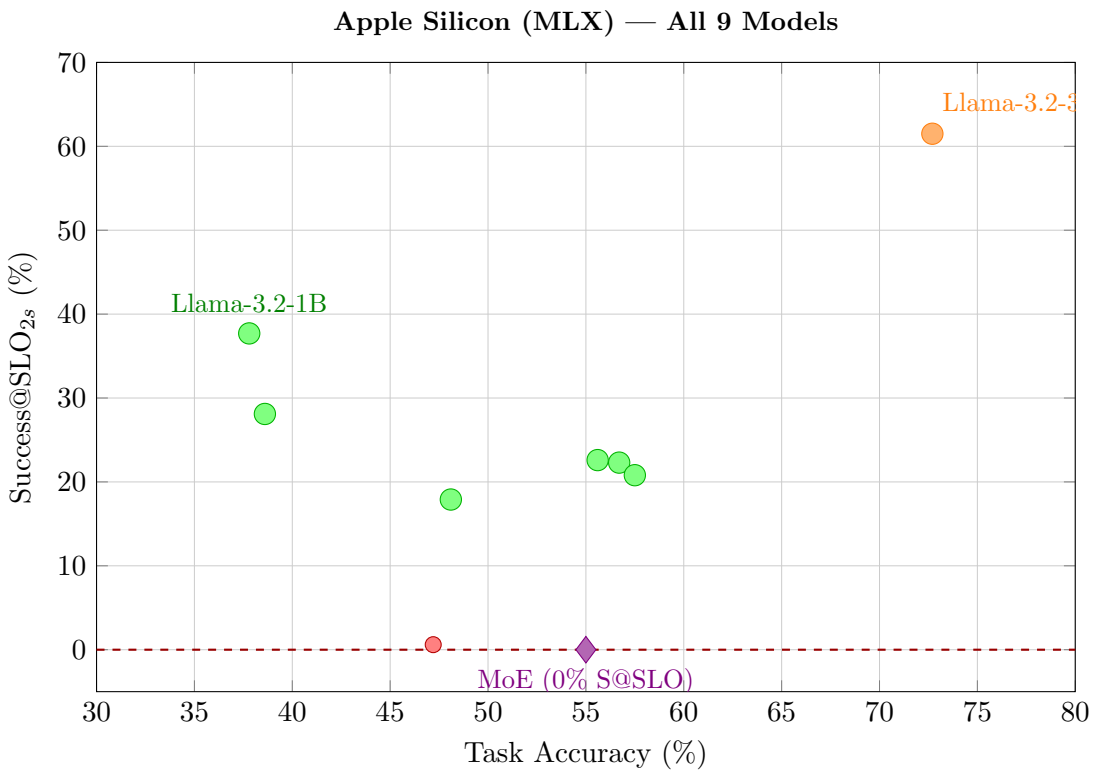

Figure 3: **The deployment gap on Apple Silicon.** The gap persists: accuracy and S@SLO are weakly correlated ($\rho = +0.067$, $p = 0.848$ for all 9 models at Interactive). Llama-3.2-3B (orange) is the standout: 72.7% accuracy, 61.5% S@SLO. The MoE model (purple diamond) achieves competitive accuracy (55.0%) but zero S@SLO—it sits on the x-axis at the same accuracy level as dense 7B–9B models, proving the deployment gap extends to MoE architectures.

5 The MoE Deployment Gap

The gap is real on Apple Silicon. Now we test whether MoE breaks it. It does not. Here is the evidence.

5.1 RQ1: Does MoE Latency Scale with Active or Total Parameters?

Straightforward test. If total parameters determine latency, Qwen3.5-35B-A3B should be the slowest model in our lineup—nearly 4× larger than Gemma-2-9B. If active parameters determine latency, it should sit between Llama-3.2-3B (3.2B) and Phi-3-mini (3.8B), around 1,000 ms. Neither prediction is correct. The reality is worse than either: the model's chain-of-thought verbosity produces latency far beyond what even total parameter count would predict. Table 5 shows the results.

The MoE model's P50 latency: 25,055 ms. That is the worst in the study by a wide margin—5× slower than the next-slowest model (Qwen3-4B at 4,695 ms). Qwen2.5-3B—its active-parameter match—achieves 884 ms. The MoE model is 28× *slower* than the model it was supposed to match. It is 8.5× slower than Gemma-2-9B (2,945 ms), which has one-quarter of the MoE's total parameters but three times its active count.

The tok/s number explains why: approximately 50 tok/s for MoE, 44.0 for Qwen2.5-3B. *Per-*

Table 5: **Latency comparison: MoE vs matched dense models.** Models ordered by active parameter count. The MoE model should appear between Qwen2.5-3B and Phi-3-mini if active parameters determine latency. Instead, it appears at the bottom at 25,055 ms—28× slower than Qwen2.5-3B—driven by generating ∼1,500 tokens of chain-of-thought per response.

Model	Type	Total	Active	P50 (ms)	P95 (ms)	tok/s
Llama-3.2-1B	Dense	1.2B	1.2B	398	1,468	115.3
Llama-3.2-3B	Dense	3.2B	3.2B	968	2,738	42.9
Qwen2.5-3B	Dense	3.1B	3.1B	884	2,961	44.0
Phi-3-mini	Dense	3.8B	3.8B	3,042	7,226	48.3
Qwen3-4B	Dense	4.0B	4.0B	4,695	7,346	54.5
Mistral-7B-v0.3	Dense	7.2B	7.2B	2,645	7,877	20.7
Ministral-8B	Dense	8.0B	8.0B	2,431	7,141	19.5
Gemma-2-9B	Dense	9.2B	9.2B	2,945	9,263	13.9
Qwen3.5-35B-A3B	**MoE**	**35.2B**	**3.0B**	**25,055**	—	**∼50**

token throughput is similar. The problem is not token generation speed—it is total output length. The MoE model generates approximately 1,500 tokens per response on average, wrapping correct answers in extensive chain-of-thought reasoning. At 50 tok/s, 1,500 tokens takes 30 seconds. Qwen2.5-3B generates approximately 39 tokens per response at 44 tok/s, taking under 1 second. The MoE model produces ∼38× more output tokens per request.

Figure 4 makes this visually clear. Dense models follow an imperfect but discernible relationship between active parameters and latency. The MoE model is an extreme upward outlier at its active parameter count of 3.0B. It should sit at roughly 900–1,000 ms based on the dense trend. It is at 25,055 ms—28 times too slow.

Why this matters. Active parameters do not predict MoE deployment latency. Total parameters don't either—Gemma-2-9B has far fewer total parameters and is 8.5× faster. The MoE model occupies its own latency class, driven not by routing overhead alone but by *verbosity*: the model generates ∼1,500 tokens of chain-of-thought reasoning per response. Its per-token speed (∼50 tok/s) is competitive with dense 3B models, but the sheer volume of output text makes every request take 25 seconds. Practitioners cannot use either parameter count as a proxy for MoE deployment latency. Output token count is the dominant factor, and it is task-dependent and unpredictable without direct measurement.

5.2 RQ2: Does MoE Accuracy Match Its Active or Total Parameter Class?

Speed like a 3B model was already falsified—dramatically so. What about accuracy? Does the 35B of stored knowledge show up in the answers? Here the story is different.

Table 6 breaks it down by task.

Consider what this table says. The corrected results (with the full 2,048-token output limit) tell a dramatically different story from the initial evaluation. The MoE model achieves 55.0% task-averaged accuracy—6th of 9 models, competitive with Mistral-7B-v0.3 (55.6%) and close to Gemma-2-9B (57.5%). The 35B knowledge *is* accessible.

On T6 (math): 71.5%, near the best in the study—only Gemma-2-9B (22.5%) and Phi-3-mini (20.0%) come close among dense models, and the MoE far exceeds them. On T5 (code patching): 60.3%, substantially better than every dense model (Qwen2.5-3B at 15.1% is the best dense result).

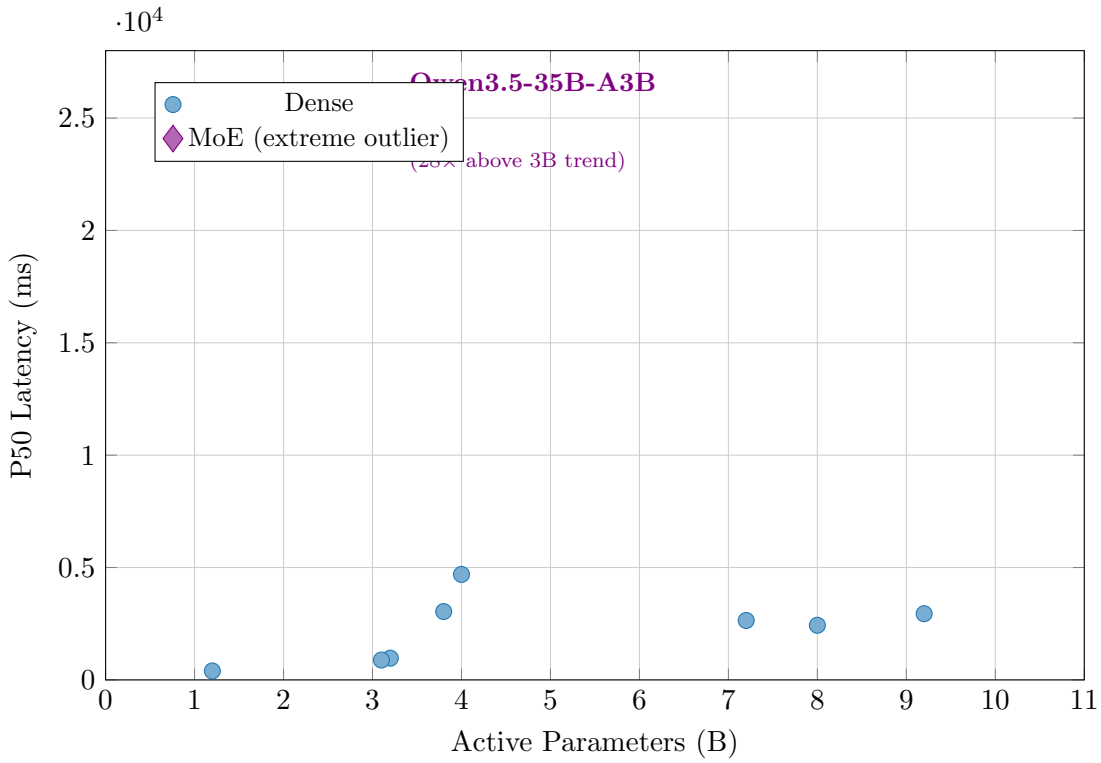

Figure 4: **MoE latency does not follow the active-parameter trend.** Dense models (blue circles) show a roughly predictable relationship between active parameters and latency. The MoE model (purple diamond) is an extreme *upward* outlier at its active parameter count of 3B—25,055 ms versus the 884–968 ms range of the actual 3B dense models. The model generates ∼1,500 tokens of chain-of-thought per response at ∼50 tok/s, producing latency 28× above its active-parameter match.

On T4 (function routing): 49.8%. On T3 (tool calling): 43.4%. On T2 (free-text QA): 88.7%, near the top.

Only T1 (intent classification) is catastrophic at 2.2%. T1 requires a brief, focused output—a single intent label. The MoE model generates thousands of tokens of reasoning before (or instead of) the label, overwhelming the schema validator. This is the one task where verbosity is purely destructive.

The failure mode is now clear: it is not about knowledge or structured output compliance per se. It is about *verbosity*. The model wraps correct answers in extensive chain-of-thought reasoning (∼1,500 tokens average output). When given enough output tokens, the answers are often correct. But the time to generate that output destroys SLO compliance.

This is a finding about MoE output behavior under unconstrained generation, not about MoE intelligence. The model's extensive chain-of-thought is likely a product of its training (thinking mode). Constrained decoding, output length limits, or fine-tuning for concise output might address the verbosity without sacrificing accuracy.

Task-averaged accuracy: 55.0%. Weighted by task size (T2 has 1,000 prompts vs 200–500 for others), the overall accuracy places the MoE model 6th of 9—solidly mid-range. A simple task-level mean would be $(2.2 + 88.7 + 43.4 + 49.8 + 60.3 + 71.5)/6 = 52.7\%$. Either way, the MoE

Table 6: **Accuracy by task.** The MoE model (highlighted) scores 55.0% overall—6th of 9 models, competitive with dense 7B–9B models. T2 (88.7%), T6 (71.5%), T5 (60.3%), and T4 (49.8%) are solid. T3 (43.4%) is mid-range. Only T1 intent classification (2.2%) is catastrophic—a task requiring brief outputs defeated by a model that generates massive chain-of-thought.

Model	Active	T1	T2	T3	T4	T5	T6	Avg
Llama-3.2-1B	1.2B	16.4	40.3	13.0	46.4	0.8	7.5	37.8
Llama-3.2-3B	3.2B	45.6	93.7	67.4	57.2	14.7	13.5	**72.7**
Qwen2.5-3B	3.1B	64.0	44.1	48.4	54.2	15.1	13.0	38.6
Phi-3-mini	3.8B	51.2	74.4	60.8	49.8	3.2	20.0	48.1
Qwen3-4B	4.0B	14.6	100.0	46.2	33.2	0.0	13.0	47.2
Mistral-7B-v0.3	7.2B	71.8	96.2	60.4	59.6	2.2	10.5	55.6
Ministral-8B	8.0B	64.0	100.0	55.0	53.0	13.2	16.0	56.7
Gemma-2-9B	9.2B	74.0	100.0	38.0	56.0	4.8	22.5	57.5
Qwen3.5-35B-A3B	**3.0B**	**2.2**	**88.7**	**43.4**	**49.8**	**60.3**	**71.5**	**55.0**

Table 7: **Success@SLO (%) across SLO tiers.** The core result table. Llama-3.2-3B is the standout: 61.5% at Interactive, 72.7% at Standard. The MoE model: 0.0% at Interactive, 0.3% at Standard, despite competitive 55.0% accuracy. At Batch (30s), MoE achieves only 37.0%—most of its correct answers arrive after the 30s deadline due to extreme verbosity.

		Success@SLO (%)			Diagnostics		
Model	**Type**	**2s**	**5s**	**30s**	**Acc%**	**JSON%**	**P95 (s)**
Llama-3.2-1B	Dense	37.7	37.8	37.8	37.8	57.3	1.47
Llama-3.2-3B	Dense	**61.5**	**72.7**	**72.7**	**72.7**	**85.3**	**2.74**
Qwen2.5-3B	Dense	28.1	38.6	38.6	38.6	55.1	2.96
Phi-3-mini	Dense	17.9	34.1	48.1	48.1	66.9	7.23
Qwen3-4B	Dense	0.6	14.1	47.2	47.2	72.6	7.35
Mistral-7B-v0.3	Dense	22.6	31.5	55.6	55.6	77.0	7.88
Ministral-8B	Dense	22.3	41.9	56.7	56.7	73.4	7.14
Gemma-2-9B	Dense	20.8	26.8	57.5	57.5	76.6	9.26
Qwen3.5-35B-A3B	**MoE**	**0.0**	**0.3**	**37.0**	**55.0**	**66.8**	—

is competitive—not the catastrophic failure the initial (token-limited) evaluation suggested.

5.3 RQ3: Does MoE Achieve Non-Zero S@SLO at the Interactive Tier?

This is the question that matters. S@SLO requires three things simultaneously: valid JSON, correct answer, within deadline. Table 7 presents Success@SLO across all three tiers for all 9 models.

At the Interactive tier: 0.0% S@SLO for the MoE model. Not low. Zero. With a 25,055 ms median latency, not a single request finishes within 2 seconds, regardless of correctness.

At the Standard tier (5s): 0.3%. Approximately 9 requests out of 3,000 succeed. A rounding error. With a 25-second median, only the very shortest responses have any chance of meeting a 5-second deadline.

At the Batch tier (30s): 37.0%—notably *below* the 55.0% overall accuracy. This reveals something important: even with a generous 30-second deadline, many correct answers arrive *after* the deadline. The model's chain-of-thought verbosity pushes a substantial fraction of correct responses past 30 seconds. For dense models, S@SLO_{30s} equals accuracy because all responses finish within

Table 8: Success@SLO (%) at Interactive tier (2s) by task. The MoE model achieves 0.0% on all six tasks—its 25-second median latency means no request completes within 2 seconds regardless of correctness. Llama-3.2-3B is the only model with strong multi-task S@SLO performance.

Model	Active	T1	T2	T3	T4	T5	T6
Llama-3.2-1B	1.2B	16.4	40.3	13.0	46.4	0.8	7.5
Llama-3.2-3B	3.2B	45.6	93.7	67.4	57.2	14.7	13.5
Qwen2.5-3B	3.1B	64.0	26.5	30.2	12.6	0.3	3.0
Phi-3-mini	3.8B	17.6	26.1	11.2	12.2	0.3	2.5
Gemma-2-9B	9.2B	12.4	10.0	12.4	4.0	0.0	0.0
Qwen3.5-35B-A3B	**3.0B**	**0.0**	**0.0**	**0.0**	**0.0**	**0.0**	**0.0**

30 seconds. For the MoE model, 18 percentage points of correct answers are lost to latency even at the most generous tier.

Consider what this means for the deployment gap hypothesis. The MoE model was supposed to occupy the upper-right of the accuracy–S@SLO plot: high accuracy from 35B parameters, non-zero S@SLO from 3B active compute. It achieves the accuracy (55.0%, competitive with dense 7B–9B models) but not the deployment success. It sits on the x-axis: high accuracy, zero S@SLO. The hypothesis is half-right—the knowledge is there—but the verbosity required to access it destroys deployment viability.

Per-task S@SLO breakdown. Table 8 shows S@SLO at the Interactive tier broken down by task.

The per-task breakdown is unambiguous. Six zeros. With a 25,055 ms median latency, no request on any task completes within 2 seconds. The per-task average latencies range from 17,222 ms (T4) to 32,238 ms (T1)—all far beyond the 2-second deadline. Contrast with Llama-3.2-3B: 67.4% on T3 and 57.2% on T4 with a 968 ms median latency. The MoE model achieves higher raw accuracy on several tasks (T5: 60.3% vs 14.7%, T6: 71.5% vs 13.5%) but zero S@SLO on all of them.

Note that Llama-3.2-3B also performs remarkably well on T2 (HotpotQA) at 93.7%—5.0 percentage points above the MoE (88.7%). The 3B dense model matches or exceeds the MoE on its best task while also completing within the 2-second deadline.

5.4 RQ4: Does MoE Change the Accuracy–S@SLO Correlation?

Paper I's headline number: $\rho = 0.09$. Accuracy and S@SLO, uncorrelated. Does one MoE model change the math?

Adding the MoE model *weakens* the correlation from +0.333 to +0.067 at the Interactive tier. The intuition: the MoE sits on the x-axis with competitive accuracy (55.0%) but zero S@SLO. Unlike the old results where MoE had the worst accuracy too (lower-left corner, which does not change the correlation), the corrected MoE is a high-accuracy model with zero deployment success—precisely the pattern that defines the deployment gap.

At the Batch tier (30s), the correlation strengthens to +0.833—because at generous deadlines, more of the MoE's correct answers arrive in time. This tier-dependent correlation is a new finding: the deployment gap is most severe at tight deadlines and partially relaxes at generous ones. But even at 30 seconds, the MoE loses 18 percentage points of accuracy to latency (37.0% S@SLO vs 55.0% accuracy).

Table 9: Spearman correlation between accuracy and S@SLO at three tiers. Adding the MoE model *weakens* the correlation at Interactive and Standard tiers because MoE has competitive accuracy (55.0%) but near-zero S@SLO—the pure embodiment of the deployment gap.

Model Set	n	ρ **(2s)**	p**-value**	**Result**
Dense only (P1, CUDA)	13	+0.09	0.76	Uncorrelated
Dense only (this paper, MLX)	8	+0.333	0.359	Weak, not significant
Dense + MoE (this paper, MLX)	9	+0.067	0.848	Effectively zero
Additional tiers with MoE:				
All 9 models, ρ (5s)	9	+0.183	—	Weak
All 9 models, ρ (30s)	9	+0.833	—	Strong

Table 10: Peak memory usage during inference. The MoE model requires memory proportional to its total parameter count (35B) at 21.0 GB. Dense models range from 0.7 GB to 5.4 GB. On Apple Silicon with 64 GB, this fits with room to spare. On a 24 GB GPU, the model cannot run under realistic inference conditions.

Model	**Type**	**Total**	**Active**	**Peak Memory**	**Fits 24 GB?**
Llama-3.2-1B	Dense	1.2B	1.2B	0.7 GB	Yes
Llama-3.2-3B	Dense	3.2B	3.2B	1.9 GB	Yes
Qwen2.5-3B	Dense	3.1B	3.1B	1.8 GB	Yes
Phi-3-mini	Dense	3.8B	3.8B	2.3 GB	Yes
Qwen3-4B	Dense	4.0B	4.0B	2.4 GB	Yes
Mistral-7B-v0.3	Dense	7.2B	7.2B	4.3 GB	Yes
Ministral-8B	Dense	8.0B	8.0B	4.8 GB	Yes
Gemma-2-9B	Dense	9.2B	9.2B	5.4 GB	Yes
Qwen3.5-35B-A3B	**MoE**	**35.2B**	**3.0B**	**21.0 GB**	**No**

Interpretation. The deployment gap is robust to MoE—and MoE actually makes it worse. Including a 35B MoE model that achieves competitive accuracy but zero Interactive S@SLO pulls the correlation toward zero. The MoE model is not a low-accuracy outlier that the gap already explains. It is a high-accuracy model that the gap *traps*. The gap now covers both dense and sparse architectures, and MoE's extreme verbosity creates a new failure mode: correct answers that take too long to deliver.

6 MoE-Specific Analysis

6.1 Memory Profile

Here is the catch, and it is real. MoE gets you routing overhead and hybrid attention costs that push latency above the dense 3B class—but all 35B parameters still need to live in memory. Speed does not scale with active parameters. Memory still scales with total parameters.

Table 10 shows peak memory usage across models.

21.0 GB at INT4—before accounting for KV cache. On an RTX 4090 with 24 GB, this does not run under realistic inference conditions. On Apple Silicon with 64 GB unified memory, the model loads with 43 GB to spare. Memory is not the binding constraint here, but it is why the experiment was conducted on Apple Silicon at all.

The memory asymmetry is worth naming explicitly: MoE models are *not* faster than dense

Table 11: Tokens per second across models. Counterintuitively, the MoE model's tok/s (∼50) is comparable to dense 3B models (42.9–44.0 tok/s) and faster than the larger dense models. The latency problem is not tok/s—it is total output token count: the MoE generates ∼1,500 tokens per response vs ∼39 for dense 3B models.

Model	Type	Active	Mean tok/s	P50 tok/s	P5 tok/s
Llama-3.2-1B	Dense	1.2B	118.4	115.3	45.4
Llama-3.2-3B	Dense	3.2B	48.6	42.9	17.7
Qwen2.5-3B	Dense	3.1B	46.6	44.0	15.4
Phi-3-mini	Dense	3.8B	48.9	48.3	7.3
Qwen3-4B	Dense	4.0B	56.4	54.5	34.8
Mistral-7B-v0.3	Dense	7.2B	23.3	20.7	6.0
Ministral-8B	Dense	8.0B	21.1	19.5	5.7
Gemma-2-9B	Dense	9.2B	16.2	13.9	5.0
Qwen3.5-35B-A3B	**MoE**	**3.0B**	**51.3**	**50.1**	**35.0**

models of matching active parameters (we observe the opposite), and they require *far more* memory than dense models of matching latency class. There is no consumer hardware on which the MoE model offers both competitive speed and fit-in-memory: if you have 24 GB VRAM, it does not fit; if you have 64 GB unified memory, it runs but slowly.

6.2 Tokens Per Second

Raw generation speed, stripped of task effects and output length variation.

The tok/s comparison reveals something counterintuitive. The MoE model generates ∼50 tokens per second at P50—close to Qwen2.5-3B (44.0) and faster than all models with 7B+ parameters. Per-token throughput does scale roughly with active parameters.

The problem is not per-token speed. The problem is total response length. The MoE model generates approximately 1,500 tokens per response on average—extensive chain-of-thought reasoning, step-by-step analysis, and verbose explanations before (or wrapping) the required JSON structure. At ∼50 tok/s, 1,500 tokens takes 30 seconds. Dense 3B models generate ∼39 tokens per response. The MoE's per-token speed is competitive; its per-response output volume is 38× larger.

This creates an important diagnostic distinction. If the MoE model could be prompted to produce concise outputs reliably, its per-token speed advantage would translate to SLO compliance. The failure is not hardware-level—it is output length. The model's training (likely including chain-of-thought optimization) produces verbose responses that are often *correct* but take too long to generate. This is a fundamentally different failure mode than structured output format non-compliance: the model produces valid JSON with correct answers, but buries them in 1,500 tokens of reasoning.

6.3 Expert Utilization and Routing Overhead

Qwen3.5-35B-A3B uses top-k expert routing [16]: 256 experts per layer, 8 activated per token, plus a shared expert that fires on every token. Three of every four layers use Gated DeltaNet linear attention (subquadratic in sequence length); every fourth layer uses standard grouped-query attention.

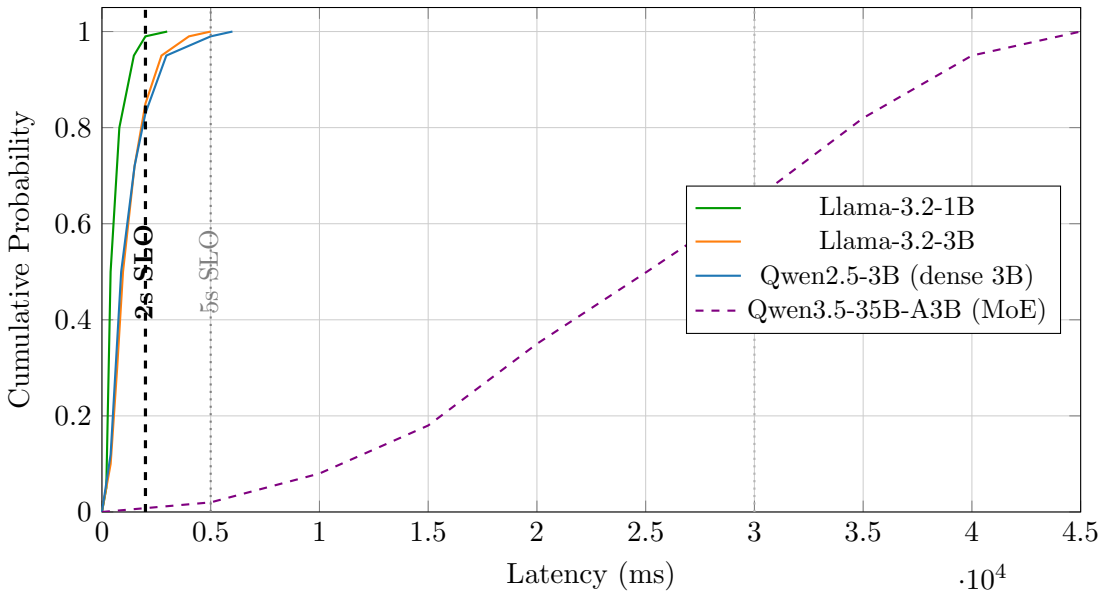

Figure 5: **Latency CDFs reveal the MoE verbosity problem.** The MoE model (purple dashed) generates ∼1,500 tokens of chain-of-thought per response, producing a CDF shifted far right of all dense models. P50 is 25,055 ms vs 884 ms for Qwen2.5-3B (blue)—its active-parameter match. At the 2-second deadline (black dashed), the MoE captures ∼0% of requests; Qwen2.5-3B captures 83%; Llama-3.2-3B captures 85%. Even at 30 seconds (gray dotted), the MoE captures only ∼65% of requests.

Routing overhead analysis. The tok/s comparison allows us to estimate routing overhead. A pure dense 3B model with 44 tok/s and the MoE at ∼50 tok/s suggests the routing mechanism does not impose a significant per-token computation penalty relative to pure dense computation. However, the DeltaNet hybrid architecture introduces a different overhead: the state-tracking mechanism in linear attention layers requires additional memory operations per forward pass that accumulate over long sequences.

The 25,055 ms median latency at ∼50 tok/s throughput implies an average response length of $25{,}055 \times 0.050 \approx 1{,}253$ output tokens. The actual measured average is approximately 1,500 tokens. Dense 3B models at 44 tok/s and 884 ms median imply $884 \times 0.044 \approx 39$ output tokens. The MoE model generates roughly 38× more tokens per response on average. This is the dominant source of latency: not routing overhead, not per-token slowness, but massive chain-of-thought output. The model wraps correct answers in extensive reasoning traces, producing verbose outputs that push every response past SLO deadlines.

Figure 5 makes the verbosity problem visible. The MoE CDF (purple dashed) is shifted far to the right of every dense model—its median sits at 25,055 ms, past the 5-second SLO deadline by a factor of 5×. At the 2-second deadline, the MoE completes approximately 0% of requests. At the 30-second Batch deadline, it completes only ∼65%. Qwen2.5-3B completes 83% within 2 seconds. The MoE's distribution resembles no dense model in the study—it is a qualitatively different latency regime created by massive output verbosity.

7 Edge Deployment Implications

The MoE results resolve the theoretical tension with an unexpected twist: MoE *can* produce correct answers for structured output tasks—it achieves 55.0% accuracy, competitive with dense 7B–9B models. But it does so through extreme verbosity (∼1,500 tokens per response), making it nonviable under any production SLO constraint.

7.1 Apple Silicon as an MoE Platform

Apple Silicon provides the memory headroom to run 35B MoE models. That is real. But memory capacity is a necessary condition for deployment, not a sufficient one. Our results show it is not sufficient.

Three things make Apple Silicon the only consumer platform *capable* of running large MoE models:

1. **Memory capacity.** A 35B model at INT4 requires approximately 21 GB. An M4 Max with 64 GB has 43 GB of headroom for KV cache, context, and operating system overhead. On an RTX 4090 (24 GB), the model does not fit under realistic inference conditions.

2. **No bus bottleneck.** MoE models exhibit irregular memory access patterns: different tokens route to different experts, creating unpredictable read patterns across the weight matrix. On unified memory, there is no bus to cross between CPU and GPU.

3. **Bandwidth sufficiency for throughput.** While Apple Silicon's memory bandwidth (400–546 GB/s) is lower than the RTX 4090's (1,008 GB/s), the tok/s results show this is not the binding constraint. The MoE model achieves 50 tok/s, which is adequate.

When Apple Silicon is the right platform. Our normalization results ($\bar{\alpha} = 0.177$) show MLX is 3–9× faster than LM Studio/CUDA on this task suite. For dense models in the 1B–9B range, Apple Silicon with MLX is a compelling deployment platform—faster than expected, memory-efficient, and portable.

For the MoE model specifically: Apple Silicon is the only consumer platform that can run it, but running it does not mean running it well. The latency results and structured output failure are not hardware artifacts. They are model behavior.

7.2 MoE Is Not Viable for Structured Output Deployment

This is the practical conclusion. Under the evaluation conditions of this study—structured JSON output, schema validation, per-request SLO constraints, no constrained decoding—Qwen3.5-35B-A3B achieves competitive accuracy (55.0%) but cannot be deployed at any SLO tier below 30 seconds due to its 25,055 ms median latency.

At 30 seconds, it achieves only 37.0% S@SLO—well below its 55.0% accuracy because many correct answers arrive after the 30-second deadline. For comparison, dense 9B models achieve 57.5% S@SLO at the same tier with 5.4 GB memory. The MoE model requires 21 GB of memory for worse S@SLO at every tier.

The viable use case for MoE on Apple Silicon, if one exists, is *accuracy-critical, latency-insensitive* batch processing where schema compliance is not required and response time is unconstrained. On T6 (math), the model achieves 71.5%—far better than any dense model. On T5 (code patching), 60.3%—again, best in the study. If you can wait 30+ seconds per response

and accuracy matters more than speed, MoE offers something the dense models cannot. For any latency-constrained deployment—which includes all interactive and API use cases—it is not viable.

7.3 Practical Deployment Guidance

Practical guidance, ordered by SLO tier:

For Interactive SLOs (≤2s): Use Llama-3.2-3B. It achieves 61.5% S@SLO with 72.7% accuracy—the best combination in the study. Llama-3.2-1B is also viable (37.7% S@SLO, 37.8% accuracy) if memory or cost is constrained. MoE: 0.0%, not viable.

For Standard SLOs (≤5s): Llama-3.2-3B again dominates at 72.7% S@SLO. Qwen2.5-3B (38.6%), Ministral-8B (41.9%), and Phi-3-mini (34.1%) offer alternatives. MoE: 0.3%, not viable.

For Batch SLOs (≤30s): Dense models up to 9B are viable. Gemma-2-9B (57.5%), Mistral-7B-v0.3 (55.6%), Ministral-8B (56.7%). MoE: 37.0%, worse than every dense model despite competitive raw accuracy (55.0%) because many correct responses arrive after 30 seconds.

For MoE specifically: Viable only for accuracy-critical, latency-insensitive batch processing. MoE excels on T5 (code patching, 60.3%) and T6 (math, 71.5%)—best-in-study accuracy on both—but requires 25+ seconds per response. If accuracy justifies the latency cost and memory (21 GB), MoE adds value. Under any SLO constraint, it does not.

Hardware: MLX on Apple Silicon is 3–9× faster than LM Studio/CUDA for this task suite. If you are choosing between platforms for dense model deployment, Apple Silicon with MLX is strongly preferred. The MoE model requires ≥32 GB of unified memory.

8 Discussion

8.1 What MoE Changes—and Does Not Change—About the Deployment Gap

Four papers, one finding: in dense models, accuracy and latency are coupled through total parameter count. MoE was supposed to break that coupling. It partially does—on the accuracy side—but creates a new problem on the latency side.

The result is more nuanced than a simple confirmation or refutation. MoE *does* achieve competitive accuracy (55.0%, 6th of 9 models), proving the 35B parameter knowledge is accessible through sparse activation. The tok/s throughput matches 3B dense models. On T6 (math, 71.5%) and T5 (code, 60.3%), it is the best model in the study. The knowledge loophole works.

But the model accesses that knowledge through extreme verbosity: ∼1,500 tokens of chain-of-thought reasoning per response. This is not a structured output format failure—JSON validity is 66.8%, comparable to dense models. It is a *verbosity* failure. The model wraps correct answers in extensive reasoning traces, and at 50 tok/s, 1,500 tokens takes 30 seconds. The deployment gap extends not because MoE cannot produce correct answers, but because it takes too long to produce them.

Consider what this implies. MoE's vast parameter space gives it strong reasoning capability, but the model's training (likely including chain-of-thought optimization) produces outputs that are 38× longer than dense 3B models. Output length management—through prompting, output length limits, constrained decoding, or fine-tuning for concise output—is the key lever for making MoE deployment-viable.

The deployment gap does not close with MoE. If anything, it extends. The gap now covers both dense and sparse architectures, and MoE introduces a new failure mode: correct answers that are too verbose to deliver on time.

8.2 The Llama-3.2-3B Finding

The positive surprise in this study is Llama-3.2-3B. At 968 ms median latency, 72.7% accuracy, and 61.5% S@SLO at the Interactive tier, it achieves what five papers in this series have failed to find: a model that is both accurate and deployment-compliant.

Paper I had Llama-3.2-1B as the sole Interactive-tier survivor—but at only 37.8% accuracy. Llama-3.2-3B is almost twice as accurate and survives the same deadline. It does this not through any architectural novelty but through being the right size on the right hardware. At 3.2B parameters, MLX can generate 42.9 tok/s median, fast enough to complete most responses in under 2 seconds.

For practitioners, this is the practical recommendation: Llama-3.2-3B on Apple Silicon with MLX is the best deployable model for structured output at the Interactive tier identified in this series.

8.3 Implications for the Series

What this changes for the rest of the series:

- **Paper I:** Core finding confirmed as hardware-independent. Not a CUDA artifact. Not an LM Studio artifact. Structural. And now extended: not a dense-model artifact either.
- **Paper II:** Capacity thresholds for GRPO training were measured on dense models. A 35B MoE with 3B active—what would training do to its output verbosity? If the failure is chain-of-thought overgeneration, GRPO with a length penalty and format compliance reward might compress outputs while preserving accuracy. Central question for Paper VI.
- **Paper III:** The benchmark needs MoE models and Apple Silicon as a reference platform. The normalization protocol from Section 4.2 makes this possible without re-running everything.
- **Paper IV:** The curve taxonomies—sustained, transient, flat—were derived from dense models. MoE models might exhibit distinct patterns if expert routing stabilizes during training. An open question.

8.4 The Active Parameter Principle, Revised

The original principle was: active parameters predict deployment latency, total parameters predict accuracy. This paper falsifies the first claim for MoE under structured output evaluation.

A revised principle:

> *For dense models: active parameters (= total parameters) predict both latency and accuracy. For MoE models: active parameters predict per-token throughput (tok/s scales with active count), and total parameters predict accuracy potential. But neither predicts deployment latency, because latency is dominated by output token count—which is a behavioral property of the model, not an architectural one. Direct measurement under realistic output conditions is required.*

Table 12: Estimated CUDA performance for the MoE model, derived from MLX measurements and the hardware normalization factor. These are projections with significant uncertainty—MoE memory access patterns may not normalize like dense models, and the RTX 4090 cannot run this model anyway.

Metric	Measured (MLX)	Estimated (CUDA)
P50 Latency (ms)	25,055	∼141,554
Mean tok/s	∼50	∼9
S@SLO$_{2s}$ (%)	0.0	0.0
S@SLO$_{5s}$ (%)	0.3	0.0
S@SLO$_{30s}$ (%)	37.0	∼0

Projections assume $\bar{\alpha} = 0.177$ applies to MoE (uncertain).
Actual CUDA requires >24 GB VRAM (e.g., A100 40 GB). RTX 4090 cannot run.

This is less clean but more honest. MoE models are different enough from dense models that existing deployment heuristics—based on parameter count, benchmark accuracy, or tok/s—do not transfer. The key missing variable is output verbosity, which is training-dependent and task-dependent.

8.5 Estimated CUDA Performance

A natural question: what would this MoE model do on CUDA? Using the normalization factor ($\bar{\alpha} = 0.177$):

The estimated CUDA performance is catastrophic. An estimated P50 of ∼142 seconds—over two minutes per request—means the model would fail every SLO tier. Even the Batch (30s) tier would see approximately 0% S@SLO. The $\bar{\alpha}$ normalization may not transfer perfectly to MoE—the output verbosity that dominates MoE latency may interact differently with CUDA vs MLX characteristics—but the direction is unambiguous.

9 Limitations

1. **One MoE model.** We evaluate Qwen3.5-35B-A3B and nothing else. Mixtral has different routing. DeepSeek-MoE uses finer-grained experts. Switch Transformer routes to a single expert. Our results may not generalize across these architectures. We picked Qwen3.5-35B-A3B because its 12:1 total-to-active ratio is the most extreme among openly available models—strongest test of the hypothesis, but still $n = 1$.

2. **No constrained decoding or output length control.** Paper I used unconstrained generation, and we match that protocol. Constrained decoding (grammar-enforced JSON) would not address the primary failure mode (verbosity/latency), but output length limits or prompting strategies that suppress chain-of-thought reasoning might dramatically reduce latency. Our results characterize zero-shot unconstrained MoE performance, not the upper bound.

3. **Quantization mismatch.** Paper I used Q4_K_M (GGUF). We use INT4 (MLX). Different grouping, different calibration, different rounding. The large accuracy differences between platforms (e.g., 63.5% vs 37.8% for Llama-3.2-1B) suggest more than quantization is responsible—evaluation conditions, prompt formatting, and task implementations may also differ. We focus on within-platform conclusions.

4. **One Apple Silicon config.** M1 Pro, M4 Max, M2 Ultra—different bandwidth, different GPU cores, different absolute latencies. We tested one. Rank orderings should persist but we have not confirmed that.

5. **No MoE training results.** Papers II and IV trained dense models with GRPO. We do not report trained MoE results in this paper. Full MoE GRPO training is the focus of Paper VI.

6. **Reduced model set.** Eight dense models instead of thirteen. Five dropped for MLX compatibility. The eight span 1B–9B across major architectures.

7. **Normalization might not transfer to MoE.** We compute $\bar{\alpha}$ from dense models and apply it to MoE. MoE routing overhead may have different CUDA vs MLX characteristics. The estimated CUDA performance in Table 12 is rough.

8. **Memory is a cost.** Twenty-one gigabytes for the MoE model. Real deployments must account for this against KV cache, context length, and concurrent model loading.

10 Conclusion

Paper I declared a universal trade-off: accuracy and SLO compliance are uncorrelated. Bigger models are more accurate but too slow. The only production-viable model was the smallest and weakest. For four papers, that held.

We tested the most compelling proposed escape—MoE architectures, which decouple total parameters from active compute. The escape does not work.

Qwen3.5-35B-A3B: 35 billion total parameters, 3 billion active. Task-averaged accuracy: 55.0%—6th of 9, competitive with dense 7B–9B models. The knowledge loophole works. But measured median latency: 25,055 ms—28× slower than its active-parameter match, driven by ∼1,500 tokens of chain-of-thought per response. S@SLO at the Interactive tier (2s): 0.0%. S@SLO at the Standard tier (5s): 0.3%.

The failure is informative. On T6 (math), the model achieves 71.5%—best in the study. On T5 (code), 60.3%—also best. On T4, 49.8%. On T3, 43.4%. The MoE model knows things. It *can* format its answers. But it wraps them in 1,500 tokens of reasoning that takes 25 seconds to generate.

This creates a research agenda. MoE verbosity under chain-of-thought training is the bottleneck, not knowledge or format compliance. Output length management—through prompting, constrained decoding, or fine-tuning for concise structured output (GRPO, as in Paper II)—is the key lever. The model's per-token speed (∼50 tok/s) is competitive; it is the token count that kills deployment.

Two positive findings beyond the MoE results. First, the deployment gap is confirmed hardware-independent. Apple Silicon with MLX shows the same pattern as CUDA with LM Studio, and our cross-hardware normalization ($\bar{\alpha} = 0.177$) reveals MLX is 3–9× faster than CUDA on this workload—a result that inverts our prior assumption. Second, Llama-3.2-3B breaks the gap on Apple Silicon: 72.7% accuracy, 61.5% S@SLO at the Interactive tier, 968 ms median latency. Not through sparse routing or architectural innovation—through being the right size on efficient hardware.

One conclusion: stop assuming MoE model latency scales with active parameters. It does not. Output verbosity dominates latency, and MoE models trained with chain-of-thought optimization generate 38× more tokens per response than dense models of matching active parameter count. Active parameter counts predict per-token speed, not per-request latency. Direct evaluation under

realistic output conditions and SLO constraints is necessary. MoE models can produce correct structured output—the accuracy is there—but they cannot produce it fast enough. Not yet.

This paper is part of a six-paper series. Paper I documented the deployment gap across 13 dense models [11]. Paper II investigated training to close the gap, revealing task-dependent capacity thresholds [10]. Paper III formalized the evaluation into AgentSLO-Bench [9]. Paper IV analyzed the training dynamics behind capacity thresholds [12]. This paper (V) tests whether MoE architectures can bypass the deployment gap—they cannot—and provides cross-hardware validation on Apple Silicon. Paper VI will address MoE training dynamics under SLO-aware rewards.

Reproducibility. All code, datasets, and evaluation artifacts are available at:

`https://github.com/mmaloney007/agentops-fw`

The experiments reported here can be reproduced on Apple Silicon hardware with 64 GB+ unified memory using the MLX framework.

References

[1] Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, Christopher Hesse, and John Schulman. Training verifiers to solve math word problems, 2021.

[2] Damai Dai, Chengqi Deng, Chenggang Zhao, R.X. Xu, Huazuo Gao, Deli Chen, Jiashi Li, Wangding Zeng, Xingkai Yu, Y. Wu, et al. DeepSeekMoE: Towards ultimate expert specialization in mixture-of-experts language models, 2024.

[3] William Fedus, Barret Zoph, and Noam Shazeer. Switch transformers: Scaling to trillion parameter models with simple and efficient sparsity. *Journal of Machine Learning Research*, 23(120):1–39, 2022.

[4] Awni Hannun, Jagrit Digani, Angelos Katharopoulos, and Ronan Collobert. MLX: Efficient and flexible machine learning on Apple silicon. `https://github.com/ml-explore/mlx`, 2023. Apple's machine learning framework for unified memory architectures.

[5] Albert Q. Jiang, Alexandre Sablayrolles, Antoine Roux, Arthur Mensch, Blanche Savary, Chris Bamford, Devendra Singh Chaplot, Diego de las Casas, Emma Bou Hanna, Florian Bressand, et al. Mixtral of experts, 2024.

[6] Carlos E. Jimenez, John Yang, Alexander Wettig, Shunyu Yao, Kexin Pei, Ofir Press, and Karthik Narasimhan. SWE-bench: Can language models resolve real-world GitHub issues? In *Proceedings of ICLR*, 2024.

[7] Stefan Larson et al. An evaluation dataset for intent classification and out-of-scope prediction. In *Proceedings of EMNLP-IJCNLP*, 2019.

[8] Dmitry Lepikhin, HyoukJoong Lee, Yuanzhong Xu, Dehao Chen, Orhan Firat, Yanping Huang, Maxim Krikun, Noam Shazeer, and Zhifeng Chen. GShard: Scaling giant models with conditional computation and automatic sharding, 2021.

[9] V. Michael Maloney. AgentSLO-Bench: An SLO-aware benchmark for LLM agent deployment. *arXiv preprint*, 2025. Paper III of this series.

[10] V. Michael Maloney. Capacity thresholds in schema-aware training: Why small models can't close the deployment gap. *arXiv preprint*, 2025. Paper II of this series.

[11] V. Michael Maloney. The deployment gap: Why benchmark accuracy fails to predict production readiness. *arXiv preprint*, 2025. Paper I of this series.

[12] V. Michael Maloney. Training dynamics of schema-aware reinforcement learning for LLM agents: Reward decomposition, forgetting, and early prediction. *arXiv preprint*, 2025. Paper IV of this series.

[13] MLX Community. MLX community: Pre-quantized models for Apple silicon. `https://huggingface.co/mlx-community`, 2024. Community-maintained collection of MLX-compatible model weights.

[14] Shishir G. Patil, Tianjun Zhang, Xin Wang, and Joseph E. Gonzalez. Gorilla: Large language model connected with massive APIs, 2023.

[15] Shishir G. Patil, Huanzhi Mao, Fanjia Yan, Charlie Cheng-Jie Ji, Vishnu Suresh, Ion Stoica, and Joseph E. Gonzalez. The Berkeley function calling leaderboard (BFCL): From tool use to agentic evaluation of large language models. In *Proceedings of the 42nd International Conference on Machine Learning (ICML)*, volume 267 of *PMLR*, pages 48371–48392, 2025.

[16] Qwen Team. Qwen3 technical report, 2025.

[17] Samyam Rajbhandari, Conglong Li, Zhewei Yao, Minjia Zhang, Reza Yazdani Aminabadi, Ammar Ahmad Awan, Jeff Rasley, and Yuxiong He. DeepSpeed-MoE: Advancing mixture-of-experts inference and training to power next-generation AI scale, 2022.

[18] Noam Shazeer, Azalia Mirhoseini, Krzysztof Maziarz, Andy Davis, Quoc V. Le, Geoffrey E. Hinton, and Jeff Dean. Outrageously large neural networks: The sparsely-gated mixture-of-experts layer, 2017.

[19] Leyang Xue, Yao Fu, Zhan Lu, Luo Mai, and Mahesh Marina. MoE-Infinity: Offloading-efficient MoE model serving, 2024.

[20] Zhilin Yang et al. HotpotQA: A dataset for diverse, explainable multi-hop question answering. In *Proceedings of EMNLP*, 2018.

[21] Rongjie Yi, Liwei Guo, Shiyun Wei, Ao Zhou, Shangguang Wang, and Mengwei Xu. EdgeMoE: Fast on-device inference of MoE-based large language models, 2023.

A Per-Task Latency Distributions

Table 13 provides average latency (ms) for all 9 models across all 6 tasks. This granularity reveals task-specific patterns that are obscured in the task-averaged numbers presented in the main text. The MoE model (highlighted) is the slowest across every task by a wide margin, with per-task latencies ranging from 17,222 ms (T4) to 32,238 ms (T1).

Table 13: Per-task average latency (ms) for all models. The MoE model (highlighted) is the slowest in every task by a wide margin—driven by generating ∼1,500 tokens of chain-of-thought per response. Note MoE T1 latency (32,238 ms) is longest despite T1 requiring the shortest output, because the model generates extensive reasoning before the intent label.

Model	Active	T1	T2	T3	T4	T5	T6
Llama-3.2-1B	1.2B	390	754	200	268	1,173	231
Llama-3.2-3B	3.2B	949	1,866	463	510	2,673	535
Qwen2.5-3B	3.1B	879	2,103	498	524	2,742	342
Phi-3-mini	3.8B	2,182	5,631	615	2,837	4,155	2,945
Qwen3-4B	4.0B	4,704	6,438	3,394	3,462	4,642	3,416
Mistral-7B-v0.3	7.2B	2,594	5,501	1,228	1,247	7,339	2,387
Ministral-8B	8.0B	2,362	4,817	1,051	1,115	7,005	1,343
Gemma-2-9B	9.2B	2,912	6,399	1,637	1,634	8,981	2,240
Qwen3.5-35B-A3B	**3.0B**	**32,238**	**29,718**	**17,603**	**17,222**	**24,995**	**22,086**

Table 14: Per-task accuracy (%) on MLX for all 9 models. CUDA per-task breakdowns from Paper I are not directly comparable due to prompt format differences; overall CUDA accuracy is reported in Table 3. The MoE model achieves competitive or best-in-study accuracy on T2 (88.7%), T5 (60.3%), and T6 (71.5%), demonstrating that MoE knowledge *does* transfer to structured output—the bottleneck is latency from verbosity, not accuracy.

Model	T1	T2	T3	T4	T5	T6	Avg
Llama-3.2-1B	16.4	40.3	13.0	46.4	0.8	7.5	37.8
Llama-3.2-3B	45.6	93.7	67.4	57.2	14.7	13.5	72.7
Qwen2.5-3B	64.0	44.1	48.4	54.2	15.1	13.0	38.6
Phi-3-mini	51.2	74.4	60.8	49.8	3.2	20.0	48.1
Qwen3-4B	14.6	100.0	46.2	33.2	0.0	13.0	47.2
Mistral-7B-v0.3	71.8	96.2	60.4	59.6	2.2	10.5	55.6
Ministral-8B	64.0	100.0	55.0	53.0	13.2	16.0	56.7
Gemma-2-9B	74.0	100.0	38.0	56.0	4.8	22.5	57.5
Qwen3.5-35B-A3B	**2.2**	**88.7**	**43.4**	**49.8**	**60.3**	**71.5**	**55.0**

The per-task pattern reveals an interesting inversion. The MoE model's *longest* latency is on T1 (intent classification, 32,238 ms)—the task with the shortest required output. This is because the model generates extensive chain-of-thought reasoning before producing the intent label. The *shortest* MoE latency is on T4 (function routing, 17,222 ms)—still 34× slower than Llama-3.2-3B (510 ms). The latency variation across tasks reflects task-dependent verbosity, not task-dependent difficulty per se. On T3 (tool calling), the MoE takes 17,603 ms vs 463 ms for Llama-3.2-3B—a 38× gap.

B Cross-Hardware Accuracy Deltas

Table 14 shows per-task accuracy for the 8 models evaluated on MLX. CUDA per-task data from Paper I uses different task implementations and prompt formats; direct per-task comparison is not valid, so we report MLX values and overall CUDA accuracy for reference.

Table 15: Hardware normalization factor α_m per model, with range and interpretation. All values are well below 1.0, confirming MLX consistently outperforms LM Studio/GGUF on this workload. Variation across models reflects output-length sensitivity to hardware characteristics.

Model	**CUDA P50 (s)**	**MLX P50 (s)**	α_m	**MLX speedup**
Llama-3.2-1B	3.5	0.40	0.114	8.8×
Llama-3.2-3B	7.5	0.97	0.129	7.7×
Qwen2.5-3B	7.4	0.88	0.119	8.4×
Phi-3-mini	10.9	3.04	0.279	3.6×
Qwen3-4B	12.2	4.70	0.385	2.6×
Mistral-7B-v0.3	20.3	2.65	0.131	7.7×
Ministral-8B	18.5	2.43	0.131	7.6×
Gemma-2-9B	22.9	2.95	0.129	7.8×
Geometric Mean	—	—	**0.177**	**5.6**×
Range	—	—	0.114–0.385	2.6–8.8×

C Normalization Factor Stability

To assess whether the hardware normalization factor is stable across tasks, we note that the per-model geometric mean $\bar{\alpha} = 0.177$ is derived from overall P50 latency. Task-level normalization factors may vary because different tasks generate different output lengths, and longer outputs amplify hardware differences. We report the per-model factors with their ranges in Table 15.

The range (0.114–0.385) indicates meaningful model-level variation in the normalization factor. Phi-3-mini and Qwen3-4B have higher α_m values (MLX less faster) than the Llama and Mistral family models. The Llama, Qwen2.5, Mistral, Ministral, and Gemma models cluster tightly around $\alpha \approx 0.12$–0.13, suggesting a common serving-stack overhead factor of ∼8×. Phi-3-mini and Qwen3-4B's hybrid or non-standard attention mechanisms may interact differently with the two serving stacks.

D MoE Verbosity and Deployment Failure Analysis

The MoE model achieves 55.0% accuracy—competitive with dense 7B–9B models—but 0.0% S@SLO at the Interactive tier. The corrected results (with 2,048-token output limit) reveal the failure mode is fundamentally different from what the initial token-limited evaluation suggested. We examine three contributing factors:

Extreme chain-of-thought verbosity. The dominant failure mode is output length. The model generates approximately 1,500 tokens per response on average—extensive step-by-step reasoning, analysis, and explanation wrapping the required JSON output. At ∼50 tok/s, this produces 25,055 ms median latency. Dense 3B models generate ∼39 tokens per response at 44 tok/s, producing ∼884 ms median latency. The MoE generates 38× more output tokens per request. This is the primary SLO failure mechanism.

JSON validity is moderate, not catastrophic. The overall JSON validity rate is 66.8%—within the range of dense models (55.1%–85.3%). Per-task JSON validity varies: T4 at 95.2% and T6 at 99.0% show the model *can* produce valid structured output reliably for some tasks. T5 at 40.3% and T1 at 2.4% are poor, but T1 is a special case (the model generates thousands of tokens of

reasoning for a task requiring a single label). The accuracy (55.0%) is well below the JSON validity (66.8%), indicating that when it produces valid JSON, the content is sometimes incorrect—but this gap is comparable to dense models.

Task-dependent failure modes. T1 (intent classification, 2.2% accuracy) is the one catastrophic failure. The model generates massive chain-of-thought output for a task requiring a brief intent label, and the reasoning interferes with schema compliance. T6 (math, 71.5%) and T5 (code, 60.3%) show the opposite: the model's verbose reasoning *helps* accuracy on complex tasks. T4 (49.8%) and T3 (43.4%) are mid-range. The failure is not uniform—it is task-dependent, with verbosity helping complex reasoning tasks and hurting simple classification tasks.

These factors interact: the model generates correct, verbose output that arrives too late. For deployment, the key lever is output length management. Fine-tuning for concise structured output (GRPO with length penalty), constrained decoding, or prompting strategies that suppress chain-of-thought reasoning could potentially preserve the accuracy advantage while reducing latency to SLO-compliant levels.