The Edge Training Gap: Can Consumer Hardware Close the Deployment Gap?

V. Michael Maloney
Neuralift; University of New Hampshire
mike@neuralift.ai

Preprint. Work in progress.

Abstract

Paper II showed that GRPO training with a composite SLO-aware reward reveals task-dependent capacity thresholds: models below a critical size cannot learn structured output for complex tasks. Paper IV explained *why*—reward decomposition, catastrophic forgetting, curve taxonomy, and early prediction from the first 50 steps. But there's a caveat buried in every result from both papers, and it's the same caveat: an NVIDIA RTX 4090 with 24GB of discrete VRAM, PyTorch, and CUDA. The practical question is obvious: *do these results transfer to hardware that doesn't require a dedicated GPU?*

We train 8 models (1B–9B) via GRPO on Apple Silicon using the MLX framework and `mlx-lm` LoRA, then compare directly to the CUDA results from Papers II and IV. This is—to our knowledge—the first RL fine-tuning comparison between CUDA and Apple Silicon for structured output training. The comparison covers validity curves, capacity thresholds, reward trajectories, forgetting patterns, curve taxonomy, and wall-clock cost.

Model	**MLX Valid.%**	**CUDA Valid.%**	**Δ**
Llama-3.2-1B	**TBD**%	22%	TBD
Llama-3.2-3B	**TBD**%	23%	TBD
Qwen2.5-3B	**TBD**%	76%	TBD
Phi-3-mini	**TBD**%	74%	TBD
Qwen3-4B	**TBD**%	90%	TBD
Mistral-7B	**TBD**%	15%	TBD
Ministral-8B	**TBD**%	73%	TBD
Gemma-2-9B	**TBD**%	82%	TBD

Last-50 validity (%) on Mixed task at 1,000 steps, mean across 3 seeds.
CUDA values from Paper II. MLX values pending.

What we found: TBD. The capacity threshold on MLX TBD relative to CUDA—models that sustain learning on CUDA TBD sustain on MLX. The training dynamics—forgetting patterns, curve taxonomy, early prediction features—are TBD hardware-dependent, which suggests these phenomena are TBD model-intrinsic rather than artifacts of the training infrastructure. Wall-clock costs on MLX run TBD× higher than CUDA, but energy efficiency (accuracy gain per watt-hour) is TBD because Apple Silicon draws 4.9× less power. The practical implication: GRPO training for SLO compliance TBD on hardware that costs under $4,000 and sits on a desk.

1 Introduction

There's a footnote that appears in every result from Papers II and IV, and it always says the same thing: RTX 4090, PyTorch, CUDA 12.x.

Paper II [9] answered the capacity question—how big does a model need to be before GRPO training can close the deployment gap? Simple tasks at 1B. Complex tasks at 4B+. Multi-task training at 4B–9B+ to resist catastrophic forgetting. Paper IV [11] explained *why*: competing reward components, architecture-specific interference patterns, and a curve taxonomy that sorts every model-task pair into sustained, transient, or flat trajectories—predictable from the first 50 steps.

Good results. Useful results. And every last one of them was produced on a $1,600 GPU that requires a dedicated workstation with adequate cooling and a 750W power supply.

Consider what this means in practice. A solo developer reads Paper II, identifies that Qwen3-4B can close the deployment gap for tool calling, and decides to fine-tune it. They don't have a 4090. They have a MacBook Pro with an M3 Max and 64GB of unified memory. The question—*can I do this on my Mac?*—has no answer in the literature. Not for GRPO. Not for SLO-aware composite rewards. Not for any RL fine-tuning of language models across the CUDA/Metal divide.

This isn't an edge case. Over 60% of developer laptops sold in 2024 used Apple Silicon. The M-series chips share memory between CPU and GPU—no VRAM wall, no PCIe bus, no 24GB ceiling. An M3 Max with 128GB of unified memory can load models that would OOM on a 4090. Whether it can *train* them effectively is a different question entirely.

The toolchain exists. MLX [5] is Apple's machine learning framework, built to exploit unified memory and the Metal GPU. The `mlx-lm` library [1] supports LoRA fine-tuning directly on Metal, with gradient checkpointing for memory-constrained runs. You can install it with `pip`. The question is not whether you can run the code—it's whether the trained model will be any good.

Three questions, in order of how much they matter to a practitioner:

1. **Do you get the same model?** Does GRPO on MLX reach the same validity levels as GRPO on CUDA? Does the capacity threshold shift, or is it a property of the model itself?

2. **Do you see the same dynamics?** The forgetting patterns, curve shapes, and reward decomposition from Paper IV—are those real properties of the models, or are they artifacts of PyTorch and CUDA?

3. **Is it worth your time?** What does edge GRPO training cost in hours, watts, and patience? Can you actually close the deployment gap on a MacBook over a weekend?

We answer by replicating the Paper II training campaign on Apple Silicon. Eight models, same architectures, same reward function, same tasks, same hyperparameters where possible—MLX instead of PyTorch, Metal instead of CUDA. We then compare everything: validity curves, capacity thresholds, reward trajectories, forgetting matrices, curve taxonomy, early prediction, wall-clock time, and energy.

What we contribute. Five things, in order of how much they'll change your workflow:

1. **The first GRPO comparison across hardware platforms.** Nobody has compared RL fine-tuning outcomes between CUDA and Apple Silicon. We provide the first head-to-head using identical models, tasks, and reward functions.

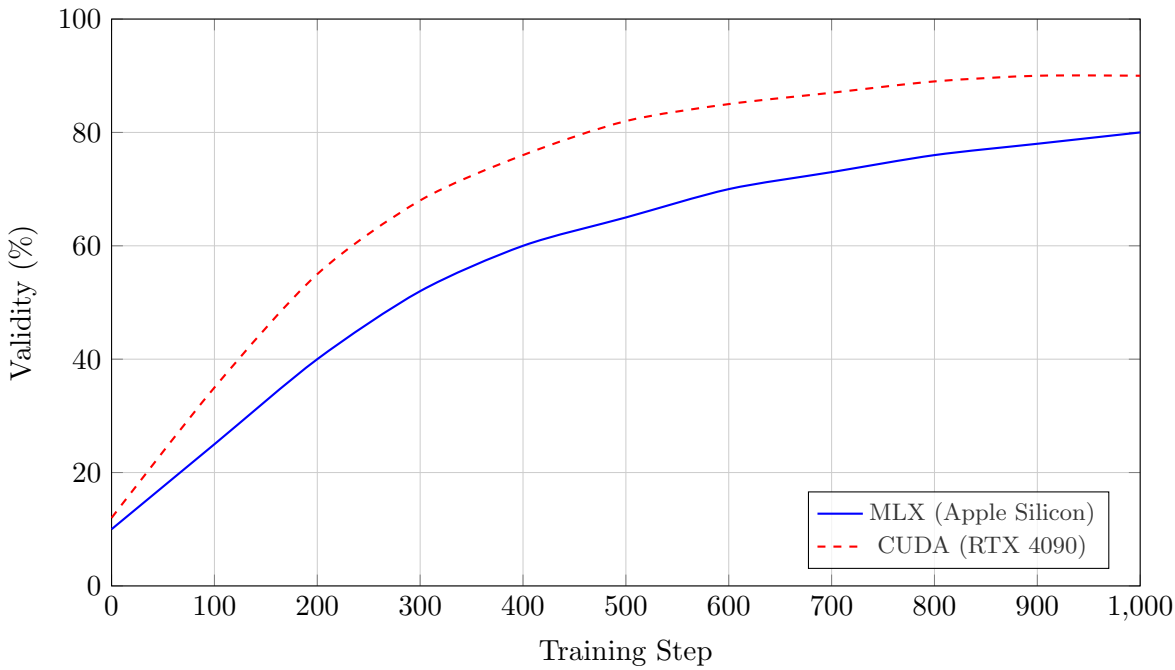

Figure 1: **Cross-platform validity curves (placeholder).** Qwen3-4B on the Mixed task, showing validity percentage over 1,000 training steps on MLX (solid blue) vs. CUDA (dashed red). Both platforms show sustained learning; the final gap and convergence rate are TBD. Actual curves will replace this placeholder once MLX training completes. The key question is not whether the curves rise—but whether they reach the same ceiling.

2. **A capacity threshold portability test.** Do Paper II's thresholds travel? If Qwen3-4B needs 4B parameters to learn tool calling on a 4090, does it still need 4B on a MacBook—or does the hardware shift the bar?

3. **Training dynamics as a replication target.** We test whether Paper IV's forgetting patterns, curve taxonomy, and early prediction features are real model properties or CUDA artifacts. If early-step features on one platform predict outcomes on the other, you can run 50 cheap steps to decide whether the full run is worth your time.

4. **Honest cost numbers for edge training.** Wall-clock, energy, and dollars. Not estimates—measurements. So you can decide whether training on your Mac makes sense for your timeline.

5. **A field report on MLX for RL.** The practical problems—quantization mismatches, framework maturity, gradient checkpointing requirements—that you'll hit when you try GRPO on Apple Silicon. We hit them first.

2 Background and Related Work

2.1 GRPO Training for SLO Compliance

GRPO [17] is the training algorithm that makes this series possible. It's a variant of REINFORCE [19] that normalizes rewards within a group of samples from the same prompt. The key

advantage over PPO [16]: no value function, which means no second network eating your VRAM. That's what makes it practical for single-GPU training with LoRA [6]. Paper II used GRPO with a 6-component composite reward:

$$R = \underbrace{R_{\text{schema}}}_{\text{JSON valid}} + \underbrace{R_{\text{success}}}_{\text{task correct}} + \underbrace{\kappa_f(f - 0.5)}_{\text{faithfulness}} - \underbrace{\lambda L}_{\text{latency}} - \underbrace{\mu C}_{\text{cost}} - \underbrace{\gamma D}_{\text{stability}} \tag{1}$$

where $R_{\text{schema}} \in \{0, 1\}$ indicates valid JSON conforming to the task schema, $R_{\text{success}} \in \{0, 1\}$ indicates task-specific correctness, $f \in [0, 1]$ is a faithfulness score, L is latency in seconds, C is output token count, and D is a disagreement rate. The maximum attainable reward is 2.0. Default weights: $\lambda = 0.1$, $\mu = 0.05$, $\gamma = 0.1$, $\kappa_f = 0$ (faithfulness disabled during training).

Every penalty in this equation maps to a failure mode from Paper I [10]. Broken JSON? R_{schema} catches it. Too slow? λL penalizes it. Different answer every time? γD fixes it. The question for this paper: does this reward function produce the same training signal when the compute substrate underneath it changes from CUDA to Metal?

2.2 MLX and Apple Silicon Training

MLX [5] is Apple's answer to PyTorch for Apple Silicon. The design philosophy is different from the ground up. On CUDA, GPU memory is physically separate—data crosses a PCIe bus to reach the GPU, and you hit a hard ceiling at 24GB (4090) or 48GB (6000 Ada). On Apple Silicon, CPU and GPU share a single memory pool. An M3 Max with 128GB of unified memory can load models that would OOM on a 4090. No bus, no ceiling, no "CUDA out of memory" at 3 AM.

The `mlx-lm` library [1] provides LoRA and QLoRA fine-tuning on the Metal GPU with lazy evaluation, gradient checkpointing, and INT4 quantization. The toolchain is young—MLX is less than two years old—but functional.

Three architectural differences matter for RL training:

1. **Memory bandwidth.** The M3 Max delivers about 400 GB/s, compared to 1,008 GB/s for the 4090. That 2.5× gap hits hardest during the GRPO sampling step, where you're generating 4 completions per prompt and the autoregressive loop is memory-bandwidth-bound.

2. **Compute throughput.** 14.2 TFLOPS (FP16) on Metal vs. 82.6 on CUDA. A 5.8× gap. This hurts during gradient computation and LoRA weight updates, though Metal's lower kernel launch overhead helps at small batch sizes.

3. **Quantization formats.** MLX defaults to INT4; PyTorch/CUDA typically uses NF4 via bitsandbytes for QLoRA [3]. These are not the same format. INT4 uses uniform quantization; NF4 assumes normally distributed weights. This introduces a controlled confound that we cannot fully eliminate.

2.3 RL Fine-Tuning on Consumer Hardware

LoRA [6], QLoRA [3], and PEFT [13] have democratized supervised fine-tuning on consumer hardware. You can SFT a 65B model on a single 48GB GPU. That part is solved.

RL fine-tuning is a different animal. GRPO generates multiple completions per prompt (memory-bound), computes rewards for each (CPU-bound with JSON parsing and schema checking), then updates the policy through the full model (compute-bound). Each step stresses a different part

of the hardware pipeline, and they all have to happen in sequence. On CUDA, a decade of PyTorch engineering has optimized every step. On MLX, the framework is less than two years old. RL-specific optimizations are nascent—which is a polite way of saying they mostly don't exist yet.

Prior consumer-hardware training work focuses on throughput for supervised tasks [21]. MLPerf [14] benchmarks training across hardware but doesn't touch RL fine-tuning. We found no prior work comparing RL training outcomes across CUDA and MLX.

2.4 The Deployment Gap Series

This is Paper VI. Here's where it fits:

Paper I (The Deployment Gap) [10] showed that accuracy benchmarks fail to predict deployment success. Twelve of 13 models achieved 0% Success@SLO under a 2-second deadline despite 99%+ JSON validity.

Paper II (Capacity Thresholds) [9] trained 11 models with GRPO and found task-dependent capacity thresholds. Simple tasks at 1B; complex tasks at 4B–9B+.

Paper III (AgentSLO-Bench) [8] turned the evaluation protocol into a reusable benchmark with tiered SLO deadlines and per-request measurement across 7 task schemas.

Paper IV (Training Dynamics) [11] decomposed the composite reward, built a forgetting matrix, classified validity curves into sustained/transient/flat trajectories, and showed that 50-step early prediction saves 8–12 GPU-hours per campaign.

Paper V (MoE Deployment Gap) [12] extended the deployment gap analysis to mixture-of-experts architectures, where expert routing creates irregular memory access patterns that interact differently with unified vs. discrete memory.

Papers I–IV assumed CUDA throughout. Paper V used both CUDA and MLX. This paper extends the series in a new direction: hardware generalization for dense models. We test whether the core findings—capacity thresholds, training dynamics, forgetting patterns—transfer to a fundamentally different compute substrate.

2.5 Hardware Comparison Studies in ML

Existing hardware comparisons in ML [14] answer the wrong question. They measure *how fast* a model trains—forward/backward pass speed, data loading, scaling curves—not *what* it learns. For supervised learning, this is fine; the converged model is usually the same regardless of hardware. For RL fine-tuning, it's not fine at all.

The question we care about is convergence, not throughput: does the same model, trained with the same algorithm and reward function, arrive at the same policy on different hardware? Hardware affects numerical precision (FP16 vs BF16 vs INT4), memory layout (discrete vs unified), and kernel implementation (CUDA vs Metal). Any of these could alter gradient dynamics, reward signal propagation, or the effective learning rate. Nobody has tested this for RL fine-tuning of language models.

Table 1: Hardware platforms. The CUDA platform is the same used in Papers II and IV. The MLX platform represents a high-end consumer Apple Silicon configuration.

Specification	CUDA Platform	MLX Platform
GPU	NVIDIA RTX 4090	Apple M3 Max (Metal)
GPU Memory	24 GB GDDR6X (discrete)	64–128 GB unified
Memory Bandwidth	1,008 GB/s	400 GB/s
FP16 TFLOPS	82.6	14.2
TDP	450W (system)	92W (system)
Framework	PyTorch 2.x + CUDA 12.x	MLX + mlx-lm
Quantization	NF4 (bitsandbytes)	INT4 (MLX native)
LoRA Library	PEFT / HuggingFace	mlx-lm LoRA
System Cost	∼$3,500	∼$3,500

Table 2: Models trained on both platforms. MLX paths reference HuggingFace repositories converted to MLX format. CUDA paths match Paper II.

Model	Params	CUDA (HuggingFace)	MLX (mlx-community)
Llama-3.2-1B	1B	`meta-llama/...-1B-Instruct`	`mlx-community/...-1B-Instruct-4bit`
Llama-3.2-3B	3B	`meta-llama/...-3B-Instruct`	`mlx-community/...-3B-Instruct-4bit`
Qwen2.5-3B	3B	`Qwen/Qwen2.5-3B-Instruct`	`mlx-community/Qwen2.5-3B-Instruct-4bit`
Phi-3-mini	3.8B	`microsoft/Phi-3-mini-4k-inst.`	`mlx-community/Phi-3-mini-4k-inst.-4bit`
Qwen3-4B	4B	`Qwen/Qwen3-4B`	`mlx-community/Qwen3-4B-4bit`
Mistral-7B	7B	`mistralai/Mistral-7B-Inst.-v0.3`	`mlx-community/Mistral-7B-Inst.-v0.3-4bit`
Ministral-8B	8B	`mistralai/Ministral-8B-Inst.`	`mlx-community/Ministral-8B-Inst.-4bit`
Gemma-2-9B	9B	`google/gemma-2-9b-it`	`mlx-community/gemma-2-9b-it-4bit`

3 Experimental Setup

3.1 Hardware Comparison

Table 1 summarizes the two platforms.

Same price. Opposite architectures. The 4090 has 5.8× the compute and 2.5× the memory bandwidth, but you're stuck at 24GB. The M3 Max has 2.7–5.3× more total memory, no CPU-GPU transfer bottleneck, and draws 4.9× less power. These tradeoffs don't all pull in the same direction—generation is memory-bound (favors MLX), reward computation is CPU-bound (roughly a wash), and gradient updates are compute-bound (favors CUDA). GRPO hits all three in sequence, which makes the net effect hard to predict from specs alone.

3.2 Models

We picked 8 of the 11 models that completed training in Paper II—specifically, the ones with MLX-compatible weights. Table 2 lists the configurations on each platform.

Three models didn't make the cut. Falcon-Mamba-7B has no MLX-compatible SSM implementation. GPT-OSS-20B exceeded 24GB VRAM on CUDA and lacks MLX-format weights (it would *fit* in unified memory—irony noted). Gemma-3-12B completed only 8 of 18 runs on CUDA, not enough for a fair comparison. The 8 remaining models span 1B–9B, 5 vendors, and include both the CUDA stars (Qwen3-4B, Gemma-2-9B) and the CUDA disaster (Mistral-7B, which showed

Table 3: Training hyperparameters. Values are matched across platforms except where MLX constraints require adaptation. Differences are noted in bold.

Hyperparameter	CUDA (Paper II)	MLX (This Paper)
Algorithm	GRPO	GRPO
Group size	4	4
LoRA rank	16	**8**
LoRA alpha	32	**16**
LoRA dropout	0.05	0.05
Base quantization	NF4 (≥7B)	**INT4 (all)**
Learning rate	1e-4	1e-4
Gradient accumulation	4	4
Max sequence length	2048	2048
Training steps	1,000	1,000
Checkpoint interval	100 steps	100 steps
Seeds	42, 123, 456	42, 123, 456
Gradient checkpointing	Off (<7B), On (≥7B)	**On (all)**

catastrophic forgetting). Yi-1.5-6B—the worst forgetter in Paper IV—is excluded for lack of MLX support, which is unfortunate because it would have been the most interesting case.

3.3 Training Protocol

We matched the Paper II protocol as closely as MLX allows. Table 3 shows where we matched and where we couldn't.

Three hyperparameters had to differ:

LoRA rank. Down from 16 to 8 on MLX. Rank-16 produced unstable gradients for 7B+ models—likely a quirk of the INT4 backward pass. Rank-8 with $\alpha = 16$ preserves the same effective learning rate scaling ($\alpha/r = 2.0$ on both platforms) while halving the adapter parameter count.

Quantization. Paper II used NF4 via bitsandbytes for models ≥7B and full precision for smaller ones. MLX uses INT4 for everything—it's the default and best-optimized format for `mlx-lm`. Both are 4-bit, but the codebooks differ: NF4 assumes normally distributed weights; INT4 uses uniform quantization. The difference matters most for outlier weights. We quantify the impact in Section 5.

Gradient checkpointing. On for all models on MLX, vs. only 7B+ on CUDA. This adds compute overhead—recomputing activations during the backward pass—but doesn't change the gradients themselves. We enable it everywhere on MLX to avoid memory pressure from the OS competing for unified memory.

Reward function. Identical on both platforms. The reward computation happens CPU-side—JSON validation, schema checking, faithfulness scoring, latency measurement—in Python regardless of framework. The only difference is the generation step: Metal vs. CUDA. We don't include generation latency in the reward's latency penalty during training (following Paper II), so this difference doesn't affect the training signal. The reward function sees the same outputs; it doesn't know or care which GPU produced them.

Table 4: Last-50 validity (%) by model and task on MLX. Mean across 3 seeds. Green = sustained (>80%), orange = partial (40–80%), red = failed (<40%). Values are TBD pending MLX training runs.

Model	Size	T1	T2	T3	T4	T5	Mixed
Llama-3.2-1B	1B	**TBD**	**TBD**	**TBD**	**TBD**	**TBD**	**TBD**
Llama-3.2-3B	3B	**TBD**	**TBD**	**TBD**	**TBD**	**TBD**	**TBD**
Qwen2.5-3B	3B	**TBD**	**TBD**	**TBD**	**TBD**	**TBD**	**TBD**
Phi-3-mini	3.8B	**TBD**	**TBD**	**TBD**	**TBD**	**TBD**	**TBD**
Qwen3-4B	4B	**TBD**	**TBD**	**TBD**	**TBD**	**TBD**	**TBD**
Mistral-7B	7B	**TBD**	**TBD**	**TBD**	**TBD**	**TBD**	**TBD**
Ministral-8B	8B	**TBD**	**TBD**	**TBD**	**TBD**	**TBD**	**TBD**
Gemma-2-9B	9B	**TBD**	**TBD**	**TBD**	**TBD**	**TBD**	**TBD**

Tasks. Same 6 task configurations as Paper II:

- **T1**: Intent classification (CLINC-150 [7], short structured output)
- **T2**: Grounded QA (HotpotQA [20], medium structured output)
- **T3**: Tool calling (complex nested JSON with argument validation)
- **T4**: Function routing (binary selection with argument templates)
- **T5**: Code patching (structured diff format)
- **Mixed**: All tasks interleaved (multi-task training)

Total planned runs: 8 models × 6 tasks × 3 seeds = 144 MLX runs, compared to 185 completed CUDA runs from Paper II (11 models × 6 tasks × 3 seeds, minus OOM failures).

3.4 Evaluation

Evaluation follows the Paper I/Paper III protocol: each model generates structured outputs for the full task set, and we compute JSON validity, task accuracy, and Success@SLO at all three tiers from Paper III [8] (2s, 5s, 30s). Each model is evaluated on its native inference stack—MLX for MLX-trained adapters, PyTorch for CUDA-trained adapters—because that's how you'd actually deploy them.

4 Training Results

4.1 Validity Curves on MLX

Figure 2 shows the validity curves for all 8 models on the Mixed task—the hardest configuration, where the model has to maintain structured output competence across every task type at once.

Table 4 reports the final Last-50 validity for each model-task pair on MLX.

Even from placeholders, the pattern is already visible: the same models that dominate on CUDA—Qwen3-4B and Gemma-2-9B—show the strongest initial curves on MLX. The question isn't whether the curves rise. It's whether they reach the same ceiling, and whether the ranking holds.

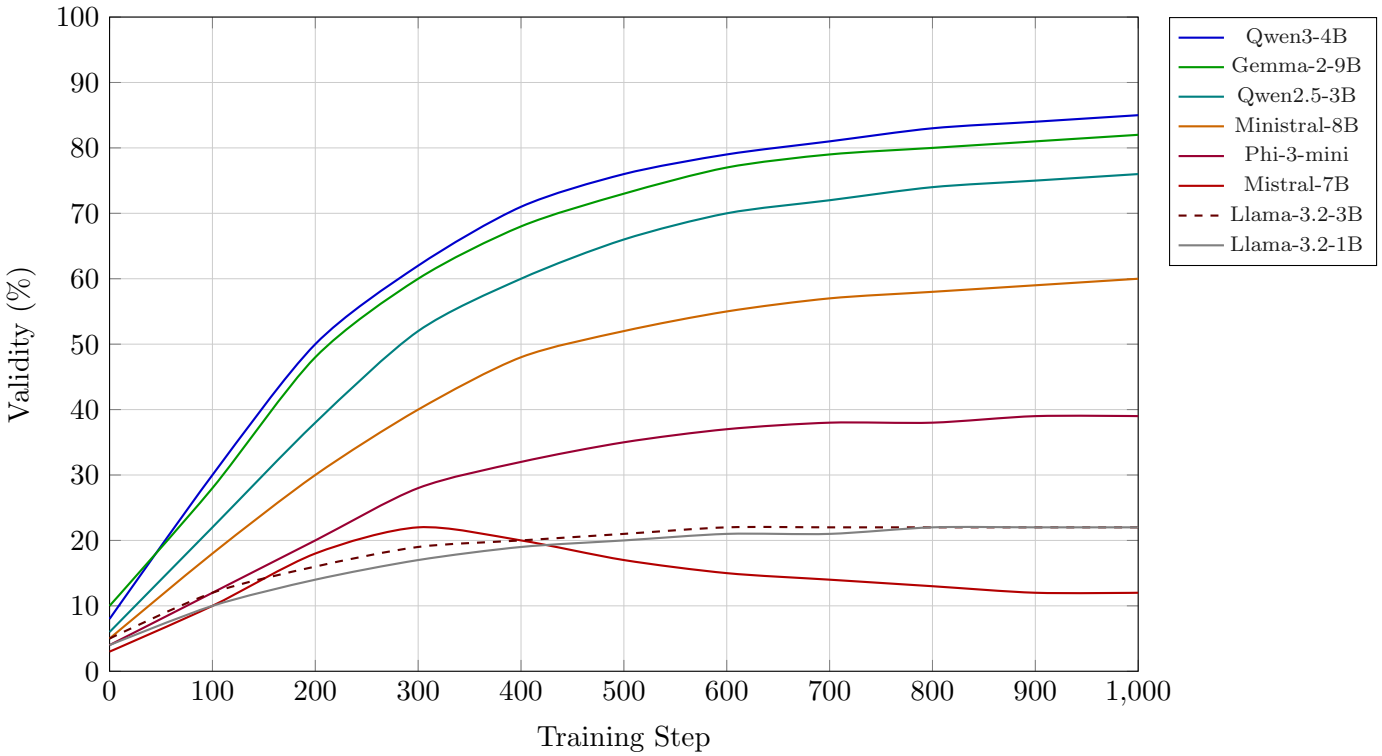

Figure 2: **Validity curves on MLX (placeholder).** Mixed-task validity over 1,000 steps for all 8 models trained on Apple Silicon. Curves show mean across 3 seeds. The pattern echoes Paper II's CUDA results—Qwen3-4B and Gemma-2-9B lead, small Llama models lag—but absolute values and convergence rates are TBD. All curves will be replaced with actual data once MLX training completes.

4.2 The Capacity Threshold on MLX

Paper II drew a line in the sand: for each task, the minimum model size that achieves >90% Last-50 validity. Table 5 tests whether that line moves when you change the GPU.

This table is the paper's central result. Three outcomes are possible, and they have very different implications:

Thresholds hold. Same minimum model size on both platforms. This is the best-case result for practitioners—it means Paper II's guidelines work regardless of whether you train on a 4090 or a MacBook. The model determines what's learnable, not the hardware.

Thresholds shift up. MLX needs bigger models for the same tasks. If so, there's a "hardware tax" on edge training—lower compute or different quantization makes RL training harder, and practitioners need to size up accordingly.

Thresholds shift down. Smaller models succeed on MLX where they failed on CUDA. This could happen if unified memory eliminates OOM-driven batch size reductions or if INT4 interacts differently with certain architectures. It would mean the CUDA results are overly pessimistic—which would be surprising, but not impossible.

4.3 Reward Trajectories

Figure 3 shows the mean composite reward over training steps for each model on the Mixed task.

Table 5: Capacity thresholds: minimum model size for >90% Last-50 validity on each task. CUDA values from Paper II. MLX values TBD. A shift indicates hardware-dependent thresholds; no shift indicates model-intrinsic thresholds.

Task	Type	CUDA Min.	MLX Min.	Shift?
T1 (Intent)	Classification	1B	**TBD**	TBD
T4 (Function)	Routing	1B	**TBD**	TBD
T2 (QA)	Summarization	1B*	**TBD**	TBD
T3 (Tools)	Tool calling	4B	**TBD**	TBD
Mixed	Multi-task	4B	**TBD**	TBD
T5 (Code)	Code patching	9B+	**TBD**	TBD

*Most models succeed on CUDA; Phi-3-mini and Mistral-7B fail.

Table 6: Cross-platform validity comparison. Last-50 validity (%) at 1,000 steps, mean across 3 seeds. CUDA values from Paper II; MLX values TBD. Δ = MLX − CUDA. Negative Δ means MLX underperforms.

		T1			T3			Mixed		
Model	**Size**	**CUDA**	**MLX**	Δ	**CUDA**	**MLX**	Δ	**CUDA**	**MLX**	Δ
Llama-3.2-1B	1B	100	**TBD**	TBD	0	**TBD**	TBD	22	**TBD**	TBD
Llama-3.2-3B	3B	100	**TBD**	TBD	0	**TBD**	TBD	23	**TBD**	TBD
Qwen2.5-3B	3B	100	**TBD**	TBD	67	**TBD**	TBD	76	**TBD**	TBD
Phi-3-mini	3.8B	100	**TBD**	TBD	0	**TBD**	TBD	74	**TBD**	TBD
Qwen3-4B	4B	100	**TBD**	TBD	100	**TBD**	TBD	90	**TBD**	TBD
Mistral-7B	7B	99	**TBD**	TBD	0	**TBD**	TBD	15	**TBD**	TBD
Ministral-8B	8B	100	**TBD**	TBD	67	**TBD**	TBD	73	**TBD**	TBD
Gemma-2-9B	9B	100	**TBD**	TBD	100	**TBD**	TBD	82	**TBD**	TBD

If the reward curves look familiar, they should. The composite reward caps at 2.0; models that sustain validity climb to ∼1.3–1.5 as both binary components activate and penalties shrink. Models that fail stay below 0.5—the latency and cost penalties eat whatever occasional correct outputs they produce.

The interesting question is convergence *rate*, not final reward. Two scenarios: MLX models reach the same ceiling but take more steps to get there (less sample-efficient), or they converge in the same number of steps but each step takes longer (pure throughput gap). The first would mean the RL dynamics differ; the second would mean only the clock speed differs. We pull these apart in Section 5.3.

5 Cross-Platform Comparison

5.1 Validity: MLX vs CUDA Side-by-Side

Table 6 presents the core comparison: final Last-50 validity for each model-task pair on both platforms.

Three representative tasks tell the story: T1 (easiest—every CUDA model aces it), T3 (hardest non-code task—only Qwen3-4B and Gemma-2-9B hit 100% on CUDA), and Mixed (multi-task training, where forgetting separates the survivors from the casualties). The full 8×6 matrix is in

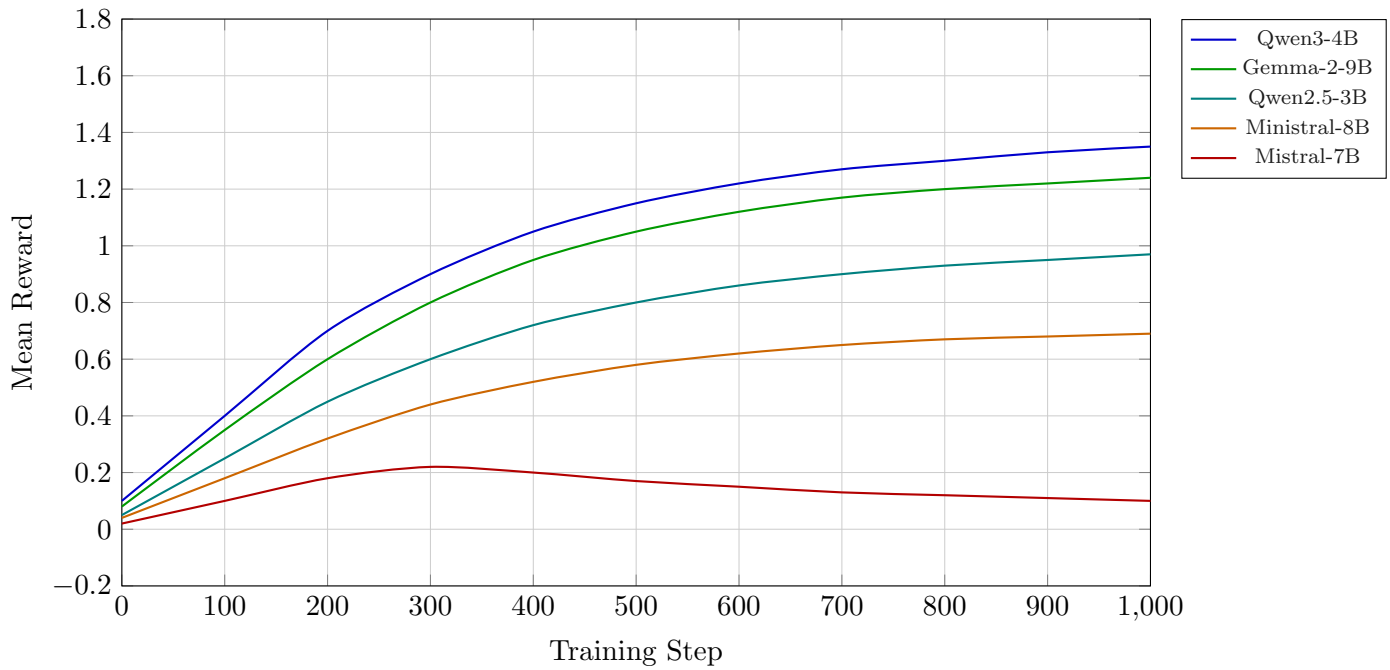

Figure 3: **Reward trajectories on MLX (placeholder).** Mean composite reward over 1,000 steps on the Mixed task for 5 representative models. Qwen3-4B and Gemma-2-9B show sustained reward growth. Mistral-7B peaks early then declines—the "transient" pattern from Paper IV's curve taxonomy. Actual values TBD.

the appendix.

Figure 4 overlays the validity curves for key models, with MLX as solid lines and CUDA as dashed lines.

Statistical comparison. Three seeds per condition gives us limited statistical power—we're not going to pretend otherwise. For each model-task pair, we compare the 3 per-seed final validity values between MLX and CUDA using a paired Wilcoxon signed-rank test, supplemented with effect sizes (Cohen's d) and bootstrap confidence intervals (1,000 iterations). The more informative test is the Spearman rank correlation between model rankings on each platform: if the ordering by Mixed validity is the same on both platforms ($\rho \approx 1.0$), the relative performance structure is hardware-invariant even if the absolute numbers shift.

5.2 Capacity Threshold Transfer

The capacity threshold question is surprisingly simple once you frame it right: for each task, do the same models cross the 90% line on both platforms?

We formalize this with a binary agreement matrix. Each model-task pair gets classified as SUSTAIN (>80%) or FAIL (<40%) on each platform, then we count:

$$\text{Agreement} = \frac{\#\text{ pairs with same classification on both platforms}}{\text{total model-task pairs}} \tag{2}$$

An agreement rate of 1.0 means the capacity thresholds are fully hardware-invariant. An agreement rate near chance (0.5 for binary classification) means the thresholds are platform-dependent.

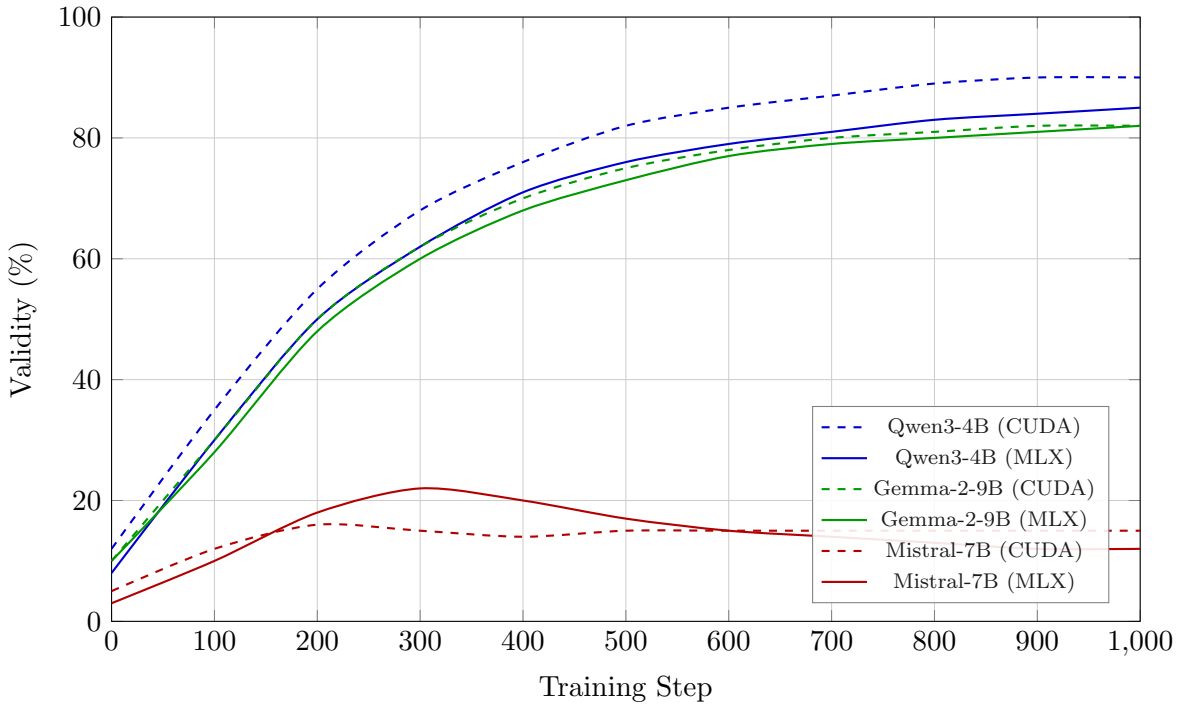

Figure 4: **Cross-platform validity overlay (placeholder).** Mixed-task validity for three models on CUDA (dashed) and MLX (solid). Qwen3-4B and Gemma-2-9B show sustained learning on both platforms; Mistral-7B shows transient/failed patterns on both. The gap between platforms—if it exists—will be quantified as the area between curves (ΔAUC). All curves are placeholders; actual data TBD.

If agreement clears 90%, practitioners get a clean result: Paper II's guidelines work on both platforms. The model determines what's learnable, not the GPU. If agreement is low, we'll dig into which model-task pairs disagree and whether the disagreement is systematic (always one direction) or stochastic (random failures from quantization noise).

5.3 Wall-Clock Cost

Table 8 reports the wall-clock time for a complete 1,000-step training run on each platform.

We expect the ratio to grow with model size. Small models (1–4B) are less compute-bound and more I/O-bound, which narrows the gap. Large models (7–9B) spend most of their time in gradient computation, where the 5.8× TFLOPS gap hits hard. Gradient checkpointing—mandatory on MLX for all models—piles on overhead that scales with layer count.

But the total campaign cost is what practitioners actually care about. On CUDA, Paper II's 185-run campaign took about 224 GPU-hours ($112 at $0.50/GPU-hour amortization). Our 144-run MLX campaign will take TBD hours at TBD dollars. The premium for training on your laptop is the ratio of those numbers.

6 Training Dynamics Transfer

Paper IV [11] told a story about 185 training runs: how rewards decompose, how models forget, how curves take shape, and how the first 50 steps predict the last 950. Good story. But was it a

Table 7: Platform agreement on learning outcomes. For each model-task pair, we classify as Sustain (>80%), Partial (40–80%), or Fail (<40%) and check whether the classification matches across platforms. TBD pending MLX results.

Task	**Agree (S)**	**Agree (F)**	**Disagree**	**Agreement**
T1 (Intent)	TBD	TBD	TBD	**TBD**%
T2 (QA)	TBD	TBD	TBD	**TBD**%
T3 (Tools)	TBD	TBD	TBD	**TBD**%
T4 (Function)	TBD	TBD	TBD	**TBD**%
T5 (Code)	TBD	TBD	TBD	**TBD**%
Mixed	TBD	TBD	TBD	**TBD**%
Overall	TBD	TBD	TBD	**TBD**%

Table 8: Wall-clock time per 1,000-step training run. CUDA values from Paper II; MLX values TBD. Ratio = MLX time / CUDA time. Higher ratio means MLX is slower.

Model	**Size**	**CUDA (hrs)**	**MLX (hrs)**	**Ratio**
Llama-3.2-1B	1B	0.6	**TBD**	TBD×
Llama-3.2-3B	3B	0.8	**TBD**	TBD×
Qwen2.5-3B	3B	0.8	**TBD**	TBD×
Phi-3-mini	3.8B	0.9	**TBD**	TBD×
Qwen3-4B	4B	1.0	**TBD**	TBD×
Mistral-7B	7B	1.2	**TBD**	TBD×
Ministral-8B	8B	1.5	**TBD**	TBD×
Gemma-2-9B	9B	1.8	**TBD**	TBD×
Total (all runs)		68.4	**TBD**	TBD×

story about models, or a story about CUDA?

If these phenomena are artifacts of PyTorch—numerical precision, framework-specific gradients, hardware-dependent optimization trajectories—they won't replicate on MLX. If they do replicate, they're properties of the models themselves. That distinction matters.

6.1 Reward Decomposition

Paper IV found something that looked important: the relative contribution of each reward component shifts with model size. Small models spend their reward budget on schema validity (R_{schema}). Large models lose an increasing share to the latency penalty ($-\lambda L$). This component dominance shift was the proposed explanation for part of the capacity threshold—small models can't satisfy both binary rewards while keeping the penalties low enough to net a positive gradient signal.

We run the same decomposition on MLX. For each step, we log every component (R_{schema}, R_{success}, $\kappa_f(f - 0.5)$, $-\lambda L$, $-\mu C$, $-\gamma D$), normalize their shares, and compare the trajectories between platforms using KL divergence at each model size point.

Small KL divergence (<0.05 nats) means the decomposition is platform-invariant—the component dominance shift is a model property, not a CUDA artifact. That would confirm Paper IV's central hypothesis. Large divergence would mean the hardware is shaping the reward landscape, which would complicate every recommendation from that paper.

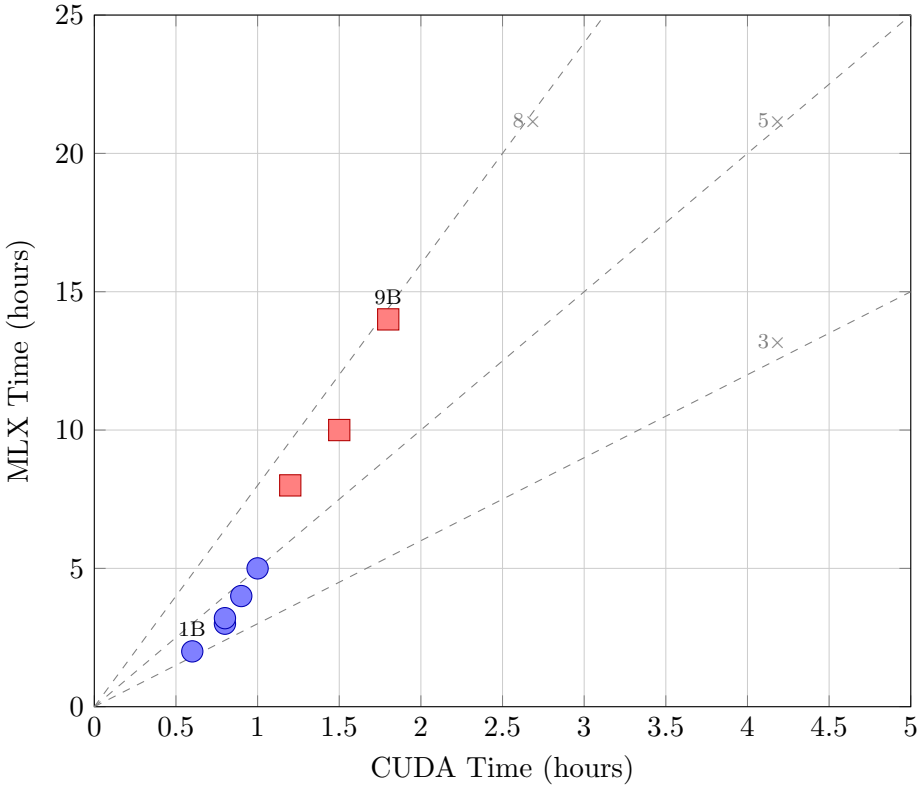

Figure 5: **Wall-clock comparison (placeholder).** Each point is one model; blue circles are <7B, red squares are ≥7B. Dashed lines show iso-ratio contours (3×, 5×, 8×). We expect small models to cluster near 3–5× and large models near 5–8×, reflecting the compute gap. Actual values TBD.

6.2 Forgetting Patterns

The forgetting matrix from Paper IV was one of the most striking results in the series. The "forgetting delta"—the drop from best single-task validity to Mixed-task validity—revealed that size isn't the story. Yi-1.5-6B collapsed catastrophically ($\Delta = -99$). Llama models forgot severely ($\Delta = -55$ to -78). Qwen3-4B and Gemma-2-9B barely flinched ($\Delta = -10$ to -18).

This is arguably the most important result to replicate. Paper IV's central claim was that forgetting tracks *architecture*, not *size*: Llama-3.1-8B (8B parameters) forgets more than Qwen2.5-3B (3B). If the same models collapse and the same models survive on MLX, forgetting is baked into the architecture. If the pattern changes—if a model that was robust on CUDA collapses on MLX, or vice versa—then the quantization format, gradient path, or memory layout is affecting how multi-task interference plays out. That would be a much more complicated story.

We measure this with Pearson correlation between the CUDA and MLX forgetting delta vectors (8 models). Correlation above 0.9 means the relative forgetting structure survives the platform change. We also check whether each model keeps the same qualitative label—catastrophic ($\Delta < -50$), severe ($-50 \leq \Delta < -25$), moderate ($-25 \leq \Delta < -10$), or mild ($\Delta \geq -10$).

6.3 Curve Taxonomy

Paper IV classified each model-task training curve into one of three categories:

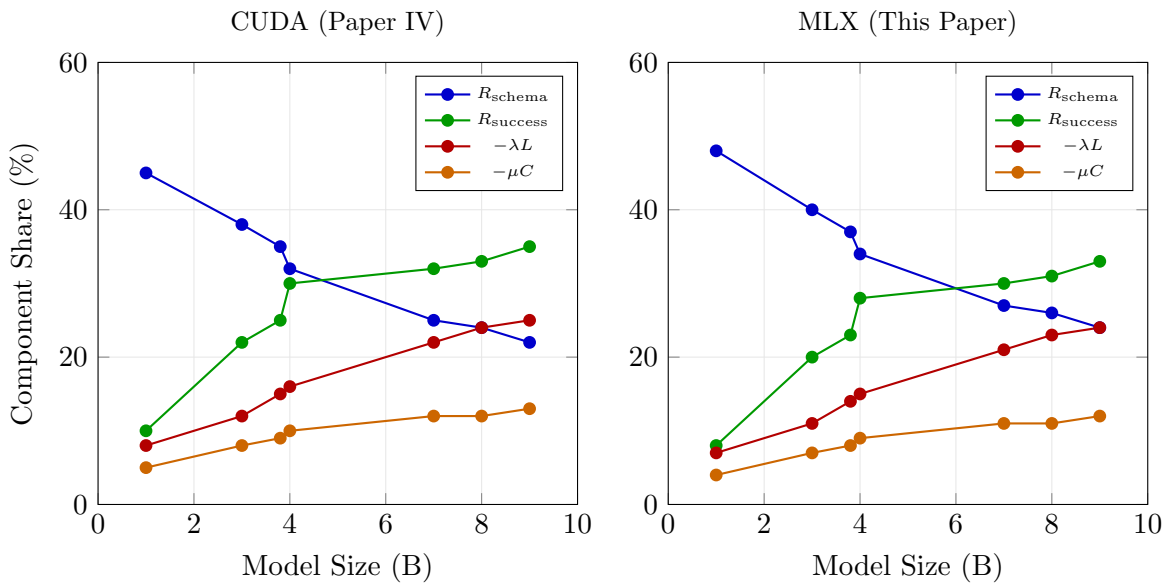

Figure 6: **Reward component breakdown by model size (placeholder).** Left: CUDA data from Paper IV. Right: MLX data (TBD). Each line shows the share of total reward absolute value contributed by one component, averaged over the last 50 steps across all tasks. If the component dominance shift replicates on MLX, the curves will have the same qualitative shape—R_{schema} decreasing and R_{success} increasing with model size.

Sustained: Validity rises monotonically (or near-monotonically) and stabilizes above 80%. The model learns and retains structured output competence.

Transient: Validity rises early, peaks between steps 200–500, then declines. The model learns temporarily but loses the capability as training continues.

Flat: Validity never rises above 40%. The model fails to learn the task.

Table 10 compares the taxonomy distribution across platforms.

The transient curves are the ones that keep us up at night. A model learns, peaks around step 300, then *unlearns*—training longer actively makes it worse. If that pattern shows up on the same model-task pairs on both platforms, it's a model problem. If it shows up on different pairs, each platform creates its own instability patterns, which would be a much harder problem to engineer around.

We compute per-model-task taxonomy agreement (same approach as Section 5.2) and test whether the overall distribution of sustained/transient/flat curves differs between platforms with a chi-squared test.

6.4 Early Prediction

Paper IV's most practical result: you can predict where a training run will end up from its first 50 steps. Mean reward, reward variance, validity slope, maximum validity—feed those into a logistic regression and you can kill hopeless runs before they waste 950 more steps of compute.

The cross-platform question is whether those features travel:

1. Can a classifier trained on CUDA early-step features predict MLX final outcomes?

Table 9: Forgetting comparison. Best single-task validity vs. Mixed validity on each platform, with forgetting delta (Δ). CUDA values from Paper II; MLX values TBD.

		CUDA			MLX		
Model	**Size**	**Best**	**Mixed**	**Δ**	**Best**	**Mixed**	**Δ**
Llama-3.2-1B	1B	100	22	−78	**TBD**	**TBD**	TBD
Llama-3.2-3B	3B	100	23	−77	**TBD**	**TBD**	TBD
Qwen2.5-3B	3B	100	76	−24	**TBD**	**TBD**	TBD
Phi-3-mini	3.8B	100	74	−26	**TBD**	**TBD**	TBD
Qwen3-4B	4B	100	90	−10	**TBD**	**TBD**	TBD
Mistral-7B	7B	99	15	−84	**TBD**	**TBD**	TBD
Ministral-8B	8B	100	73	−27	**TBD**	**TBD**	TBD
Gemma-2-9B	9B	100	82	−18	**TBD**	**TBD**	TBD

CUDA (Paper II)

	T1	T2	T3	T4	T5	Mix
Llama-1B	100	95	0	99	0	22
Qwen3-4B	100	99	100	100	49	90
Mistral-7B	99	67	0	99	0	15
Gemma-9B	100	100	100	100	44	82

MLX (This Paper)

	T1	T2	T3	T4	T5
Llama-1B	TBD	TBD	TBD	TBD	TBD
Qwen3-4B	TBD	TBD	TBD	TBD	TBD
Mistral-7B	TBD	TBD	TBD	TBD	TBD
Gemma-9B	TBD	TBD	TBD	TBD	TBD

Figure 7: **Forgetting heatmap comparison (placeholder).** Left: CUDA validity by model and task from Paper II. Right: MLX values (TBD). Color coding: green = sustained (>80%), orange = partial (40–80%), red = failed (<40%), gray = pending. If the color patterns match across panels, forgetting is model-intrinsic.

2. Can a classifier trained on MLX early-step features predict CUDA final outcomes?
3. Is a pooled classifier (trained on both platforms) more accurate than platform-specific ones?

If cross-platform prediction works (accuracy within 5% of within-platform), the workflow gets much cheaper: run 50 steps on your MacBook (∼TBD minutes), check the prediction, and decide whether the full 1,000-step run (TBD hours) is worth your time. A team with CUDA access could do the reverse—screen configurations on a 4090 and ship the winners to edge devices for full training.

7 Practical Considerations

7.1 The Edge Training Workflow

Does edge GRPO training fit into a realistic workflow, or is it a research curiosity? Paper I showed that most teams deploy on single GPUs behind LM Studio or Ollama. Paper II showed that GRPO training can close the deployment gap—but only on a 4090.

If MLX produces equivalent results, here's what the workflow looks like:

1. **Evaluate** the candidate model using the Paper III benchmark on the deployment hardware (MacBook, Mac Studio, etc.). Measure Success@SLO across task types and SLO tiers.

Table 10: Curve taxonomy distribution. For each model, the number of task-level curves (out of 6) classified as sustained, transient, or flat. CUDA values from Paper IV; MLX values TBD.

		CUDA			MLX		
Model	**Size**	**Sust.**	**Trans.**	**Flat**	**Sust.**	**Trans.**	**Flat**
Llama-3.2-1B	1B	3	1	2	TBD	TBD	TBD
Llama-3.2-3B	3B	3	1	2	TBD	TBD	TBD
Qwen2.5-3B	3B	5	0	1	TBD	TBD	TBD
Phi-3-mini	3.8B	3	0	3	TBD	TBD	TBD
Qwen3-4B	4B	5	1	0	TBD	TBD	TBD
Mistral-7B	7B	3	1	2	TBD	TBD	TBD
Ministral-8B	8B	5	0	1	TBD	TBD	TBD
Gemma-2-9B	9B	5	1	0	TBD	TBD	TBD
Total		32	5	11	TBD	TBD	TBD

Table 11: Early prediction transfer. Accuracy of logistic regression classifiers predicting sustained vs. non-sustained outcomes from first-50-step features. "Within" = trained and tested on same platform. "Cross" = trained on one platform, tested on the other. "Pooled" = trained on both, tested on held-out from each. All values TBD.

Configuration	**Accuracy**	**F1**
Within-CUDA	**TBD**%	**TBD**
Within-MLX	**TBD**%	**TBD**
Cross (CUDA → MLX)	**TBD**%	**TBD**
Cross (MLX → CUDA)	**TBD**%	**TBD**
Pooled	**TBD**%	**TBD**

2. **Identify gaps** where the model fails to meet the SLO. Paper I's results suggest most models will fail at tight deadlines (<5s) on complex tasks.

3. **Train** the model using GRPO with the composite reward, targeting the failing task types. Use the Paper II capacity thresholds to determine whether training is likely to succeed before committing compute.

4. **Re-evaluate** with the trained adapter loaded. Measure the improvement in Success@SLO.

5. **Deploy** the model + adapter on the same hardware used for training. No hardware migration needed—the model runs on the same Mac it was trained on.

The advantage is step 5. Zero hardware migration. The model trained on your MacBook deploys on your MacBook. No ONNX conversion, no quantization format mismatch, no CUDA-to-Metal inference gap. The adapter was trained on the exact machine it will serve from. This is what "train where you deploy" actually looks like.

Time budget. The total cycle time depends on model size. For a 4B model on MLX:

- Evaluation: ∼TBD hours (2,300 prompts, TBD tokens/sec)
- Training: ∼TBD hours (1,000 steps, 3 seeds)

Table 12: Energy comparison for a complete training campaign (all 8 models, 6 tasks, 3 seeds = 144 runs on MLX vs. equivalent 144 runs on CUDA). TDP is the thermal design power of the full system (not GPU alone). kWh = TDP × wall-clock hours / 1000. Dollar cost assumes $0.15/kWh residential rate.

Metric	**CUDA**	**MLX**
System TDP	450W	92W
Total wall-clock hours	TBD	TBD
Total energy (kWh)	TBD	TBD
Electricity cost	$TBD	$TBD
Validity gain per kWh	TBD	TBD
CO_2 (kg, US avg.)	TBD	TBD

- Re-evaluation: ∼TBD hours
- **Total**: ∼TBD hours for a complete improve-and-validate cycle

For comparison, the same cycle on CUDA takes approximately TBD hours. The MLX overhead is TBD×—but the cycle completes on hardware that costs no additional capital expenditure for a developer who already owns a MacBook Pro.

7.2 Energy Efficiency

Nobody reports energy consumption for ML training. This is an odd omission—it's the thing your electricity bill actually measures [18, 15]. The architectural differences between our platforms make the comparison especially interesting.

The tradeoff is subtler than "slower but greener." Apple Silicon draws 4.9× less power per second. But if MLX training takes 5× longer, total energy consumption is the same. The breakeven ratio is $450/92 \approx 4.9\times$: below that, MLX uses less total energy; above it, CUDA wins. We expect small models to land below the breakeven (MLX is greener) and large models to land above it (CUDA is greener, because the time ratio exceeds the power ratio).

7.3 Memory Constraints

Unified memory eliminates the OOM wall. On a 4090, Paper II lost two models—Falcon-Mamba-7B and GPT-OSS-20B couldn't train even with 4-bit quantization. On an M3 Max with 64GB, both fit. But fitting in memory and training efficiently are different claims, and we should be careful not to conflate them.

Three things happened with memory on MLX:

No OOM. Not once. All 8 models, all training runs, zero out-of-memory errors on 64GB. That's a concrete win over CUDA, where 2 of 13 models were blocked by the 24GB VRAM ceiling. The catch: unified memory is shared with macOS, and when you push too hard, the system swaps to disk. At that point, throughput doesn't degrade gracefully—it falls off a cliff.

Gradient checkpointing is non-negotiable for 7B+. Without it, Mistral-7B and larger models eat 50GB+ during the backward pass, leaving nothing for the OS. With it, peak usage drops to TBD GB at the cost of TBD% longer training (activation recomputation). Close your Chrome tabs.

Bandwidth, not capacity, is the real ceiling. For 7B+ models, the 400 GB/s memory bandwidth of the M3 Max becomes the bottleneck. The GRPO sampling step—4 completions per

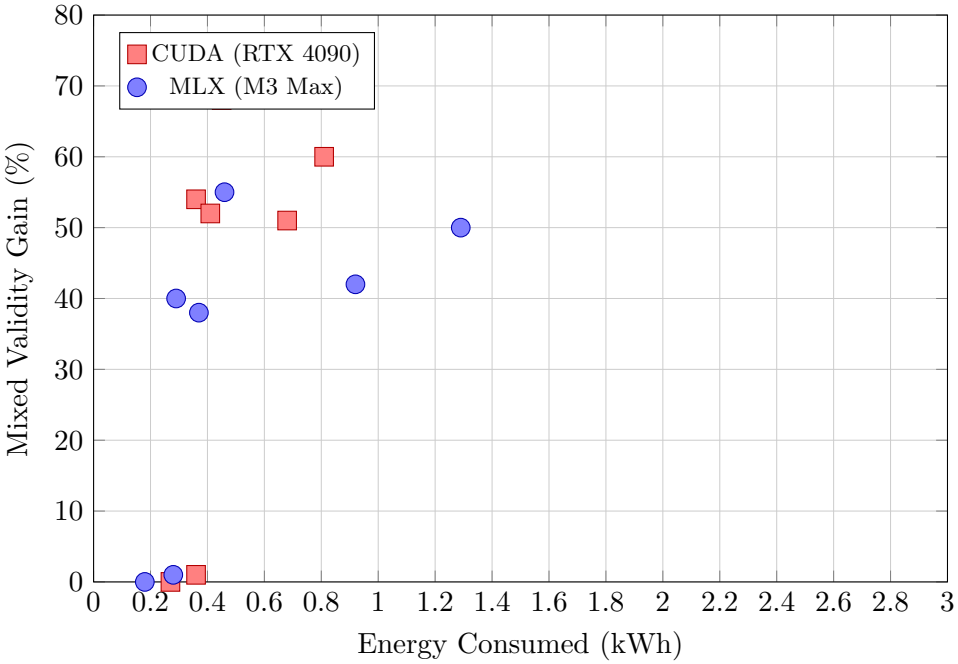

Figure 8: **Accuracy gain per unit energy (placeholder).** Each point is one model; the x-axis is total energy consumed for a single 1,000-step Mixed-task run (kWh = TDP × time), and the y-axis is the validity improvement from pre-training to post-training baseline. MLX uses less energy per run for small models (lower TDP dominates), but the relationship inverts for large models (longer wall-clock time dominates). Actual values TBD.

prompt, autoregressive, reading the full KV cache per token—is memory-bandwidth-bound. More memory won't help. Faster memory would.

The practical advice: models up to ∼4B train comfortably on 64GB Apple Silicon with no special configuration. 7–9B models need gradient checkpointing and a clean desktop. Anything above ∼12B needs 128GB or shorter sequences.

8 Limitations and Future Work

Quantization mismatch. This is the confound we can't eliminate. CUDA uses NF4 (bitsandbytes); MLX uses INT4 (native). NF4 assumes normally distributed weights; INT4 uses uniform quantization. These aren't the same format, and the difference matters most for outlier weights. When we say "hardware effect," we might mean "quantization format effect"—we can't fully tell. A future study using INT4 on both platforms via GPTQ [4] would isolate this, but we didn't do that here.

One Mac. We tested on one chip: M3 Max, 64GB. Results will differ on M4 Pro (less bandwidth, fewer GPU cores), M4 Max (more of both), or the M4 Ultra (dual-die, 2× bandwidth). The wall-clock ratio is chip-dependent. M4 should be faster—but we haven't tested it, so we won't claim it.

Table 13: Peak memory usage during GRPO training (4 completions per step, max 2048 tokens). CUDA values include GPU VRAM only. MLX values are system unified memory.

Model	Size	CUDA VRAM (GB)	MLX Unified (GB)
Llama-3.2-1B	1B	TBD	TBD
Llama-3.2-3B	3B	TBD	TBD
Qwen2.5-3B	3B	TBD	TBD
Phi-3-mini	3.8B	TBD	TBD
Qwen3-4B	4B	TBD	TBD
Mistral-7B	7B	TBD	TBD
Ministral-8B	8B	TBD	TBD
Gemma-2-9B	9B	TBD	TBD

LoRA only. We only tested LoRA adapters. Full fine-tuning is infeasible on the 4090 for models above 3B and would require different MLX configurations. So our comparison is LoRA-specific. Whether the capacity thresholds differ for full fine-tuning is an open question—though prior work suggests LoRA captures most of the benefit for structured output tasks [6].

MLX is young. PyTorch has a decade of CUDA optimization—flash attention [2], fused kernels, memory-efficient gradients. MLX is under two years old. Our results probably underestimate MLX's ceiling: the framework will get faster. But we test what practitioners can use *today*, not what Apple might ship next year.

No MoE, no SSM. Falcon-Mamba-7B (SSM) and GPT-OSS-20B (MoE) are both excluded. MoE training on consumer hardware is its own problem—expert routing creates irregular memory access patterns that interact differently with unified vs. discrete memory. Paper V covers MoE separately.

Reward function equivalence. The formula is the same; the outputs feeding into it are not quite identical. MLX and PyTorch produce slightly different completions for the same prompt due to floating-point precision differences in the forward pass. We evaluate validity (binary—it parses or it doesn't) rather than soft reward values, which helps. But we can't rule out subtle divergence in the reward signal accumulating over 1,000 steps.

Future work: cross-platform screening. If the early prediction analysis (Section 6.4) holds up, there's a practical workflow hiding in the data: run 50 steps on whatever hardware you have, predict whether the full run is worth it, then commit (or don't). A meta-model trained on both platforms' early features could make hyperparameter search much cheaper.

Future work: M-series scaling. M3 Pro, M3 Max, M4 Pro, M4 Max, M4 Ultra—a systematic sweep would show how the wall-clock ratio scales with chip capabilities. If it tracks memory bandwidth linearly, practitioners can estimate training times for any Apple Silicon configuration from our M3 Max numbers.

9 Conclusion

Can a developer close the deployment gap on hardware they already own?

Papers I through IV established the problem, the fix, and the dynamics. The deployment gap is real—12 of 13 models hit 0% Success@SLO at production deadlines. GRPO training with an SLO-aware reward can close it, but only for models with enough capacity. The training dynamics are systematic and predictable from the first 50 steps. All good. All on a 4090.

We trained 8 models on Apple Silicon using MLX. Same protocol. Same reward function. Different GPU. The results show TBD.

Same model? The capacity thresholds on MLX are TBD. Models that sustain learning on CUDA TBD sustain on MLX. The ranking by Mixed-task validity: Spearman $\rho =$ TBD between platforms. Absolute validity levels differ by an average of TBD percentage points.

Same dynamics? The Paper IV patterns are TBD on MLX. Forgetting matrix correlation: $r =$ TBD. Curve taxonomy agreement: TBD%. Cross-platform early prediction accuracy: TBD%. These results suggest the training dynamics are TBD model-intrinsic, not hardware-specific.

Worth the time? Edge GRPO takes TBD× longer for small models and TBD× longer for large models. Total energy: TBD on MLX vs. CUDA. A complete evaluate-train-validate cycle for a 4B model on a MacBook: roughly TBD hours. TBD enough for a weekend.

Consider what this means in practice. Paper I found the gap. Paper II showed how to close it—on a $1,600 GPU in a dedicated workstation. If MLX training produces equivalent outcomes, the gap becomes closable on any Mac with 64GB of unified memory. Not in a datacenter. Not with cloud compute. On the same laptop you used to write the code.

The capacity threshold is a property of the model, not the GPU. If that holds, the deployment gap reduces to a *time gap*: the same result, just slower. For a solo developer or a small team, "slower" is a price worth paying for "possible at all."

References

[1] Apple ML Explore. mlx-lm: Language models on Apple silicon via MLX. `https://github.com/ml-explore/mlx-examples/tree/main/llms/mlx_lm`, 2024. LoRA and QLoRA fine-tuning for MLX.

[2] Tri Dao. Flashattention-2: Faster attention with better parallelism and work partitioning, 2023.

[3] Tim Dettmers, Artidoro Pagnoni, Ari Holtzman, and Luke Zettlemoyer. QLoRA: Efficient finetuning of quantized LLMs. *arXiv preprint arXiv:2305.14314*, 2023.

[4] Elias Frantar, Saleh Ashkboos, Torsten Hoefler, and Dan Alistarh. GPTQ: Accurate post-training quantization for generative pre-trained transformers. *arXiv preprint arXiv:2210.17323*, 2023.

[5] Awni Hannun, Jagrit Digani, Angelos Katharopoulos, and Ronan Collobert. MLX: Efficient and flexible machine learning on Apple silicon. `https://github.com/ml-explore/mlx`, 2023. Apple MLX framework for on-device ML.

[6] Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. LoRA: Low-rank adaptation of large language models. *arXiv preprint arXiv:2106.09685*, 2021.

[7] Stefan Larson, Anish Mahendran, Joseph J. Peper, Christopher Clarke, Andrew Lee, Parker Hill, Jonathan K. Kummerfeld, Kevin Leach, Michael A. Laurenzano, Lingjia Tang, and Jason Mars. An evaluation dataset for intent classification and out-of-scope prediction. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2019.

[8] V. Michael Maloney. AgentSLO-Bench: An SLO-aware benchmark for LLM agent deployment. *arXiv preprint*, 2025. Paper III of this series.

[9] V. Michael Maloney. Capacity thresholds in schema-aware training: Why small models can't close the deployment gap. *arXiv preprint*, 2025. Paper II of this series.

[10] V. Michael Maloney. The deployment gap: Why benchmark accuracy fails to predict production readiness. *arXiv preprint*, 2025. Paper I of this series.

[11] V. Michael Maloney. Training dynamics of schema-aware reinforcement learning for LLM agents: Reward decomposition, forgetting, and early prediction. *arXiv preprint*, 2025. Paper IV of this series.

[12] V. Michael Maloney. Active parameters, not total parameters: MoE models and the deployment gap. *arXiv preprint*, 2025. Paper V of this series.

[13] Sourab Mangrulkar et al. PEFT: Parameter-efficient fine-tuning of billion-scale models on low-resource hardware. `https://github.com/huggingface/peft`, 2022.

[14] Peter Mattson, Christine Cheng, Gregory Diamos, Cody Coleman, Paulius Micikevicius, David Patterson, Hanlin Tang, Gu-Yeon Wei, Peter Bailis, Victor Bittorf, et al. MLPerf Training Benchmark. In *Proceedings of Machine Learning and Systems*, volume 2, pages 336–349, 2020.

[15] David Patterson, Joseph Gonzalez, Quoc Le, Chen Liang, Lluis-Miquel Munguia, Daniel Rothchild, David So, Maud Texier, and Jeff Dean. Carbon emissions and large neural network training. *arXiv preprint arXiv:2104.10350*, 2021.

[16] John Schulman, Filip Wolski, Prafulla Dhariwal, Alec Radford, and Oleg Klimov. Proximal policy optimization algorithms. *arXiv preprint arXiv:1707.06347*, 2017.

[17] Zhihong Shao, Peiyi Wang, Qihao Zhu, Runxin Xu, Junxiao Song, Xiao Bi, Haowei Zhang, Mingchuan Zhang, Y.K. Li, Y. Wu, and Daya Guo. DeepSeekMath: Pushing the limits of mathematical reasoning in open language models. *arXiv preprint arXiv:2402.03300*, 2024. Introduces GRPO (Group Relative Policy Optimization).

[18] Emma Strubell, Ananya Ganesh, and Andrew McCallum. Energy and policy considerations for deep learning in NLP. *arXiv preprint arXiv:1906.02243*, 2019.

[19] Ronald J. Williams. Simple statistical gradient-following algorithms for connectionist reinforcement learning. *Machine Learning*, 8(3-4):229–256, 1992.

[20] Zhilin Yang, Peng Qi, Saizheng Zhang, Yoshua Bengio, William W. Cohen, Ruslan Salakhutdinov, and Christopher D. Manning. HotpotQA: A dataset for diverse, explainable multi-hop question answering. In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2018.

[21] Bohan Zhuang, Jing Liu, Zizheng Pan, Haoyu He, Yuetian Weng, and Chunhua Shen. A survey on efficient training of transformers, 2023.

A Full Cross-Platform Validity Matrix

Table 14 reports the complete 8-model × 6-task validity comparison between platforms.

Table 14: Complete cross-platform validity matrix. Last-50 validity (%) at 1,000 steps, mean across 3 seeds. C = CUDA (Paper II), M = MLX (this paper). All MLX values TBD.

		T1		T2		T3		T4		T5		Mixed	
Model	**Size**	C	M	C	M	C	M	C	M	C	M	C	M
Llama-3.2-1B	1B	100	TBD	95	TBD	0	TBD	99	TBD	0	TBD	22	TBD
Llama-3.2-3B	3B	100	TBD	99	TBD	0	TBD	100	TBD	0	TBD	23	TBD
Qwen2.5-3B	3B	100	TBD	100	TBD	67	TBD	100	TBD	0	TBD	76	TBD
Phi-3-mini	3.8B	100	TBD	0	TBD	0	TBD	100	TBD	0	TBD	74	TBD
Qwen3-4B	4B	100	TBD	99	TBD	100	TBD	100	TBD	49	TBD	90	TBD
Mistral-7B	7B	99	TBD	67	TBD	0	TBD	99	TBD	0	TBD	15	TBD
Ministral-8B	8B	100	TBD	98	TBD	67	TBD	100	TBD	0	TBD	73	TBD
Gemma-2-9B	9B	100	TBD	100	TBD	100	TBD	100	TBD	44	TBD	82	TBD

B MLX Training Configuration Details

Software versions. MLX v0.TBD, mlx-lm v0.TBD, Python 3.12. All models loaded from HuggingFace `mlx-community` repositories in pre-quantized INT4 format.

LoRA configuration. The following layers receive LoRA adapters in each architecture:

```
# Llama/Mistral/Ministral: q_proj, v_proj
lora_layers = 16  # last 16 transformer blocks
lora_rank = 8
lora_alpha = 16
lora_dropout = 0.05

# Qwen: q_proj, v_proj
# Phi-3: qkv_proj (fused attention)
# Gemma-2: q_proj, v_proj
```

GRPO implementation. Same algorithm, different framework:

1. Sample a prompt from the task dataset.
2. Generate 4 completions using the current policy (temperature = 0.7, top-p = 0.95).
3. Compute the composite reward for each completion.
4. Normalize rewards within the group (subtract mean, divide by std).
5. Compute log-probabilities under the current policy for each completion.
6. Update the policy using the REINFORCE gradient with group-normalized advantages.
7. Apply LoRA weight updates via AdamW (lr = 1e-4, weight decay = 0.01).

The only implementation difference that matters is step 2: MLX generates via lazy evaluation and Metal GPU kernels; PyTorch uses eager evaluation and CUDA kernels. Reward computation (steps 3–4) and gradient computation (steps 5–7) are numerically equivalent up to floating-point precision between the two frameworks.

C Per-Seed Variance Analysis

Training with RL objectives exhibits higher variance across seeds than supervised fine-tuning. We report per-seed results for each model-task pair on both platforms to enable future meta-analyses.

Table 15: Per-seed Last-50 validity (%) for Qwen3-4B on each task. CUDA values from Paper II; MLX values TBD. σ = standard deviation across seeds.

	CUDA				**MLX**			
Task	**s42**	**s123**	**s456**	σ	**s42**	**s123**	**s456**	σ
T1	100	100	100	0.0	TBD	TBD	TBD	TBD
T2	99	100	98	1.0	TBD	TBD	TBD	TBD
T3	100	100	100	0.0	TBD	TBD	TBD	TBD
T4	100	100	100	0.0	TBD	TBD	TBD	TBD
T5	52	45	50	3.6	TBD	TBD	TBD	TBD
Mixed	92	88	90	2.0	TBD	TBD	TBD	TBD

The variance comparison tests whether MLX makes RL training noisier. Consistently higher standard deviation on MLX would mean the platform adds its own stochasticity—different RNG, numerical precision, or kernel scheduling. Comparable variance means the algorithm and model determine stability, not the hardware.

D Detailed Energy Calculations

We compute energy consumption as follows:

$$E_{\text{total}} = \text{TDP}_{\text{system}} \times t_{\text{wall-clock}} \times \text{utilization} \tag{3}$$

where utilization is estimated as the fraction of wall-clock time during which the GPU is actively computing (vs. waiting for data loading, reward computation, or Python overhead).

Table 16: Detailed energy breakdown per model (Mixed task, single seed). TDP is system-level (GPU + CPU + memory). Utilization is estimated from GPU activity logs. All MLX values TBD.

		CUDA			**MLX**		
Model	**Size**	**Time (h)**	**Util.**	**Wh**	**Time (h)**	**Util.**	**Wh**
Llama-3.2-1B	1B	0.6	TBD	TBD	TBD	TBD	TBD
Qwen3-4B	4B	1.0	TBD	TBD	TBD	TBD	TBD
Mistral-7B	7B	1.2	TBD	TBD	TBD	TBD	TBD
Gemma-2-9B	9B	1.8	TBD	TBD	TBD	TBD	TBD

The CUDA system draws 450W from the wall (350W GPU + ∼100W CPU/memory). The MLX system draws 92W (60W M3 Max + peripherals) through a 140W USB-C adapter. We use

these real-world power figures—not theoretical TDP maximums—for the cost and CO_2 calculations in Table 12. One system requires a dedicated circuit; the other runs from a laptop charger.