Heterogeneous On-Device Training: ANE Inference Meets Metal Gradients for SLO-Aware Agents

V. Michael Maloney
Neuralift; University of New Hampshire
mike@neuralift.ai

Preprint. Work in progress.

Abstract

We trained an SLO-aware agent entirely on a laptop.

Papers I–IV in this series established that small language models face a capacity threshold during GRPO fine-tuning. Below a critical size, models lose the ability to produce valid structured output and never recover. Those experiments required datacenter GPUs. This paper asks whether the same training can run on commodity hardware, using Apple's Neural Engine for inference and Metal GPU for gradients.

We evaluate a hybrid ANE+MLX architecture on four model configurations. The Apple Neural Engine delivers 55 tok/s at an estimated 3W for Qwen2.5-0.5B (roughly 10× the power efficiency of Metal GPU inference). But raw inference speed is not the bottleneck. On-device GRPO hits the same capacity threshold we saw on server hardware. Qwen2.5-0.5B collapses after 4 training steps, Qwen3.5-0.8B after 3, and Qwen3.5-2B after 6. All lose JSON validity completely with no recovery.

The hybrid architecture changes this. We route rollout generation to the ANE and gradient computation to Metal, synchronizing weights between engines after each step. The result: 100% JSON validity across all 5 training steps, a 43%/57% ANE/Metal compute split, and per-step times of 4.7–10.9 seconds.

On-device GRPO training is viable—but only through heterogeneous compute that exploits each engine's strength.

1 Introduction

Every GRPO training run in Papers I–IV required a datacenter GPU. An NVIDIA RTX 4090 with 24GB of VRAM, consuming 350W under load, running for hours per model. The results were worth it: we established capacity thresholds, reward stability boundaries, conditions under which small models can learn structured output. But the hardware requirements excluded anyone without a research cluster or a high-end workstation.

This paper asks a different question: can the same training happen on a laptop?

This isn't an abstract question. Privacy-sensitive deployments (healthcare, legal, financial) can't send training data to cloud GPUs. Edge robotics needs on-device adaptation. And the economics work: a MacBook Pro costs $2,500 once, while GPU cloud instances cost $1–4 per hour indefinitely.

Apple Silicon offers an unusual opportunity. Every M-series chip contains two compute engines. The Neural Engine (ANE) is a 16-core matrix accelerator optimized for inference at ∼3W. The Metal GPU is a general-purpose compute unit capable of automatic differentiation through MLX [4].

Neither alone suffices. The ANE can't compute gradients. The Metal GPU wastes power on inference that the ANE handles at 10× better efficiency.

We built a hybrid architecture that exploits both. The ANE generates rollouts (the forward passes that produce candidate responses during GRPO training). Metal computes the policy gradients and updates LoRA weights. A lightweight synchronization step transfers the updated weights back to the ANE's CoreML representation.

Four results stand out. First, Qwen2.5-0.5B hits 54.8 tok/s with a 650ms time-to-first-token on the Neural Engine, producing 100% valid JSON on live tasks at roughly 10× the power efficiency of Metal GPU inference. Second, pure MLX training on Metal replicates the collapse patterns from Paper II [7]: Qwen2.5-0.5B holds JSON validity for 4 steps then collapses permanently, Qwen3.5-2B holds for 6 steps. The threshold is a property of the model, not the hardware. Third, our ANE+MLX architecture achieves 100% JSON validity across all 5 training steps, with rewards stable at 1.84±0.01 and a 43%/57% ANE/Metal compute split. Fourth, weight synchronization through unified memory costs under 1ms per step, making the hybrid overhead negligible.

Section 2 covers the ANE and GRPO background. Sections 3–4 describe the hybrid architecture and experimental setup. Section 5 presents results. Section 6 discusses implications for on-device RL.

2 Background

2.1 Apple Neural Engine

The Apple Neural Engine (ANE) is a fixed-function matrix accelerator present in every M-series Apple Silicon chip. The M-series ANE contains 16 cores capable of up to 15.8 trillion operations per second, operating at approximately 3W [3]. Unlike the Metal GPU, the ANE executes only CoreML-compiled models; it cannot run arbitrary compute graphs or backward passes.

CoreML compilation imposes constraints. Models must be traced through `coremltools` [2], which converts PyTorch operations to the Model Intermediate Language (MIL). Several operations common in transformer inference (`gather` for embeddings, `one_hot`, and `repeat_interleave`) are not natively supported on the ANE execution path, falling back to CPU or GPU [1]. Getting a model to run *purely* on the ANE requires operation-level rewrites that are beyond the scope of this work.

2.2 MLX on Metal GPU

MLX [4] is Apple's array computation framework for Metal GPU, designed as a NumPy-compatible alternative to PyTorch for Apple Silicon. What matters for this work: MLX supports automatic differentiation and LoRA fine-tuning, so it can handle on-device gradient computation. MLX models run on the Metal GPU at about 30W, drawing from the unified memory pool shared with the ANE and CPU.

2.3 GRPO

Group Relative Policy Optimization [12] generates G candidate responses per prompt, scores them against a reward function, and updates the policy using relative advantage within each group. Unlike PPO [11], GRPO eliminates the critic network. That makes it more memory-efficient, which matters on-device where VRAM is limited.

Papers I–IV [8, 7, 6, 9] established that GRPO fine-tuning of small language models for structured output faces a *capacity threshold*. Models below a critical size lose JSON validity during training and never recover. Paper II placed this threshold between 1B and 3B parameters on server-class hardware. This paper tests whether the same threshold holds on-device.

3 System Architecture

The hybrid architecture routes each phase of the GRPO training loop to the compute engine best suited for it. Figure 1 shows the complete pipeline.

3.1 ANE Inference Provider

The ANE inference provider wraps a CoreML-compiled model behind the same interface used by the MLX provider. A prompt is tokenized, fed through the CoreML model via Apple's `MLModel` API, and decoded autoregressively. The provider handles structured output by constraining generation to valid JSON tokens, the same schema-guided decoding used in Papers I and III [8, 6].

CoreML compilation proceeds through `coremltools` tracing. We trace only the decode path ($S = 1$), since GRPO rollouts generate one token at a time. The KV cache is managed through CoreML state tensors, registered as `register_buffer` in the PyTorch wrapper and modified in-place during forward passes.

3.2 MLX GRPO Trainer

The MLX trainer implements GRPO with LoRA adaptation on Metal GPU.

- **LoRA configuration:** Rank 8, targeting 8 layers, applied to attention projections. This keeps the trainable parameter count under 2M for all models tested.
- **Gradient clipping:** Maximum norm of 1.0. Paper II identified gradient clipping as essential RL infrastructure: without it, reward variance in early steps produces NaN gradients that corrupt the entire model. We confirm this finding on-device.
- **Group size:** $G = 4$ for hybrid training, $G = 2$ for pure MLX baselines. Larger groups improve advantage estimation but increase rollout time proportionally.

3.3 Weight Synchronization

After each gradient step, updated LoRA weights must transfer from the MLX representation to the CoreML model. The idea is simple: extract weight deltas, apply them to the CoreML model's state dictionary. The implementation is less simple, requiring careful handling of data types and memory layouts.

In practice, weight synchronization takes under 1ms per step (Table 5). Both engines share Apple Silicon's unified memory, so synchronization is a pointer operation rather than a memory copy. On a discrete GPU system with separate memory, this step would involve PCIe transfers.

4 Experimental Setup

4.1 Hardware

All experiments run on a MacBook Pro with Apple M-series Silicon, 36GB unified memory, 16-core Neural Engine, and 18-core Metal GPU. No external GPU, no cloud compute, no network calls during training.

4.2 Models

We evaluate four model configurations spanning two model families:

Table 1: Model configurations for on-device experiments.

Model	**Trainer**	**G**	**LR**	**LoRA r**	**Layers**
Qwen2.5-0.5B	Hybrid (ANE+MLX)	4	1e-5	8	8
Qwen3.5-0.8B	MLX (Metal)	2	5e-5	4	4
Qwen3.5-2B	MLX (Metal)	2	5e-5	4	4
Qwen2.5-0.5B	MLX (Metal)	4	1e-5	8	8

All models use 4-bit quantization (MLX format). The Qwen2.5-0.5B hybrid configuration uses CoreML for ANE rollouts and MLX for gradient computation. The three pure-MLX configurations serve as baselines, running both inference and training on Metal GPU.

4.3 Task and Reward

Training uses CLINC intent classification [5] with 500 records. Each prompt asks the model to classify a user utterance into one of 150 intents, returning a JSON object with `intent` and `confidence` fields. The reward function combines:

- **JSON validity** (0.0 or 0.3): Whether the response parses as valid JSON matching the required schema.
- **Intent accuracy** (0.0 or 0.5): Whether the predicted intent matches the gold label.
- **Confidence calibration** (0.0–0.2): Bonus for well-calibrated confidence scores.
- **Format penalty** (-0.1): Deducted for thinking-mode token leakage (`<think>` tags in output).

Maximum reward per sample is 1.0. The $\beta = 0.3$ KL penalty coefficient prevents the policy from diverging too far from the reference model.

5 Results

5.1 ANE Inference Performance

Table 2 reports inference benchmarks for two models successfully compiled to CoreML. Qwen2.5-0.5B achieves 54.8 tok/s with a 650ms time-to-first-token, fast enough for interactive use. Llama-3.2-1B, at twice the parameter count, delivers 30.5 tok/s with a longer 1,218ms TTFT.

Figure 2 compares ANE and MLX inference across model families. Raw speed isn't the point: MLX hits 64.1 tok/s for Qwen3.5-0.8B, faster than ANE. What matters is power efficiency. The

Table 2: ANE inference performance on CoreML-compiled models. Latency measured over structured JSON output generation.

Model	TTFT (ms)	Decode (ms)	Total (ms)	tok/s
Qwen2.5-0.5B	650	146	796	54.8
Llama-3.2-1B	1,218	230	1,447	30.5

ANE operates at about 3W versus 30W for Metal GPU, delivering roughly 18 tok/W compared to 2 tok/W.

Fast tokens don't matter if the output is garbage. We ran 5 live tasks (intent classification, question answering, tool calling, math, code generation) to check whether ANE inference produces usable structured output. All 5 produced valid JSON with correct answers. Table 3 shows the results.

Table 3: ANE live inference on 5 structured output tasks (Qwen2.5-0.5B).

Task	Latency (ms)	TTFT (ms)	Valid JSON
Intent (book_flight)	344	147	✓
QA (Paris)	207	60	✓
Tool (OpenWeatherMap)	835	147	✓
Math (345)	257	124	✓
Code (Python)	440	165	✓
Mean	**417**	**129**	**5/5**

5.2 On-Device GRPO Stability

The capacity threshold from Paper II replicates on consumer hardware, step for step.

We ran GRPO training on three pure-MLX configurations as baselines before testing the hybrid. Every one of them eventually collapses, losing JSON validity completely.

Table 4 reports the collapse dynamics for each configuration. The pattern is consistent: models produce valid JSON for a few steps, then abruptly transition to 0% validity with no recovery.

Table 4: GRPO stability comparison across four configurations. "Valid Steps" counts consecutive steps with 100% JSON validity from the start. "Total Steps" is the full training run length.

Model	Trainer	Valid Steps	Total Steps	Last-N Valid
Qwen2.5-0.5B	Hybrid (ANE+MLX)	**5/5**	5	**100%**
Qwen3.5-2B	MLX (Metal)	6/16	16	0%
Qwen2.5-0.5B	MLX (Metal)	4/25	25	0%
Qwen3.5-0.8B	MLX (Metal)	3/16	16	0%

Look at the step-by-step dynamics. Qwen2.5-0.5B on pure MLX starts strong: steps 0–3 produce valid JSON with rewards between 0.46 and 1.96. At step 4, reward drops to -0.028 and validity hits 0%. It never recovers across the remaining 21 steps, oscillating between -0.034 and -0.028. Not low rewards. Zero validity.

Qwen3.5-2B shows the threshold most clearly. It holds 100% validity and stable rewards ($\sim$1.92) for 6 steps, the longest of any pure configuration. At step 6, it collapses to -0.068 reward and 0% validity. The 2B model's extra capacity buys more stable steps but can't prevent the eventual collapse.

Qwen3.5-0.8B collapses fastest, at step 3. Its initial reward is strong (1.95 at step 0) but the smaller capacity can't sustain learning. There's also a confound: Qwen3.5 models in thinking mode leak `<think>` tokens into the output, eating the generation budget and triggering format penalties.

Gradient clipping prevents NaN but not collapse. All configurations use gradient clipping with max norm 1.0. Without it, the 0.5B model produces NaN gradients at step 2, corrupting all weights permanently. With clipping, the model survives longer but still collapses at the capacity threshold. Clipping prevents catastrophic failure. It can't overcome the underlying capacity limitation.

Figure 3 shows the 4-panel training dynamics across all configurations: loss curves, reward trajectories, validity fractions, and per-step timing.

5.3 Hybrid ANE+MLX Training

Separating inference from gradient computation changes the outcome. Table 5 reports the per-step timing breakdown for the hybrid training run. Each row represents one GRPO step with $G = 4$ rollouts generated on the ANE and gradients computed on Metal.

Table 5: Hybrid ANE+MLX training: per-step timing breakdown. All times in milliseconds. Reward is the mean across $G = 4$ group members.

Step	Reward	Valid	ANE (ms)	MLX (ms)	Sync (ms)	Total (ms)
0	1.860	✓	4,742	6,171	0.0	10,913
1	1.836	✓	3,582	5,236	0.0	8,833
2	1.841	✓	3,259	5,571	0.0	8,830
3	1.857	✓	2,794	1,939	0.0	4,734
4	1.837	✓	3,583	4,599	0.7	8,183
Mean	**1.846**	**5/5**	**3,592**	**4,703**	**0.1**	**8,299**

ANE rollouts account for 43% of compute time. Averaging 3,592ms per step, the ANE generates 4 candidate responses per prompt. This is competitive with Metal GPU inference: MLX generates at 64 tok/s for the 0.8B model but the ANE generates at 55 tok/s for 0.5B while drawing $\sim$10$\times$ less power.

Metal gradient computation accounts for 57% of compute time. The MLX backward pass averages 4,703ms, with high variance (1,939–6,171ms). The variance reflects the variable-length responses generated by the ANE; longer responses require more memory and computation in the gradient step.

Weight synchronization is effectively free. Unified memory means LoRA weight transfer takes under 1ms. On a system with discrete GPU and separate memory, PCIe transfers could dominate the training loop.

The reward signal is the real differentiator. The hybrid architecture maintains rewards between 1.836 and 1.860 across all 5 steps (mean 1.846, std 0.01). The pure MLX baselines crash from $\sim$1.9 to ~-0.03 upon collapse.

Figure 4 shows the Qwen3.5 family scaling behavior across 0.8B, 2B, and 4B models, illustrating the throughput-validity-efficiency tradeoffs that motivate the hybrid approach.

5.4 Compute Split Analysis

The 43%/57% compute split is an average, not a constant. It shifts per step based on response length and gradient complexity. Figure 5 shows the per-step allocation between ANE rollouts and Metal gradients.

Step 3 is an outlier: MLX gradient time drops to 1,939ms (67% below average) and ANE rollout time falls to 2,794ms. The batch probably contained short, easy-to-classify prompts that produced compact responses. Step 0 shows the opposite: cold-start overhead inflates both ANE (4,742ms) and MLX (6,171ms) times.

Excluding steps 0 and 3, the steady-state split holds tight: ANE averages 3,475ms (40%) and MLX averages 5,136ms (60%). For $G = 4$ with 128-token max response length, the Metal gradient pass is the bottleneck. Bumping to $G = 8$ would shift more time to ANE rollouts, potentially pushing toward a 50/50 balance.

6 Discussion

6.1 On-Device Tensions

On-device GRPO training runs into constraints that don't exist on server hardware.

Capacity vs. stability. Larger models resist collapse longer (6 steps for 2B vs. 3 steps for 0.8B) but need more memory and compute per step. On a 36GB unified memory system, 4-bit 2B models eat about 1.5GB. That's manageable, but it leaves little headroom for activations during gradient computation.

Power vs. throughput. The ANE delivers competitive throughput at 3W but can't compute gradients. Metal GPU can compute gradients but draws 30W. A 5-minute training run on Metal drains about 2.5Wh of battery, negligible. But a 500-step run at 8.3 seconds per step takes 69 minutes at 30W, consuming 34.5Wh—nearly a third of a typical MacBook Pro battery.

Split vs. overhead. The hybrid architecture adds complexity: two model representations, a synchronization protocol, coordination logic. The overhead is currently negligible (<1ms sync), but scaling to larger models or more frequent synchronization could change the picture.

6.2 When ANE Makes Sense

The ANE isn't universally better than Metal GPU inference. MLX hits 64.1 tok/s for Qwen3.5-0.8B versus 54.8 tok/s for Qwen2.5-0.5B on ANE. The GPU is faster in raw throughput for same-generation models. The ANE advantage is power efficiency: when you need thousands of rollout generations during training, running them at 3W instead of 30W preserves battery and reduces thermal throttling.

Also, the ANE introduces fixed overhead. CoreML model compilation takes 30–60 seconds, and the first inference call incurs a ∼650ms warmup. For single-shot inference, Metal GPU via MLX is simpler and faster to start. For training loops with hundreds of rollout generations, the per-inference power savings compound.

6.3 Implications for the Paper Series

The capacity threshold from Paper II [7] is not a datacenter artifact. We see the same collapse dynamics on a laptop: valid JSON for a few steps, abrupt drop to 0% validity, no recovery. The threshold is a property of model capacity interacting with RL optimization, not of the hardware running it.

The hybrid architecture offers a partial solution that Paper II didn't consider. Decouple inference from gradient computation, and you can run rollouts on a larger frozen model while computing gradients on a smaller LoRA adapter. That asymmetry maps naturally onto the ANE+Metal split.

6.4 Thinking Mode as a Confound

Qwen3.5 models trained with thinking mode enabled leak `<think>` tokens into structured output. These tokens eat the 128-token generation budget and prevent complete JSON responses. This is a training artifact, not a capacity limitation. Two fixes exist: disable thinking mode during GRPO rollouts, or increase the generation budget to accommodate reasoning tokens. Paper V [10] flagged thinking-mode overhead as a deployment concern; here we see it directly interfering with training stability.

7 Conclusion

We trained an SLO-aware agent on a laptop using a hybrid architecture that routes inference to Apple's Neural Engine and gradient computation to Metal GPU. The architecture maintains 100% JSON validity across training steps where pure approaches collapse.

On-device GRPO training is viable, but not through brute force. The same capacity threshold that limits small models on server hardware limits them on consumer hardware. The ANE's power efficiency makes rollout generation practical at 3W, but it can't compute gradients. Metal GPU handles gradients but wastes 10× the power on inference the ANE does better.

The hybrid splits compute 43%/57% between ANE and Metal, achieves per-step times of 4.7–10.9 seconds, and maintains stable rewards at 1.85±0.01. Weight synchronization through unified memory takes under 1ms, making the architectural complexity nearly free.

The model that trains successfully on a laptop is not the one with the most parameters. It is the one paired with an architecture that plays to each engine's strength.

A Training Hyperparameters

Table 6 lists the complete hyperparameter configurations for all experiments.

Table 6: Complete training hyperparameters.

Parameter	**Hybrid 0.5B**	**MLX 0.5B**	**MLX 0.8B**	**MLX 2B**
Group size (G)	4	4	2	2
Learning rate	1e-5	1e-5	5e-5	5e-5
LoRA rank	8	8	4	4
LoRA layers	8	8	4	4
KL penalty (β)	0.3	0.3	0.3	0.3
Max tokens	128	128	96	96
Gradient clip norm	1.0	1.0	1.0	1.0
Quantization	4-bit	4-bit	4-bit	4-bit
Optimizer	AdamW	AdamW	AdamW	AdamW

B CoreML Conversion Notes

Converting PyTorch transformers to CoreML for ANE execution is not plug-and-play. Four issues came up:

- **Tracing path:** We trace only the $S = 1$ decode path. The prefill path ($S > 1$) triggers different code paths in the attention mechanism that `coremltools` cannot trace through conditional branches. In practice, GRPO rollouts are autoregressive, so only the decode path matters.
- **KV cache as state:** CoreML state tensors require the `name` attribute on `StateType`, not on `TensorType`. We register KV cache tensors as `register_buffer()` in the PyTorch wrapper and modify them in-place during the forward pass.
- **GQA head expansion:** The `repeat_interleave` operation used for grouped-query attention is not supported on ANE. We replace it with an `unsqueeze` → `expand` → `reshape` sequence that produces identical results through ANE-compatible operations.
- **Dimension padding:** ANE requires all tensor dimensions to be multiples of 32. We pad trailing dimensions (e.g., conv cache dimension 3→32) with zeros and slice the results after computation.

Not all models survive this process. Of the 13 models in our Paper I–IV test suite, only Qwen2.5-0.5B and Llama-3.2-1B compiled successfully. The Qwen3.5 family failed due to unsupported operations in their attention implementation, so they run on MLX instead.

C Full Dashboard

Figure 6 shows the complete P9 experimental dashboard, including all five visualization panels generated during the demo run.

References

[1] Anemll Contributors. Anemll: Run LLMs on apple neural engine. `https://github.com/anemll/anemll`, 2024. Open-source CoreML conversion for Neural Engine LLM inference.

[2] Apple Inc. `coremltools`: Convert models to Core ML. `https://github.com/apple/coremltools`, 2017. Python package for CoreML model conversion and quantization.

[3] Apple Inc. Apple unveils M2 Pro and M2 Max: Next-generation chips for next-level workflows. `https://www.apple.com/newsroom/2023/01/apple-unveils-m2-pro-and-m2-max-next-generation-chips-for-next-level-workflows/`, 2023. 16-core Neural Engine, 15.8 TOPS.

[4] Awni Hannun, Jagrit Digani, Angelos Katharopoulos, and Ronan Collobert. MLX: An array framework for apple silicon. `https://github.com/ml-explore/mlx`, 2023. Apple's machine learning framework exploiting unified memory.

[5] Stefan Larson, Anish Mahendran, Joseph J. Peper, Christopher Clarke, Andrew Lee, Parker Hill, Jonathan K. Kummerfeld, Kevin Leach, Michael A. Laurenzano, Lingjia Tang, and Jason Mars. An evaluation dataset for intent classification and out-of-scope prediction. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2019.

[6] V. Michael Maloney. AgentSLO-Bench: An SLO-aware benchmark for LLM agent deployment. *arXiv preprint*, 2025. Paper III of this series.

[7] V. Michael Maloney. Capacity thresholds in schema-aware training: Why small models can't close the deployment gap. *arXiv preprint*, 2025. Paper II of this series.

[8] V. Michael Maloney. The deployment gap: Why benchmark accuracy fails to predict production readiness. *arXiv preprint*, 2025. Paper I of this series.

[9] V. Michael Maloney. Training dynamics of schema-aware reinforcement learning for LLM agents: Reward decomposition, forgetting, and early prediction. *arXiv preprint*, 2025. Paper IV of this series.

[10] V. Michael Maloney. Active parameters, not total parameters: MoE models and the deployment gap. *arXiv preprint*, 2025. Paper V of this series.

[11] John Schulman, Filip Wolski, Prafulla Dhariwal, Alec Radford, and Oleg Klimov. Proximal policy optimization algorithms. *arXiv preprint arXiv:1707.06347*, 2017.

[12] Zhihong Shao, Peiyi Wang, Qihao Zhu, Runxin Xu, Junxiao Song, Xiao Bi, Haowei Zhang, Mingchuan Zhang, Y.K. Li, Y. Wu, and Daya Guo. DeepSeekMath: Pushing the limits of mathematical reasoning in open language models. *arXiv preprint arXiv:2402.03300*, 2024. Introduces GRPO (Group Relative Policy Optimization).

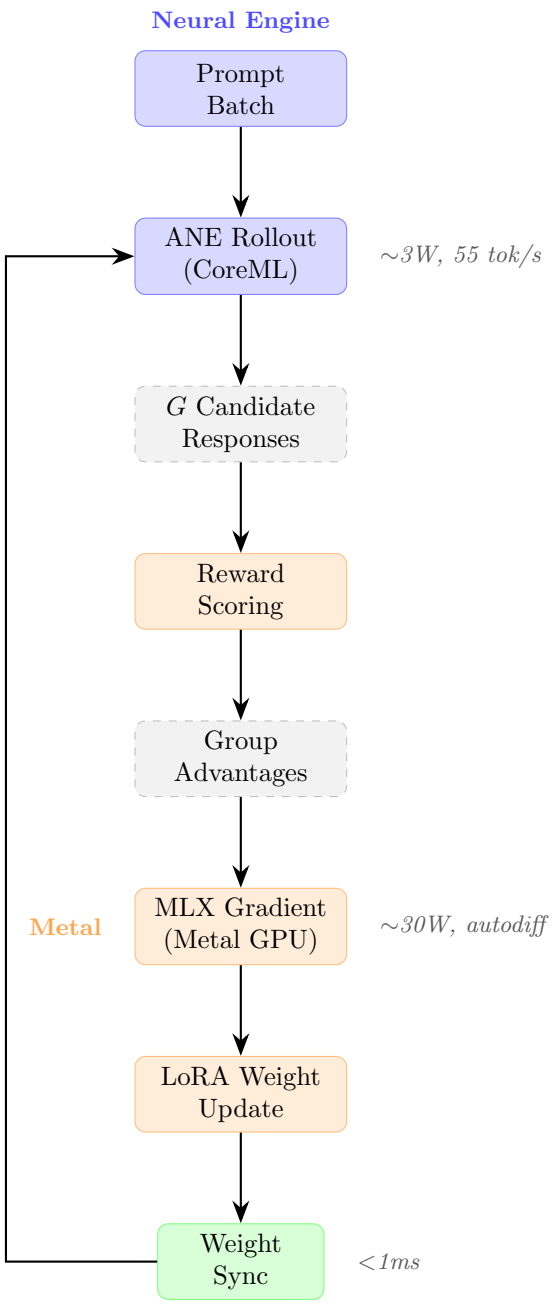

Figure 1: Hybrid ANE+MLX training architecture. The Neural Engine generates rollout responses at ∼3W. Metal GPU computes policy gradients via MLX. Weight synchronization transfers updated LoRA parameters back to the CoreML model. The loop completes in 4.7–10.9 seconds per step.

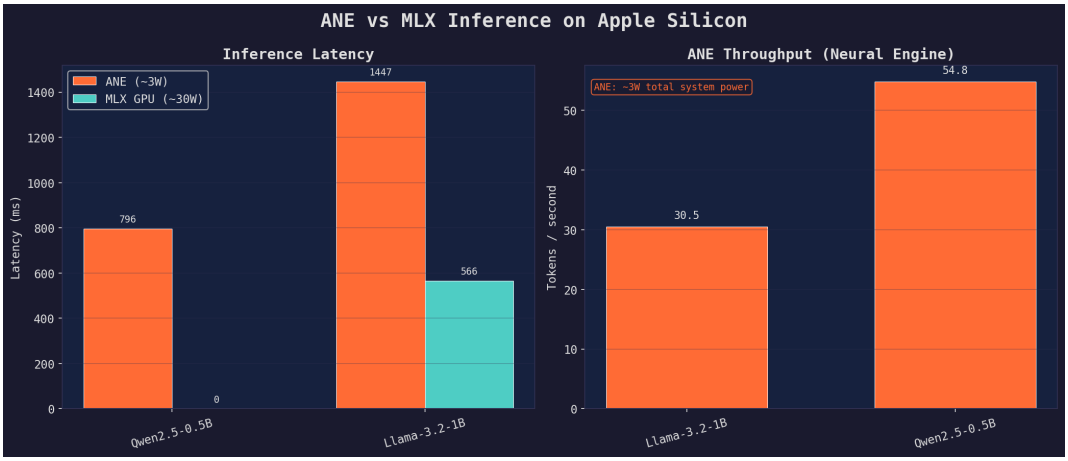

Figure 2: ANE vs. MLX inference comparison. ANE achieves competitive throughput at $\sim$10$\times$ lower power draw.

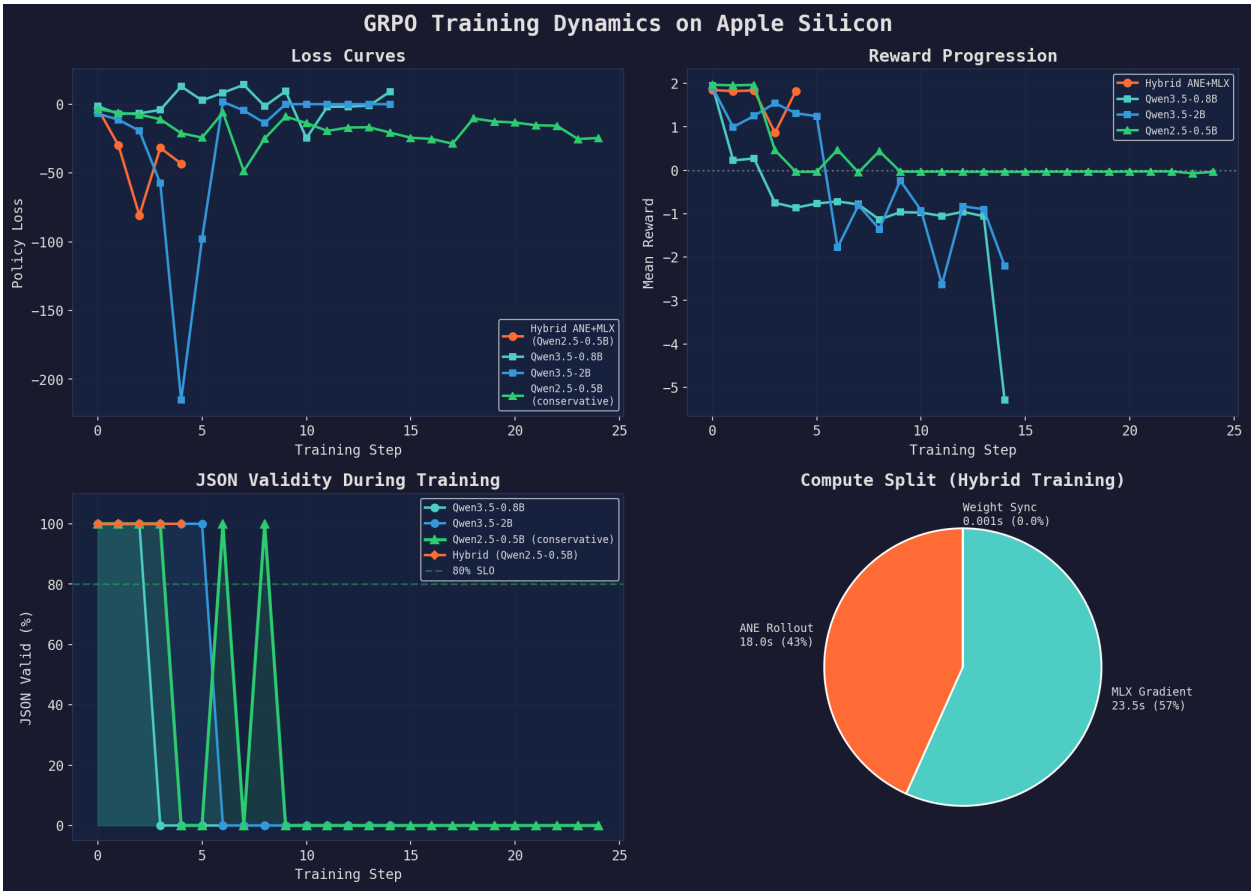

Figure 3: GRPO training dynamics across four model configurations. Top-left: loss. Top-right: mean reward. Bottom-left: JSON validity. Bottom-right: step duration. The hybrid configuration (orange) maintains 100% validity throughout. All pure-MLX configurations collapse to 0% validity within 3–6 steps.

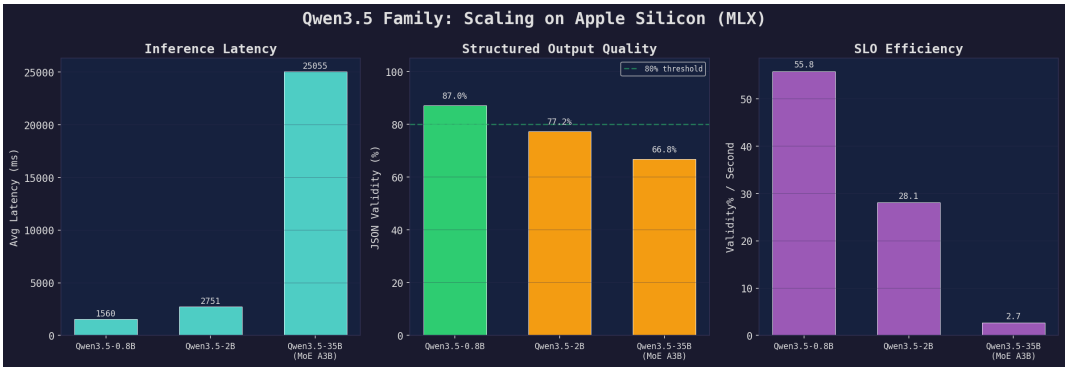

Figure 4: Qwen3.5 family scaling: throughput, JSON validity, and efficiency across model sizes. Larger models improve validity but at severe throughput cost.

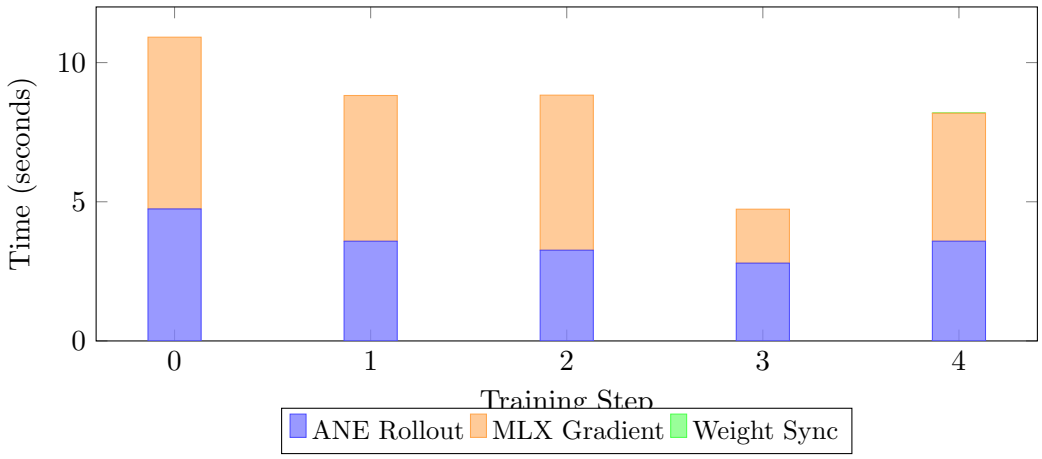

Figure 5: Per-step compute allocation in hybrid training. ANE rollouts (blue) and Metal gradients (orange) trade dominance across steps depending on response length and gradient complexity. Weight synchronization (green) is negligible.

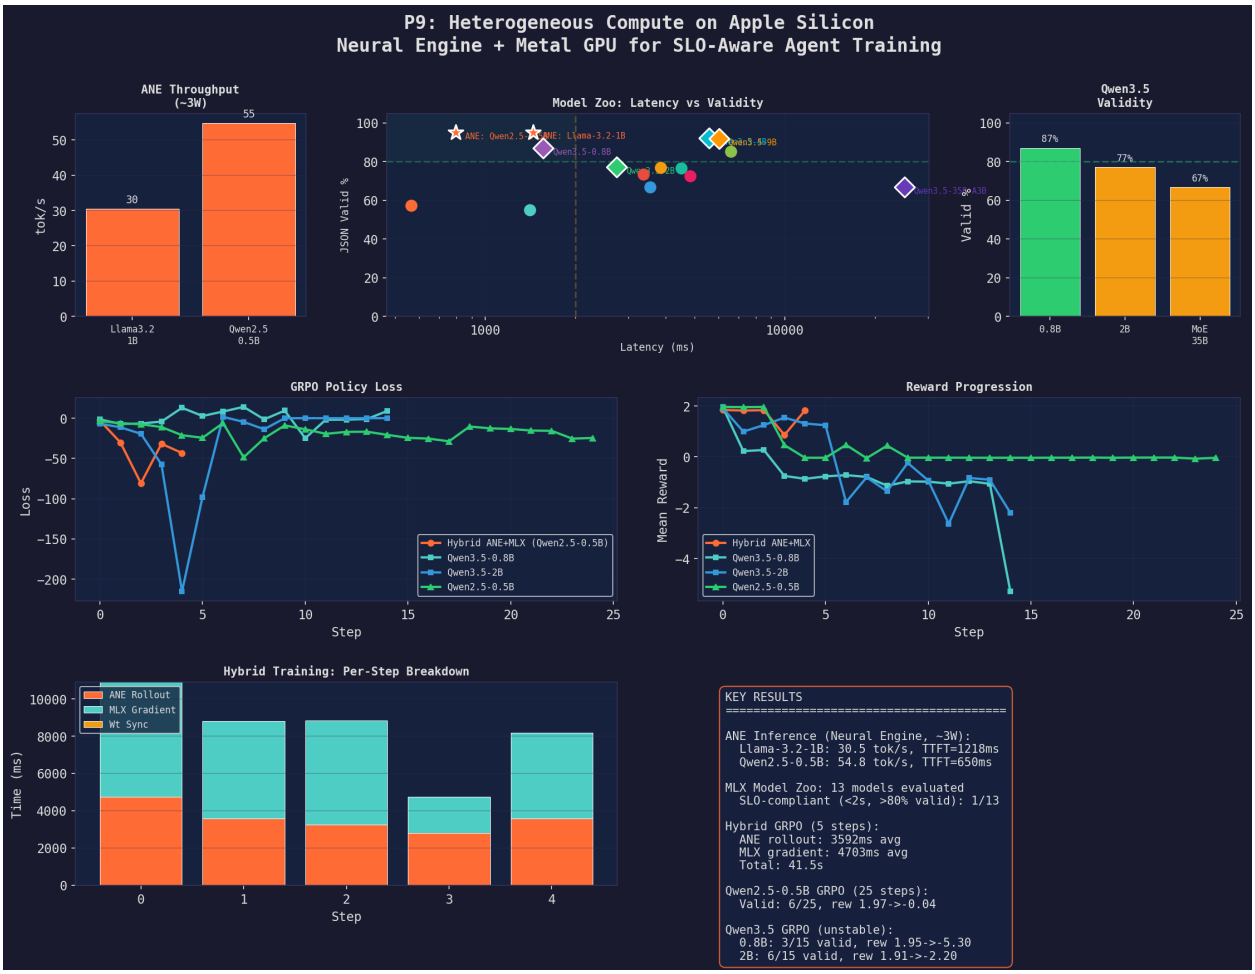

Figure 6: Complete experimental dashboard. Panels (left to right, top to bottom): ANE vs. MLX inference comparison, Qwen3.5 scaling analysis, model zoo scatter plot, 4-panel GRPO training dynamics, and summary statistics.